

Advanced Programming for Scientific Computing (PACS)

Lecture title: Template metaprogramming,
type_traits and concepts

Luca Formaggia

MOX
Dipartimento di Matematica
Politecnico di Milano

A.A. 2022/2023

Working with types

Generic programming, which in C++ translates in template programming allows to treat types as sort-of "variables".

This fact on the one hand requires tools to interrogate or manipulating types and, on the other hand, opens up the possibility of implementing more efficient code by "letting the compiler do the work".

Indeed, all operations that involve templates must be resolved at the moment of the instance of the template during the compilation process, this fact can be exploited by letting the compiler perform tasks that would be otherwise done at run-time.

Some elements of Metaprogramming

Metaprogramming is the design of computer programs that write or manipulate other programs (or themselves) as their data, or that **do part of the work at compile time that would otherwise be done at runtime**. This may allow programmers to minimize the number of lines of code to express a solution by the user (hence reducing development time), or to generate more efficient code.

In C++ we indicate with **template metaprogramming** techniques by which the “intermediate code” (that you never see!) is in fact generated by the compiler thanks to the use of C++ templates.

Template metaprogramming

Metaprogramming uses only constructs that can be resolved at compile time (static constructs). It is a small subset of C++ constructs, but sufficient to do quite a lot of stuff.

Our main aims are:

- ▶ To allow a more generic programming, by implementing a better control on types (type_traits and concepts).
- ▶ Allow to increase the efficiency of some algebraic manipulation by letting the compiler do some operations at compile time (numerical metaprogramming).

Some metaprogramming techniques relies to the **SFINAE** paradigm. This rule applies during overload resolution of function templates: *when substituting the **explicitly specified** or **deduced type** for the template parameter fails, the specialization is discarded from the overload set instead of causing a compile error.*

SFINAE

The rules for SFINAE are complex (if you wish look [here](#)), but in fact the mechanism is not so complicated, we give here an example

```
template <class T> int f(typename T::value_type);  
template <class T> int f(T);  
...  
int i = f<int>(10); // uses 2nd overload  
int j = f<vector<int>>> f(10); // uses 1st overload
```

Here, when the compiler tries in the first overload to resolve `int::value_type` has a **failure**, since `int` does not define any type called `value_type` (indeed it does not define any type at all). But this failure **is not an error** since there is an overload that works fine: indeed `std::vector` stores a type alias called `value_type`!

SFINAE is essential for generic programming to work! **remember: SFINAE works only during template parameter substitution.**

Unevaluated context and static evaluation

Other two important concepts are those of **unevaluated context** and **static** evaluation. Normally an expression is evaluated, i.e. it expresses a value:

```
c = a + b;
```

the expression "a+b" is **evaluated** and the resulting **value** is assigned to c.

But there are cases where expressions are not used to provide a value, but only information about a type. We are in an **unevaluated context**

Unevaluated context

For instance, the operands of **sizeof**, `decltype` (as well as `alignof`) and **typeid()** are **never evaluated**:

- ▶ No code is generated;
- ▶ The operand may not be an expression, but just a type.
- ▶ We need only declaration of functions or classes.

```
struct A {  
double M_x,  
int M_n  
double fun(double);  
}; // only declaration
```

```
int n = sizeof(A); //# of bytes of an object of type A  
int m = sizeof a; //# of bytes of an object of type A
```

In the last use of **sizeof**, `a` is **not evaluated**.

Static evaluation

The compiler may perform arithmetic operations if the operands are constant expressions.

```
constexpr double f(const double x){return x*x;}
```

```
...
```

```
constexpr double z=9.0;
```

```
double c = 5.0 +3.0 // it's 8.0
```

```
double d = z + f(z); // its 90!
```

The right-hand-side of last two last statements may be computed **at compile-time** (static evaluation), while here

```
std::cin >>c;
```

```
d = fun(c);
```

the computation may only be performed run-time.

Tools for template metaprogramming

We will give description of some techniques for metaprogramming. But let's start with a brief explanation of `decltype`, `decltype<T>` and `invoke_result`. `decltype` is defined the header `<utility>`, `invoke_result` in `<type_traits>`.

decltype and declval<T>

decltype(expr) returns the type of expression expr. declval<T> (in <utility>) makes it possible to deduce the return type of member functions of T using decltype, **without the need to go through constructors!**

```
class A{  
    ...  
    public:  
    double foo();  
}
```

```
...  
decltype(std::declval<A>().foo()) n; // n is an double
```

No object of type A is constructed and foo() is not called. We just get the type of the return value!. declval can only be used in an unevaluated context.

An example of use of Declval

in [MetaProgramming/DecltypeDeclVal/main_decltype.cpp](#) you find an example of the use of `decltype` and `declval<T>`.

In that example you find also an use of the type traits (see next slides) `is_constructible` and `is_default_constructible`, and of the `typeid` utility!

invoke_result, formerly result_of, now declared **deprecated**, can be used to get the return type of functions, particularly useful in the case of overloaded functions. It can also be used as an alternative to the decltype/decltype pair for member functions (with a more complex syntax).

The class

```
template< class F, class... ArgTypes> class invoke_result;
```

contains a type alias, called type with the type returned by the **callable objects** of class type F when used with arguments Args... (the list may also be empty). The type alias **is not present** if the callable object cannot be called with those argument types!

You can use invoke_result_t<> in place of invoke_result<>::type (as explained later).

Let's see an example in the next slide.

Use of `invoke_result`

```
struct S { // a callable object of class type  
// two overloads of the call operator  
    double operator()(char, int&) const;  
    int operator()(int)const;  
};  
double fun(int x);  
...  
std::invoke_result_t<S, char, int&> b; // b is a double  
std::invoke_result<S, int>::type x; // x is an int  
std::invoke_result_t<decltype(fun), int> z; // z is a double
```

Note the use of `decltype` since `fun` is not a type but the name of a function. In fact, `decltype(fun)` returns `double (*)(int)`, but maybe you prefer not to know it :).

Use of `invoke_result`

`invoke_result` may also be used to interrogate the return type of member function different from the call operator. The syntax is a bit weird:

```
struct C {  
    double Func(char, int&);  
    ...};  
...  
std::invoke_result_t<decltype(&C::Func), C, char, int&> y;  
// y is a double.
```

Explanation of the arguments: `&C::Func` is a pointer to the (non static) member function `Func` of `C`, called by an object of type `C` (the second argument) and taking `char` and `int` & as arguments.

Note: Why the second argument? Well, it accounts for inheritance. That argument may be a class derived from that indicated in the first argument!.

A note on `static_assert` VS `assert`

The differences:

`assert` available in `<cassert>` is a **cpp** macro that allows assertion checks **at run time**, if the predicate is false the program terminates. `static_assert` operates at **compile time** and, if the predicate is false you have a compiler error.

`assert` is disabled if the preprocessor variable `NDEBUG` is set. `static_assert` is (obviously) always enabled.

You cannot add a message in `assert`, but see [Utilities/ExtendedAssert.hpp](#) for some home-made extensions of `assert`

Another note: enriching messages

The preprocessor adds two macro variables: `__FILE__` and `__LINE__`, that contain the name of the file and the current line number! So you can write

```
static_assert(condition, "Error_at_line___LINE___of_file___FILE__");
```

or

```
std::cerr<<"Got_stuck_at_line___LINE__\n";
```

The [Utilities/ExtendedAssert.hpp](#) utilities take advantage of that.

Type Traits

Type traits are tools that operate on types. They are very useful in template metaprogramming and in C++ we have an extensive set of type traits into the language. [Look here](#) for the full list. You should include `<type_traits>`.

C++ Standard Type traits

C++ type-traits are template classes (actually structs), It is not important to know the details, just that they either provide a value (in fact a constant expression) or a type (or both).

If it provides a value, the trait has a static variable called `value`, containing the value.

If it provides a type, it has a type alias called `type` which gives the type.

Moreover, if the trait returns a value, it also provides `value_type` (the type of the value), and is implicitly convertible to that type (providing value).

Two goodies

To simplify life, we have `type_trait_t<T>` as an alias to `typename type_trait<T>::type`:

```
template<class T>
using type_trait_t=typename type_trait_t<T>::type;
```

It spares us the hideous `typename` for template dependent types

We also have `type_trait_v<T>`, which is equivalent to `type_trait<T>::value`:

```
template<class T>
constexpr type_trait<T>::value_type type_trait_v=
    type_trait<T>::value;
```

An example

`is_pointer<T>` may be used to interrogate if a type is a pointer.

```
#include <type_traits>
template<class T>
auto fun(T x)
{
    if constexpr(is_pointer_v<T>) // or is_pointer<T>::value
        // do something
    else
        // do something else
}
```

This function behaves differently whenever the argument is a pointer. Note the use of **if constexpr**, so the compiler will actually compile only the true block.

This use of type-traits spare us the need of writing an overload to specialize the function for pointers.

Organization of standard type_traits

- ▶ Primary type categories: `is_int<T>`, `is_pointer<T>`, `is_function<T>`, `is_rvalue_reference<T>`, `is_enum<T>`...
- ▶ Composite type categories: `is_scalar<T>`, `is_reference<T>`, `is_member_pointer<T>` ...
- ▶ Type properties: `is_const<T>`, `is_trivial<T>`, `is_abstract<T>`, `is_polymorphic<T>`...
- ▶ Supported operations: `is_copy_constructible<T>`, `is_assignable<T>`, `has_virtual_destructor<T>`...
- ▶ Property queries: `rank<T>`, `extent<T>`,...
- ▶ Type relationships: `is_base_of<B,D>`, `is_convertible<From,To>`..
- ▶ Type modifications: `remove_const<T>`, `remove_reference<T>`, `decay<T>`, `add_const<T>`, `make_unsigned<T>`...

and more... The full list in cppreference.com.

`common_type<T...>`

It determines the common type among all types `T...`, that is the type all `T...` can be implicitly converted to. If such a type exists, the member `type` names that type.

Otherwise, there is no member `type`. Useful in mixed mode arithmetic. Its use however, is less crucial now that we have automatic function where the return type is automatically deduced.

An example of `common_type<T...>`

The common type of the arguments and **double**.

```
template<class T, class U>
struct Foo{
    std::common_type_t<double, T, U> v;
    ...
};
..
```

`Foo<int, int> a;` *// a.v is a double*

`Foo<std::complex<double>, double> b;` *// b.v is complex<double>*

`Foo<long double, double> c;` *// c.v is long double*

The rules for `common_type` are rather complex, but usually it does what you expect (at least for POD).

A complex type trait: `enable_if`

```
template< bool B, class T = void > struct enable_if;
```

It is a convenient way to leverage SFINAE to conditionally remove functions from overload resolution based on type traits and to provide separate function overloads and specializations for different type traits.

`std::enable_if` can be used as an additional function argument (not applicable to operator overloads), as a return type (not applicable to constructors and destructors), or as a class template or function template parameter.

How does enable_if work?

```
enable_if<b, MyType>
```

If **b** is **true** it defines a **type** in `enable_if<b, MyType>::type`, equal to `MyType` (or **void** if `MyType` has not been given). If **b** is **false** then **it does not define ::type**. So, `enable_if<false, MyType>::type` may be used to generate a substitution failure.

A possible implementation (just to show how it can be done) is

```
template<bool B, class T = void>
struct enable_if {};
// partial specialization for B=true
template<class T>
struct enable_if<true, T>
{ using type=T; };
```

A possible use in [MetaProgramming/Enable_if/Norms.hpp](#).

An example of `enable_if`

I want that the gcd template function **exist only** if the arguments are integers. Trying to do `gcd(5.0,6.0)` should generate the compiler error: `gcd(double, double)` is not defined!.

```
template
<class T,
    std::enable_if_t
        <std::is_integral<T>::value,T > =0 >
    // NOTE THE SPACE BETWEEN > AND = !!
>
T gcd(T a, T b){ return (b == 0)? a: gcd(b, a % b);}
```

If `T` is an integer the second template parameter is an **integer defaulted to 0**. If not, `enable_if` does not define a type, and the function **is not generated**. The default value to enable calling `gcd` without specifying the second template parameter, which is used only to activate `enable_if`.

Note the use of `enable_if_t` alias to avoid **typename**.

An explanation

When a template parameter is not actually used, I can avoid giving it a name

In the above example, then, if I instantiate the function with `T=int` (it means that the argument is an `int`) the template parameters expands to

`<int , int=0>`

which is admissible.

I can give a name to the second parameter if I want to, but it is irrelevant since it is not used. The second template parameter is indeed a sort of “dummy” parameter, necessary only to activate the mechanism of `enable_if`.

Another example of `enable_if`

I have a class `transposeView<Matrix>` that provides a **view** of a `Matrix` as its transpose. I want to make sure that the non-const version of **`operator(int,int)`** used to address elements of the transposed `Matrix` **is not generated if `Matrix` is const!**

```
template<class Matrix>
class MatrixView
{
    // the type of the matrix element
    using value_type=typename Matrix::value_type;
    ...
    template<typename T=Matrix>
    std::enable_if_t<!std::is_const_v<T>,value_type &>
        operator()(size_type r, size_type c){return matrix_(c,r);}
}
```

Note the trick of using a defaulted template to enforce template parameter substitution (**otherwise it will not work: I need SFINAE!**).

The full example in [MetaProgramming/transposeView](#)

A note

The use of generic programming in C++ is increasing, since it has several advantages. Consequently, the language is evolving in order to make "template metaprogramming" techniques simpler.

Indeed, the introduction of the "compile time if" **if constexpr** can sometimes (but not always) avoid the need of `enable_if`. Also the tag dispatching technique that we will present later can now be replaced by **if constexpr**.

And the introduction of concepts may simplify life further.

Testing memory layout

Some types of traits may be used to check the memory layout of a type. We need to clarify two important concepts: **trivially-copyable** and **standard layout**.

I give just a general explanation of the terms, for the details consult any good manual or the web, like cppreference.com.

Trivially copyable and standard layout

If a type is **trivially copyable** it is possible to copy data via **memcpy**, rather than having to use the standard copy operations. In general, a trivially copyable type is a type for which the underlying bytes **can be copied directly to a buffer and then back into another object of the same type**. So you can define more efficient copy operations and **serialize the object directly**. This is also useful in an Message Passing context.

A type has a **standard-layout** if it orders and packs its members in a way that is compatible with the C language. A pointer to a standard-layout class may be converted (with `reinterpret_cast`) to a pointer to its first non-static data member and vice versa. **But the most important thing is: you can communicate with C code.**

See the example in [MetaProgramming/Trivial](#).

Some notes

The type traits `is_trivially_copyable` and `has_standard_layout` do what they say but they give sufficient, not necessary conditions. The compiler is not psychic and cannot examine the semantic of your code:

```
struct S{  
    S& operator=(const S& r){i=r.i; a=r.a; return *this;}  
    int i;  
    double a;  
};
```

This class is “morally” trivially copyable, since the copy-assignment I have defined **does exactly what the synthetic (trivial) one would do!**. But the compiler cannot know it, so `is_trivially_copyable <S>::value` is **false**.

Some advice

- ▶ Don't define your own constructor/move-copy assignment operators if the synthetic ones are provided and do what you want (and it happens in most of the cases).
- ▶ Containers and algorithms of modern implementation of the standard library make use of the properties of contained types to optimize copy of objects. Indeed, they may use `memcpy` (or `memmove`), which are more efficient than the usual copy operations. **Another reason to use standard containers and algorithms!**

Defining your type-traits

So far we have seen some used of type-traits provided by the C++ standard library. But maybe you may want to define yours!. We will present here only a few basic examples, also because with the introduction of concepts some things can (sometimes) be made simpler.

First, we introduce `std::integral_constant<T,V>` and its most important specialization, `std::false_type<T>` and `std::true_type<T>`

`integral_constant<T,V>`

This little trait expresses and integral constant expression of value V and type T (clearly a integral type!). It contained value=V and is **implicitly convertible to T** and the conversion provides the value V. You can think it as a map from an (integral) value to a type. An example of **type tagging**

```
enum ORDERING{ ROWMAJOR, COLUMNMAJOR };
template <ORDERING O> // a template alias for simplicity
using OT=std::integral_constant<ORDERING,O>;
template<ORDERING O>
Class Matrix{
public:
...
    auto getIndex(size_t i, size_t j){ return getIndex(i, j, OT<O>{});}
private:
    auto getIndex(size_t i, size_t j, OT<ROWMAJOR>);// version for ROWMAJOR
    auto getIndex(size_t i, size_t j, OT<COLUMNMAJOR>);// version for COLUMNMAJOR
};
```

the correct version is obtained automatically by choosing the correct overload (tag dispatching)

dispatching with `if constexpr`

Clearly you can obtain the same result with an `if constexpr` (more similar to ordinary programming)

```
enum ORDERING{ ROWMAJOR, COLUMNMAJOR };  
template<ORDERING O>  
Class Matrix{  
public:  
...  
    auto getIndex(size_t i, size_t j){  
        if constexpr (O == ROWMAJOR)  
            // row major version  
        else  
            // columnmajor version  
        }  
};
```

Simpler, but I like also the solution with tag dispatch which is also working if you compile with a non c++17 compliant compiler.

true_type and false_type

These specializations of integral_constant. They represent bool and they convert implicitly to the corresponding value

```
bool a= true_type{}; // a is true  
auto c=false_type::value; // c is a bool with value false
```

We will see a possible use in the following handcrafted examples.

A simple example

We want to assess if a type is an `std::complex<T>`

```
#include <type_traits>
#include <complex>

template<typename T>
struct is_complex : std::false_type {};// primary template
// partial specialization for complex
template<typename T>
struct is_complex<std::complex<T> > : std::true_type {};

// Example of usage
template<typename T> auto modulo(T x){
    if constexpr (is_complex<T>{}){
        return std::sqrt(x.real()*x.real())...
    }
    else
        return std::abs(x)
    ...
}
```

This example of usage is useless since `std::abs` has already an overload for complex that computes the modulo.

Testing for the presence of a member function

It would be nice if I could make the compiler check if a type contains the method `clone()`. It is possible, but it is not so easy, there are different ways of doing it. I present one in [Utilities/CloningUtilities.hpp](#). It is rather complex. I leave it to the nerds. In the folder [MetaProgramming/IsClonable](#) you have a simpler version, which however does not check the return type. We omit the detail since now with the introduction of concepts and requires expressions the test can be made much more easily.

Concepts

Concepts have been introduced in C++20 to make template programming safer and simplify the creation of type traits. Concepts, in fact, introduce semantics to types. The main objectives are

- ▶ Have a safer and more understandable template code;
- ▶ Simplify some template metaprogramming constructs;
- ▶ Provide better error messages;

Let start with an example

I want to create a template function that however accepts only integral types. Before C++20 I could do

```
#include <type_traits>
template <typename T>
T absdiff (T const & a, T const & b)
{
    std::static_assert(std::is_integral_v<T>,"Must_be_integral" );
    return a>b? a-b: b-a;
}
```

This is fine and still valid, but it may become tedious and not scalable if the condition on the type is more complex.

The concept solution

Using a concept with a **requires clause**

```
#include <concepts>
template<class T>
requires std::integral<T>
T absdiff (T const & a, T const & b)
{return a>b? a-b: b-a;}
```

or, if you prefer a trailing requires,

```
#include <concepts>
template<class T>
T absdiff (T const & a, T const & b) requires std::integral<T>
{return a>b? a-b: b-a;}
```

or, **even nicer**, just

```
#include <concepts>
template <std::integral T>
T absdiff (T const & a, T const & b)
{return a>b? a-b: b-a;}
```

Simpler and clearer, isn't it! `std::integral` is a **predefined concept** provided by the standard library that represent the semantic of *integral type*, the (almost) full list is **here**.

Other examples

If the concept takes just one template parameter you may use a **constrained template parameter**

```
template<std::floating_point T>  
T foo(T const &);
```

The **requires clause** allows more flexibility

```
template<typename B, typename D>  
requires std::derived_from<B,D> // D must derive from B  
void foo(B const &, D const &);
```

You can combine concepts

```
template<typename T>  
requires std::integral<T> || std::floating_point<T>  
void foo(B const &)
```

Constrained **auto**

In many places where you can use **auto** you can now use constrained **auto**:

```
// Error if the function does not return and integral type  
template<class T>  
std::integral auto foo(T val);  
// This is equivalent to  
// template<std::movable T> void foo(T&& a)  
void foo(std::movable auto&& a); // ok only if arg can be moved  
// this lambda accepts only floating points  
auto f=[](std::floating_point auto x){return 3*x;}
```

Use with class templates

```
template<std::floating_point T, typename U>  
class C{  
public:  
    void push_back(const U & e) requires std::copyable<U>;  
    ...  
}
```

Also for class you have also the requires clause syntax

```
template<typename U>  
requires std::swappable<U, double>  
class C{  
    ...  
}
```

Overloading with concepts

Concepts participate to overloading/specialization mechanism

```
#include <iterator>
template <std::forward_iterator I>
auto foo(I start, I end); // version 1

template <std::random_access_iterator I>
auto foo(I start, I end); // version 2
...
...
std::vector<double> v;
std::set<double> s;
foo(std::begin(v), std::end(v)); // calls version 2
foo(std::begin(s), std::end(s)); // calls version 1
```

Specialization with concepts

Another example with classes

```
#include <concepts>
template <typename T>
class MyClass{....}
```

```
// Specialization for floating points
```

```
template <std::floating_point T>
class MyClass{....}
....
```

```
MyClass<float> a; // uses specialization
MyClass<int> b; // uses primary template
```

Writing your concepts

The general syntax is

```
template<typename T> // you can have more the 1 param.  
concept Name = constraint_expression;
```

A `constraint_expression` can be a **constexpr** boolean expression, logical combination of other concepts or a **requires expression**.

```
// using a user defined trait (defined in is_complex.hpp)  
template <class T>  
concept complex = apsc::TypeTraits::is_complex_v<T>;  
// combining two concepts  
template <class T>  
concept numeric = std::floating_point<T> || complex<T>  
// using requires expression  
template <typename T, typename Q>  
concept multipliable = requires( T v, Q q)  
{ v*q; } // multiplication btw T and Q must be valid
```

Let's look in more details the **requires expression** (not to be confused with the `requires` clause seen in the previous slides).

The `requires` expression

The `requires` expression is a sort of "template function" whose content is however not evaluated but just checked for consistency. If we have a failure the `requires` returns **false**. A simple example

```
template <typename T>
concept addable = requires (T a, T b) {a+b};
```

```
template <typename T>
concept has_value_type = requires {
typename T::value_type; // Type member value_type must exist
};
```

```
...
template <has_value_type T> // fails if T has no value_type
class C{...};
void function(addable auto x){...} // fails if x not addable
```

```
...
std::set<double> s;
function(1); // OK
function(s); // Fails: operator + not defined for sets
C<std::set<double>> c; //ok
C<int> d; //Fails. int does not contain a type called value_type
}
```

Testing the presence of a member function

Inside the requires expression the `{}` block (called compound expression) may be used to create a unevaluated context for an expression. It tests its validity and may also check the expression using a concept. Here, a possible concept for clonable classes :

```
template <typename T>
concept clonable = requires( T v)
{
    {v.clone()} -> std::convertible_to<std::unique_ptr<T>>;
};
```

Here an alternative way of testing if a type is `std::complex` number

```
template<class T>
concept Complex=
requires(T x)
{
    {x.imag()} -> std::same_as<T>;
    {x.real()} -> std::same_as<T>;
};
```

Note the use of the `same_as` concept to test the type.

A note

Note the difference between the type-trait `is_complex` defined at the beginning of this lecture, where the construction relies on class specialization and SFINAE mechanism, and the concept `Complex<T>` in a previous slide.

The main difference is that the type traits `is_complex<T>` checks whether `T` is in fact a `complex<C>`, for any `C`; the concept `Complex<T>` tests whether `T` satisfies a specific semantic: having two function members of a given name and return type.

The nested `requires` expression

Look at this concept that wants to test if a class defines objects of size greater than 16 bytes:

```
#include <concepts>
template <typename T>
concept bigsize = requires (T a) {
    requires sizeof(a) > 16; // size of object must be >16bytes
};
```

If I had written just

```
sizeof(a) > 16;
```

the concept will always return **true** since that expression is **always valid**. I need a nested **requires clause** to test the value of the resulting constant expression.

Overview of the synopsis of requires expressions

```
template<typename T> /*...*/  
requires (T x) // optional set of fictional parameter(s)  
{  
    // simple requirement: expression must be valid  
    x++;    // expression must be valid  
  
    // type requirement: 'typename T', T type must be a valid type  
    typename T::value_type;  
    typename S<T>;  
  
    // compound requirement: {expression}[noexcept][-> Concept];  
    // {expression} -> Concept<A1, A2, ...> is equivalent to  
    // requires Concept<decltype((expression)), A1, A2, ...>  
    {*x}; // dereference must be valid  
    {*x} noexcept; // dereference must be noexcept  
    // dereference must return T::value_type  
    {*x} noexcept -> std::same_as<typename T::value_type>;  
  
    // nested requirement: requires ConceptName<...>;  
    requires integral<T>; // constraint integral<T> must be satisfied  
};
```

Concepts and type-traits

Concepts and type traits are very much linked. We have seen that a type-trait can be easily transformed into concepts. For instance here I exploit the `std::is_integral` trait:

```
template<typename T>  
concept integral = std::is_integral_v<T>;
```

At the same time, concepts can be used to test condition on types, similarly to type traits. This code snippet

```
if constexpr(integral<T>)  
{...
```

is equivalent to

```
if constexpr(is_integral_v<T>)  
{...
```

Beware: Sometimes there are subtle differences between type-traits and related concept. Always consult good references.

Conclusions

What we have seen are the major feature of concepts. On references on the web you find more details, but what we have said so far is quite sufficient to enable you to use constraints in your program if you wish.

Some other examples are in the folder [C++20/Concepts/](#).

To know more

If you want to know more, besides [cppreference.com](#) you have some nice blogs, [here](#), [here](#), [here](#) and [here](#) (in order of complexity), plenty of examples.