



Parallel Patterns in OpenMP and CUDA

Parallel Computing

Serena Curzel



Politecnico di Milano

Dipartimento di Elettronica, Informazione e Bioingegneria

serena.curzel@polimi.it



- ❑ We can use the OpenMP API to implement parallel patterns
- ❑ Some of them come for free!
- ❑ Notebook (includes link to download code examples):

https://colab.research.google.com/drive/18aKcU9i11pqz629vP-kWOC9bhJtvBfc?usp=share_link




- ❑ SIMD, no dependencies between operations
- ❑ Comes for free

`#pragma omp parallel for`





□ Also comes for free

```
#pragma omp parallel for  
reduction(+:var)
```




- ❑ The structure of the work is not regular, it can change at runtime
- ❑ Assumes that all executions are independent
- ❑ Typical application: tree searches

```
#pragma omp task
```




- ❑ Dedicated clause available from OpenMP 5.0
- ❑ Three-phase approach

- ❑ Dedicated clause available from OpenMP 5.0



```
#pragma omp simd reduction(inscan,  
+:scan_a)  
for(int i = 0; i < N; i++){  
    simd_scan[i] = scan_a;  
    #pragma omp scan exclusive(scan_a)  
    scan_a += a[i];  
}
```



Three-phase tiled approach

1 - Build intermediate results in parallel

2 - Sum intermediate results in pairs

3 - Build output results in parallel



- ❑ The histogram is an important and useful computation pattern
- ❑ Allows us to illustrate the concepts of
 - ▶ Partitioning
 - ▶ Atomic operations
 - ▶ Privatization
- ❑ Key issue to solve: **output interference**

- ❑ The histogram is a method for extracting notable features and patterns from large data sets.


Examples:

- ▶ Feature extraction in images
- ▶ Fraud detection in credit card transactions
- ▶ Speech recognition

- ❑ Basic histograms - for each element in the data set, use the value to identify a "bin counter" to increment



- Histogram that counts the occurrence of letters in a string (grouped in 4-letter bins)
 - ▶ For each character in the input string, increment the appropriate bin counter
- Example – histogram for the string “programming massively parallel processors”

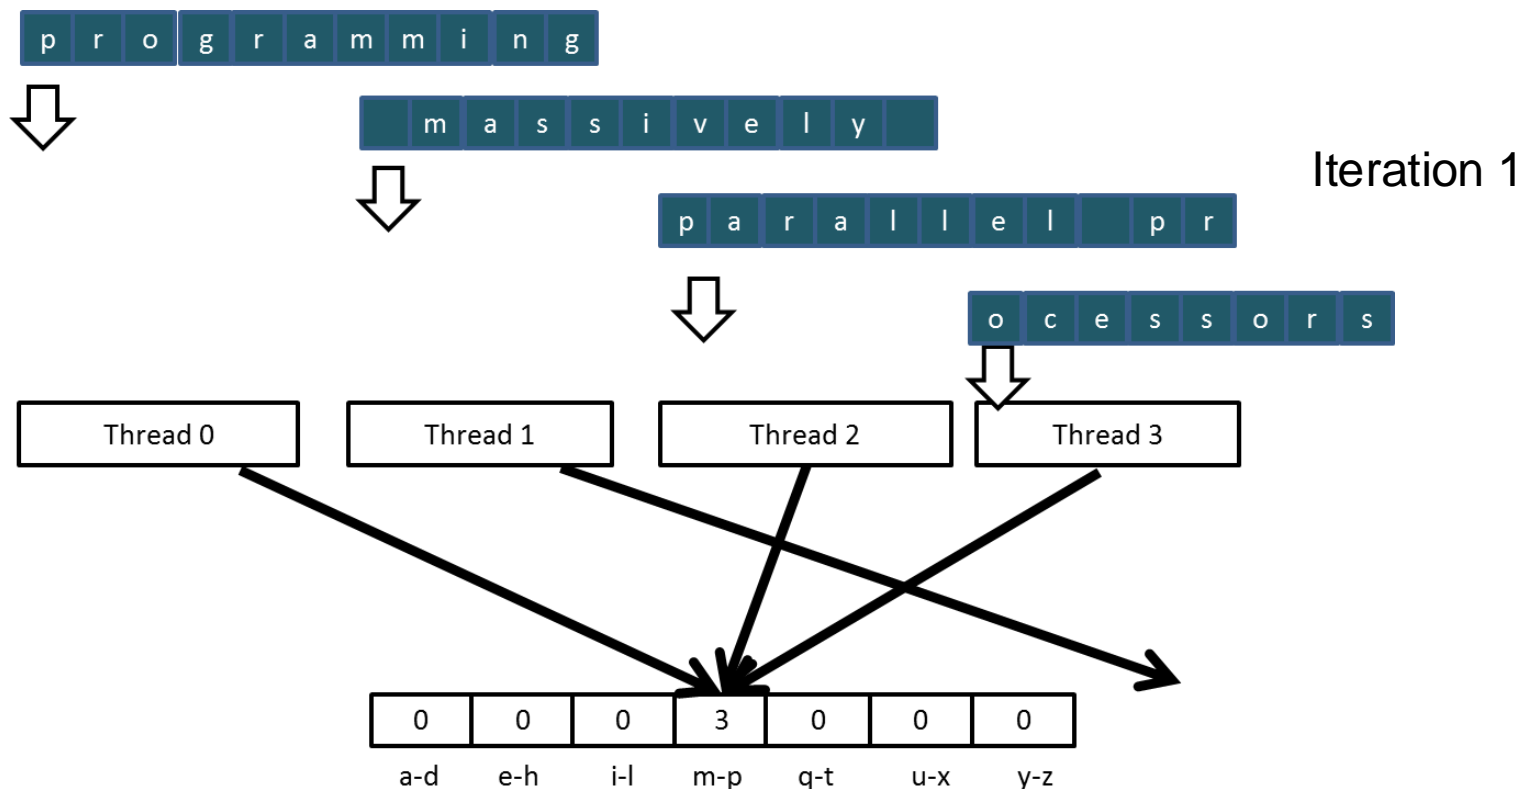


□ Sequential implementation in C

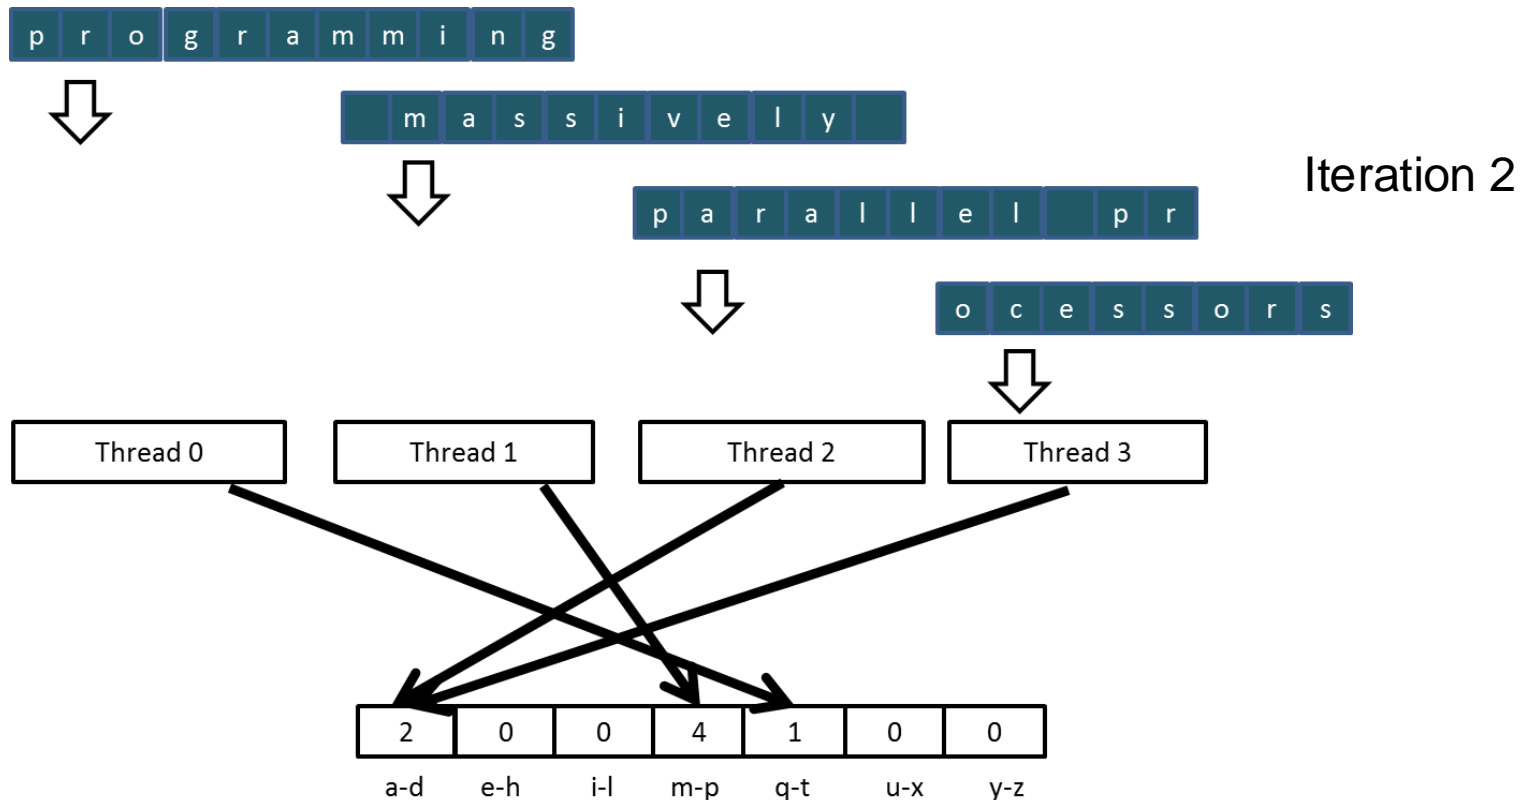
```
sequential_Histogram(char *data, int length, int *histo)
{
    for (int i = 0; i < length; i++) {
        int alphabet_position = data[i] - 'a';
        if (alphabet_position >= 0 && alphabet_position < 26)
        {
            histo[alphabet_position/4]++
        }
    }
}
```

□ Simplest parallel version:


- ▶ Partition the input into sections
- ▶ Each thread iterates through one section




- Simplest parallel version:
 - ▶ Partition the input into sections
 - ▶ Each thread iterates through one section



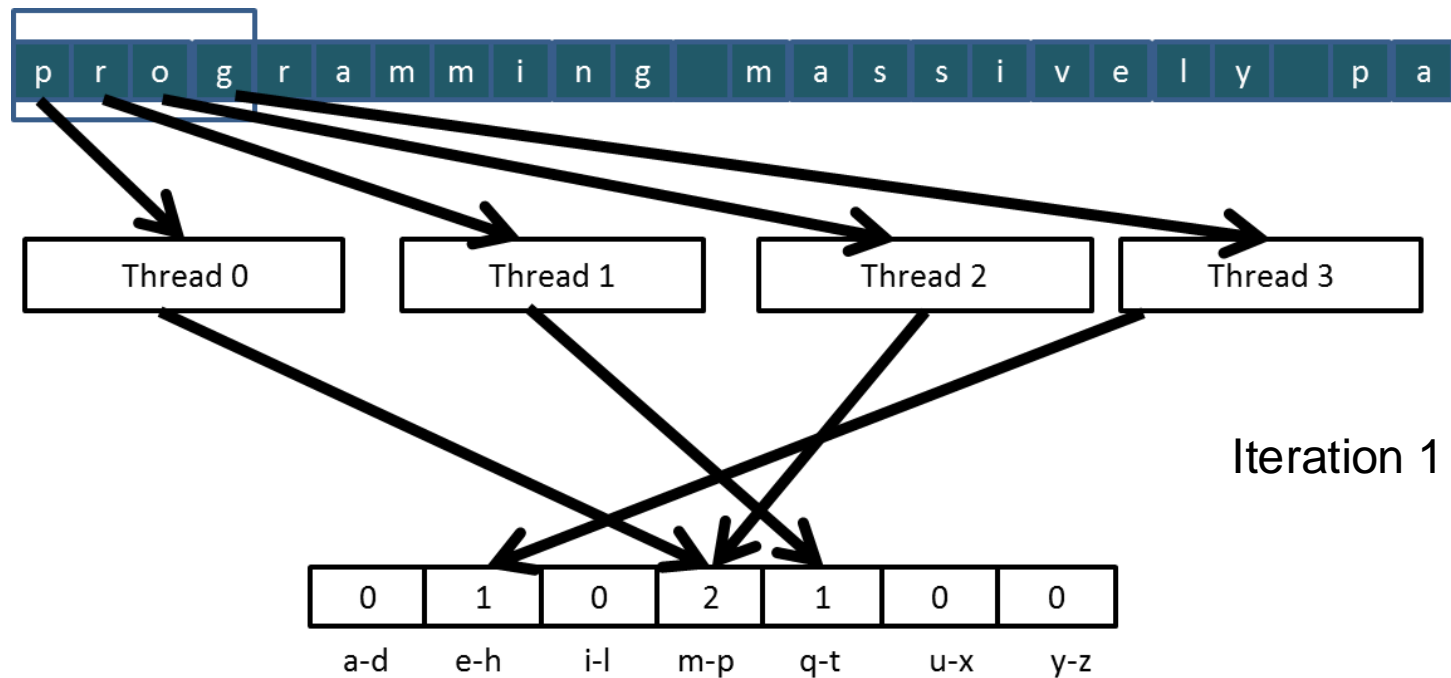
- ❑ Issue: input partitioning affects memory access efficiency
 - ▶ Adjacent threads do not access adjacent memory locations
 - ▶ Accesses are not coalesced
 - ▶ DRAM bandwidth is poorly utilized



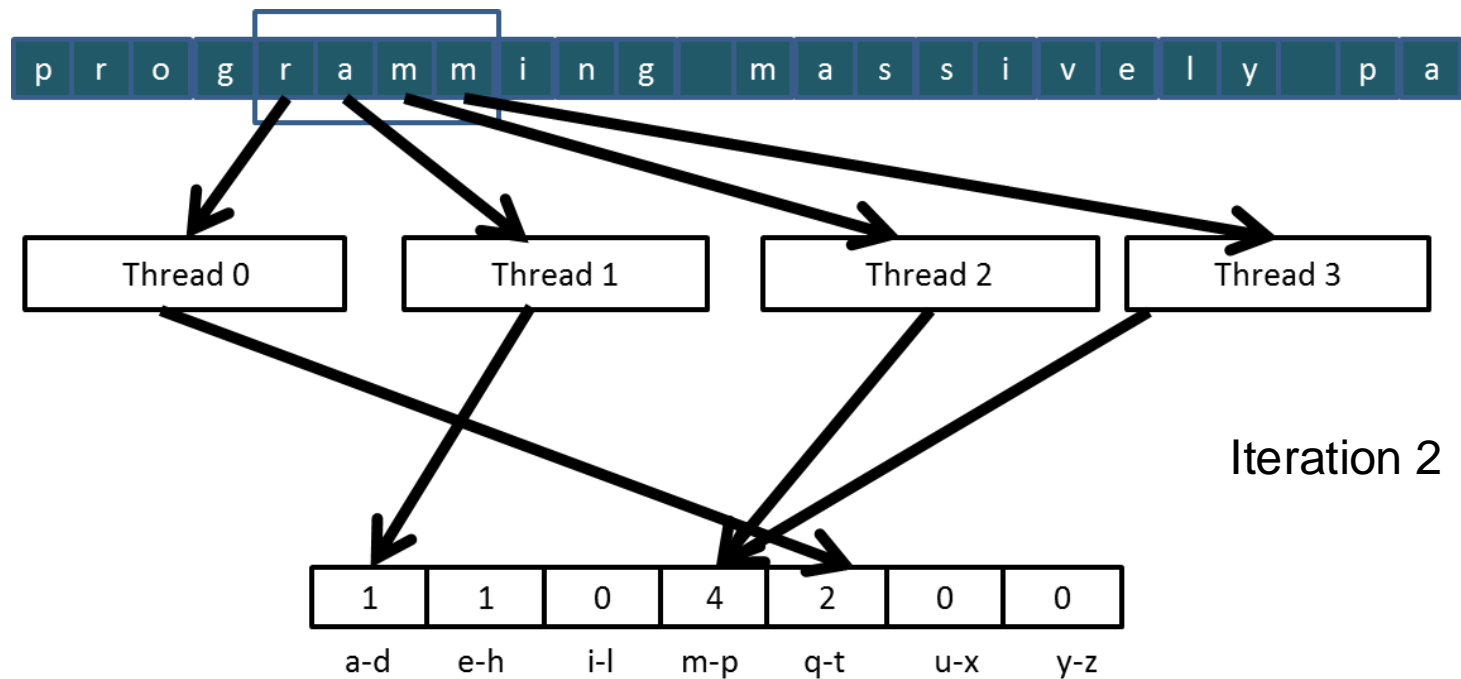
- ❑ Change to interleaved partitioning
 - ▶ All threads process a contiguous section of elements
 - ▶ They all move to the next section and repeat
 - ▶ Memory accesses are coalesced



□ Parallel histogram with interleaved partitioning



Parallel histogram with interleaved partitioning



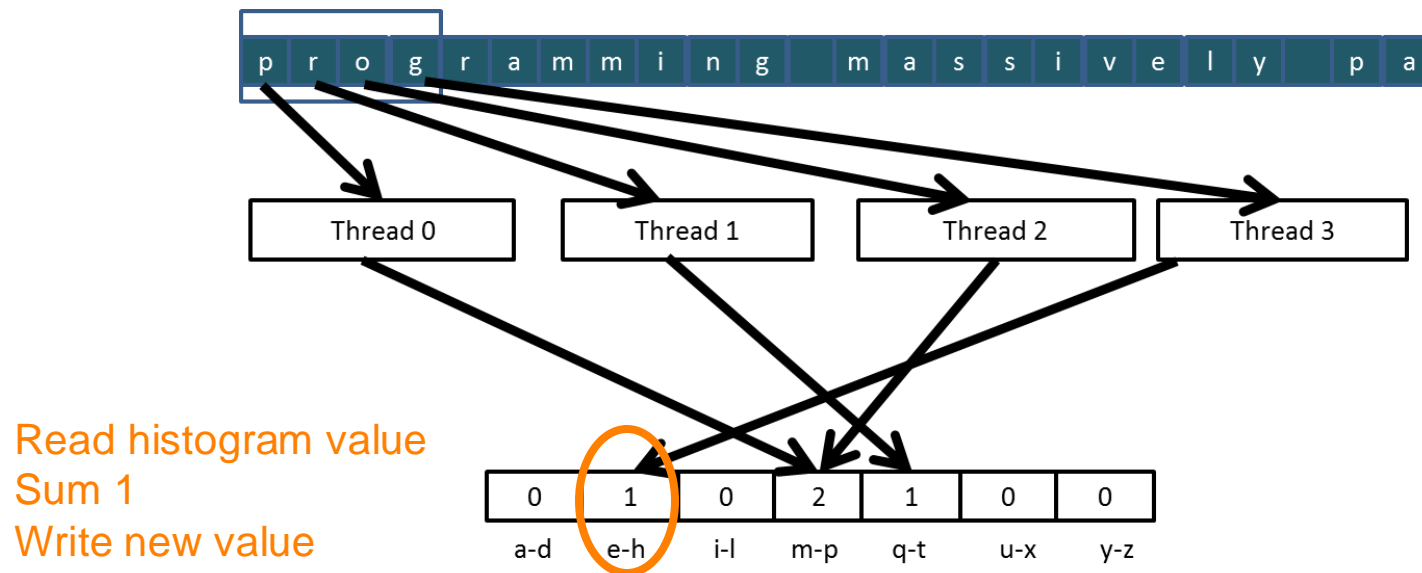
□ Parallel histogram in CUDA with interleaved partitioning

```
__global__ void histo_kernel(unsigned char *buffer, long
size, unsigned int *histo)
{
    int tid = threadIdx.x + blockIdx.x * blockDim.x;
    int stride = blockDim.x*gridDim.x;

    for (unsigned int i = tid; i < size; i += stride ) {
        int alphabet_position = buffer[i] - 'a';
        if (alphabet_position >= 0 && alpha_position < 26)
            histo[alphabet_position/4] += 1;
    }
}
```

❑ Issue: data races

- ▶ Data races occur in read-modify-write operations
- ▶ Data races cause errors that are hard to reproduce
- ▶ Histograms (and other collaboration patterns) include read-modify-write operations



thread1:

Old = Mem[x]
New = Old + 1
Mem[x] = New

Read
Modify
Write

thread2:

Old = Mem[x]
New = Old + 1
Mem[x] = New

□ Data race:

- ▶ If Mem[x] was initially 0, what is its value when threads 1 and 2 have completed?
- ▶ What does each thread get in its Old variable?
- ▶ The answers depend on the relative execution order!

Time	Thread 1	Thread 2
1	(0) Old \leftarrow Mem[x]	
2	(1) New \leftarrow Old + 1	
3		(0) Old \leftarrow Mem[x]
4	(1) Mem[x] \leftarrow New	
5		(1) New \leftarrow Old + 1
6		(1) Mem[x] \leftarrow New

Time	Thread 1	Thread 2
1		(0) Old \leftarrow Mem[x]
2		(1) New \leftarrow Old + 1
3	(0) Old \leftarrow Mem[x]	
4		(1) Mem[x] \leftarrow New
5	(1) New \leftarrow Old + 1	
6	(1) Mem[x] \leftarrow New	

Wrong result

Time	Thread 1	Thread 2
1	(0) Old \leftarrow Mem[x]	
2	(1) New \leftarrow Old + 1	
3	(1) Mem[x] \leftarrow New	
4		(1) Old \leftarrow Mem[x]
5		(2) New \leftarrow Old + 1
6		(2) Mem[x] \leftarrow New

Time	Thread 1	Thread 2
1		(0) Old \leftarrow Mem[x]
2		(1) New \leftarrow Old + 1
3		(1) Mem[x] \leftarrow New
4	(1) Old \leftarrow Mem[x]	
5	(2) New \leftarrow Old + 1	
6	(2) Mem[x] \leftarrow New	

Correct result

- ❑ Atomic operations perform read-modify-write as a single hardware instruction on a memory address
 - ▶ The hardware ensures that no other thread can perform another read-modify-write operation on the same location until the current atomic operation is complete
 - ▶ Any other threads that attempt to perform an atomic operation on the same location will be held in a queue
 - ▶ All threads perform their atomic operations serially on the same location

- ❑ Atomic operations in CUDA are performed by calling functions that are translated into single instructions (*intrinsic functions* or *intrinsics*)
 - ▶ Available intrinsics for atomic add, sub, inc, dec, min, max, exch (exchange), CAS (compare and swap)
- ❑ Example: atomic add
 - ▶ `int atomicAdd(int* address, int val);`
 - ▶ reads a 32-bit word from `address`, computes (old value + `val`), and stores the result back to memory at the same address
 - ▶ the three operations are performed in one atomic transaction

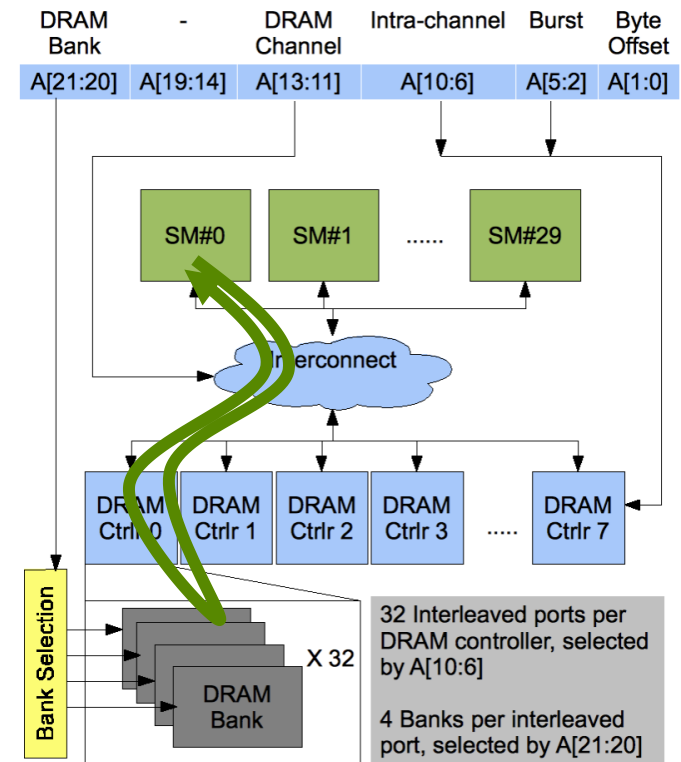
□ Parallel histogram in CUDA with atomic operations

```
__global__ void histo_kernel(unsigned char *buffer,
                             long size, unsigned int *histo)
{
    int tid = threadIdx.x + blockIdx.x * blockDim.x;
    int stride = blockDim.x * gridDim.x;

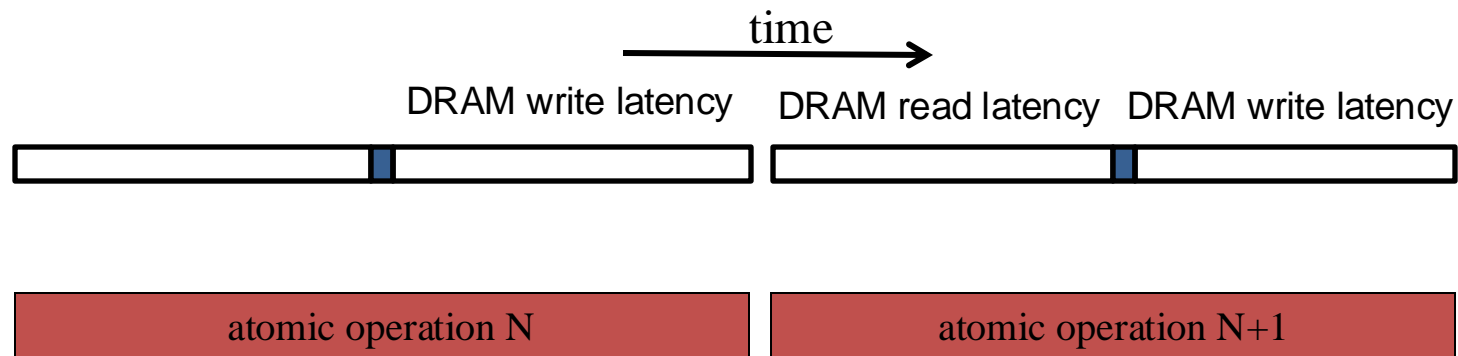
    for (unsigned int i = tid; i < size; i += stride ) {
        int alphabet_position = buffer[i] - "a";
        if (alphabet_position >= 0 && alpha_position < 26)
            atomicAdd(&(histo[alphabet_position/4]), 1);
    }
}
```

□ Performance considerations:



- ▶ Atomic operations can access different types of memory (global, shared, caches)
- ▶ Atomic operations impact latency and throughput
- ▶ For example, a DRAM access has a latency of a few hundred cycles




- ❑ The total latency for an atomic operation on global memory (DRAM) is typically more than 1000 cycles
- ❑ If many threads attempt to do an atomic operation on the same location, the throughput is reduced to 1/1000 of the peak memory throughput
 - ▶ Throughput = rate at which the application can execute an atomic operation



- ❑ Atomic operations on L2 cache
 - ▶ About 1/10 of the DRAM latency
 - ▶ Shared among all blocks
- ❑ Atomic operations on shared memory
 - ▶ Very short latency
 - ▶ Private to each thread block





❑ Cost

- ▶ Overhead for creating and initializing private copies
- ▶ Overhead for accumulating the contents of private copies into the final shared copy

❑ Benefits

- ▶ Much less contention and serialization in accessing both the private copies and the final copy
- ▶ The overall performance can often be improved more than 10x

- ❑ Privatization is a powerful and frequently used technique for parallelizing applications
- ❑ The operation needs to be associative and commutative
 - ▶ Histogram add operation is associative and commutative
- ❑ The private data must fit into shared memory
 - ▶ Works for small histograms
 - ▶ It is possible to partially privatize the histogram and only go to the copy in global memory when needed

- ❑ The histogram pattern is an important use case for privatization
 - ▶ All threads working on an input section are in the same block
 - ▶ Threads in the same block can access shared memory
 - ▶ Throughput for atomics on shared memory is 10-100x better than on DRAM or L2 cache

□ Parallel histogram in CUDA with privatization (part 1)

```
__global__ void histogram_privatized_kernel(unsigned char*
input, unsigned int* bins, unsigned int num_elements,
unsigned int num_bins) {

    unsigned int tid = blockIdx.x*blockDim.x + threadIdx.x;
    extern __shared__ unsigned int histo_s[];

    for(unsigned int binIdx = threadIdx.x; binIdx < num_bins;
binIdx +=blockDim.x) {
        histo_s[binIdx] = 0u;
    }

    __syncthreads();
```

□ Parallel histogram in CUDA with privatization (part 2)


```
for (unsigned int i = tid; i < num_elements; i +=
blockDim.x*gridDim.x) {

    int alphabet_position = buffer[i] - "a";
    if (alphabet_position >= 0 && alpha_position < 26)
        atomicAdd(&(histo_s[alphabet_position/4]), 1);
}
__syncthreads();


for(unsigned int binIdx = threadIdx.x; binIdx <
num_bins; binIdx += blockDim.x) {
    atomicAdd(&(histo[binIdx]), histo_s[binIdx]);
}
}
```

- ❑ The reduction is another typical computation pattern that can be parallelized
- ❑ We can use it to analyze
 - ▶ Work efficiency
 - ▶ Resource efficiency
- ❑ Key issue to solve: **control divergence**


- ❑ Reduction is the final step in “partition and summarize” used to process large data sets
- ❑ Needs an associative (and commutative) combiner function
- ❑ Sequential reduction has $O(N)$ complexity




- ❑ Parallel reduction tree
 - ▶ Assuming there are enough resources, it can perform $N-1$ operations in $\log_2(N)$ steps
 - ▶ Average Parallelism $(N-1)/\log_2(N)$
 - ▶ Work efficient, but not resource efficient!



- ❑ Simple parallel implementation
 - ▶ $N/2$ threads
 - ▶ $\log_2(N)$ steps
 - ▶ At each step each thread operates on 2 elements, then the number of threads is halved
- ❑ In-place reduction in shared memory
 - ▶ Reduced global memory traffic
- ❑ N has to be lower than the maximum number of threads in a block



- ❑ Thread to data mapping:
 - ▶ Each thread writes to an even-index location in the partial sum
 - ▶ One of the inputs comes from the same index
 - ▶ The second input is increasingly far away



□ Baseline CUDA implementation

```
__shared__ float partialSum[SIZE];


partialSum[threadIdx.x] =
X[blockIdx.x*blockDim.x+threadIdx.x];

unsigned int t = threadIdx.x;

for (unsigned int stride = 1; stride < blockDim.x; stride
*= 2)
{
    __syncthreads();
    if (t % (2*stride) == 0)
        partialSum[t] += partialSum[t+stride];
}
```


- ❑ In each iteration, two control flow paths will be sequentially traversed for each warp:
 - ▶ Threads that perform addition and threads that do not
 - ▶ Threads that do not perform addition still consume execution resources
- ❑ Half or fewer of threads will be executing after the first step
 - ▶ After the 5th step, entire warps in each block will fail the if test, poor resource utilization but no divergence

- ❑ Change index usage to reduce divergence
 - ▶ Possible because reduction operators are commutative and associative
- ❑ Partial sums are stored in the front of the array



❑ Improved CUDA implementation

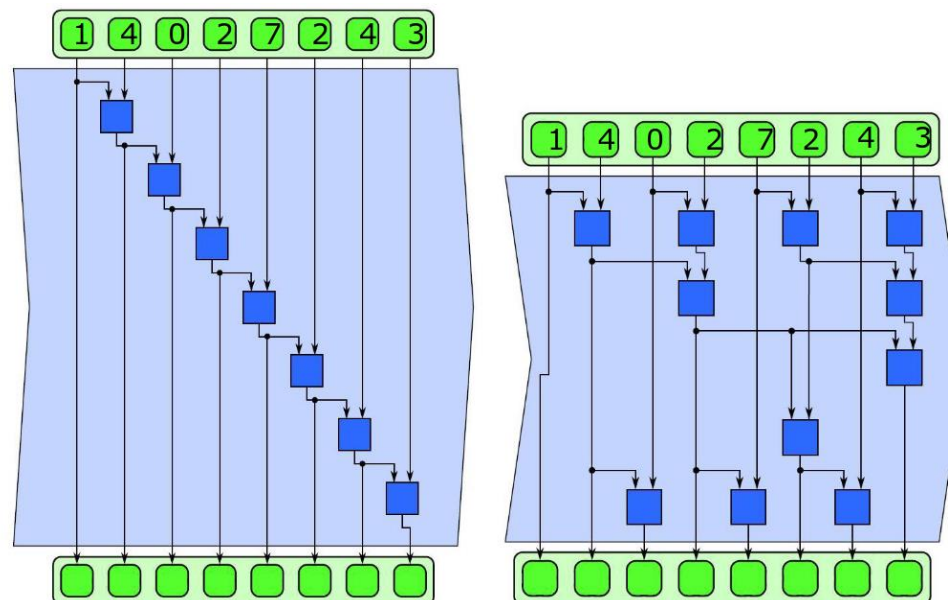
```
__shared__ float partialSum[SIZE];

partialSum[threadIdx.x] =
X[blockIdx.x*blockDim.x+threadIdx.x];

unsigned int t = threadIdx.x;

for (unsigned int stride = blockDim.x/2; stride >= 1;
stride = stride>>1)
{
    __syncthreads();
    if (t < stride)
        partialSum[t] += partialSum[t+stride];
}
```

- ❑ Scan is frequently used to parallelize sequential recursive algorithms
- ❑ Sequential scan is extremely efficient, parallel scan can be slower in some cases
- ❑ Key issue to solve: **work efficiency**



```
y[0] = x[0];  
for (i = 1; i < Max_i; i++)  
    y[i] = y[i-1] + x[i];
```


□ Sequential C implementation of an inclusive addition scan

- ▶ N additions for N elements, $O(N)$ complexity
- ▶ Only 1 operation more than sequential reduction
- ▶ Computationally efficient!

$$\begin{aligned}y_0 &= x_0 \\y_1 &= x_0 + x_1 \\y_2 &= x_0 + x_1 + x_2 \\&\dots\end{aligned}$$

- ❑ Naive parallel implementation of an inclusive addition scan
 - ▶ Threads can work independently on y elements
 - ▶ A lot of redundant computation!
 - ▶ The last thread is performing a sequential reduction, so execution time is the same or worse as the sequential scan

- Parallel implementation with interleaved reduction trees
 - ▶ Read input into shared memory
 - ▶ Thread j adds elements j and $j\text{-stride}$ and writes result into element j
 - ▶ Double stride and repeat



❑ Work-inefficient CUDA implementation


```
__shared__ float XY[SECTION_SIZE];
int i = blockIdx.x * blockDim.x + threadIdx.x;

if (i < InputSize) {
    XY[threadIdx.x] = X[i]; }

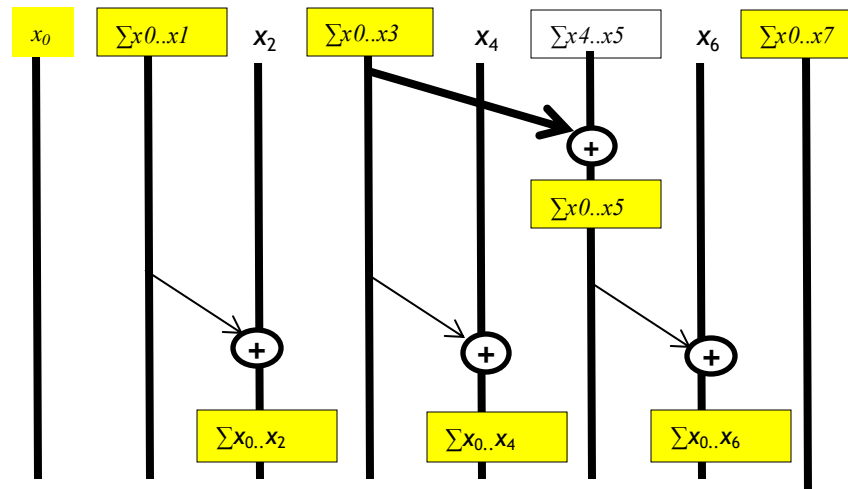
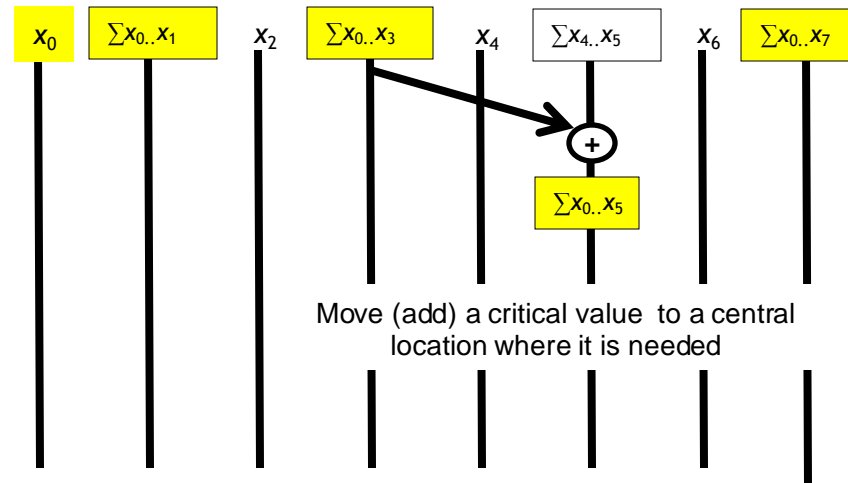
// the code below performs iterative scan on XY
for (unsigned int stride = 1; stride < blockDim.x;
stride *= 2) {
    __syncthreads();
    if (stride <= threadIdx.x) {
        XY[threadIdx.x] += XY[threadIdx.x - stride]; }
}
```


- ❑ The parallel scan executes $\log_2(N)$ iterations
- ❑ The iterations do $(N-1), (N-2), (N-4), \dots (N-N/2)$ additions each
 - ▶ Total number of additions: $O(N \cdot \log_2(N))$
- ❑ Compare to $O(N)$ for the sequential implementation!
 - ▶ 10x overhead with $N=1024$
- ❑ Work inefficient and consumes extra energy due to inactive hardware resources

- ❑ Improving efficiency with:
 - ▶ Two-phased approach
 - ▶ Reuse of intermediate results
 - ▶ More complex thread index-data index mapping
- ❑ The two-phased approach is based on balanced trees
 - ▶ Not an actual data structure, just to describe what each thread does at each step
 - ▶ Create partial sums (leaves to root)
 - ▶ Build output values (root to leaves)



Post-
reduction
reverse
phase



□ Reduction phase

```
// XY[2*BLOCK_SIZE] is in shared memory

for (unsigned int stride = 1; stride <= BLOCK_SIZE;
     stride *= 2) {
    int index = (threadIdx.x+1)*stride*2 - 1;
    if(index < 2*BLOCK_SIZE)
        XY[index] += XY[index-stride];
    __syncthreads();
}
```

□ Distribution phase

```
for (unsigned int stride = BLOCK_SIZE/2; stride > 0;
     stride /= 2) {
    __syncthreads();
    int index = (threadIdx.x+1)*stride*2 - 1;
    if(index+stride < 2*BLOCK_SIZE) {
        XY[index + stride] += XY[index];
    }
}
__syncthreads();
if (i < InputSize) Y[i] = XY[threadIdx.x];
```

- ❑ The work efficient version executes $\log_2(N)$ iterations in the reduction step
 - ▶ Total number of operations: $N-1$
- ❑ In the distribution phase the iterations do $(2-1)$, $(4-1)$, $(N/2-1)$ operations
 - ▶ Total: $(N-2) - (\log_2(N)-1)$
- ❑ Both phases are $O(N)$ and they do not perform more than $2(N-1)$ operations
- ❑ With enough hardware resources, the 2x extra work is compensated by the reduced execution time

- ❑ NVIDIA Deep Learning Institute, "Accelerated Computing" Teaching Kit
- ❑ D.B. Kirk and W.W. Hwu, "Programming Massively Parallel Processors. A hands-on approach", 3rd edition, Morgan Kaufmann 2017
- ❑ Mark Harris, Parallel Prefix Sum with CUDA
https://developer.nvidia.com/gpugems/GPUGems3/gpugems3_ch39.html