



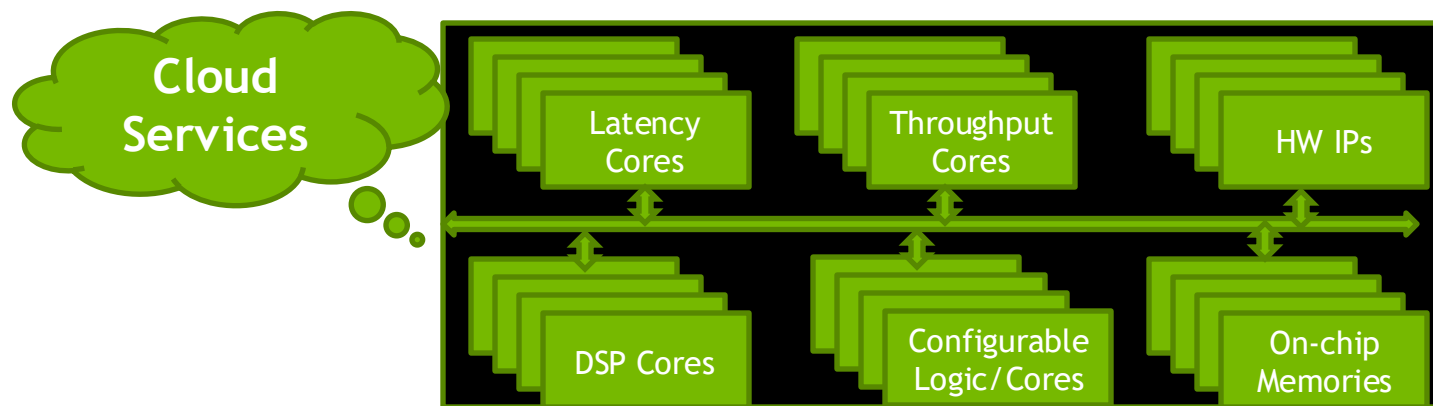
# **Heterogeneous computing: Domain-Specific Languages and High-Level Synthesis**

Parallel Computing

**Serena Curzel**

Politecnico di Milano  
Dipartimento di Elettronica , Informazione e Bioingegneria  
[serena.curzel@polimi.it](mailto:serena.curzel@polimi.it)

- ❑ Observation: most “real world” applications have complex workload characteristics
  - Parallelizable & hard to parallelize
  - SIMD patterns & divergent control flows
  - Predictable & unpredictable data accesses
- ❑ The best system is a heterogeneous one




- ❑ Given a fixed power budget, the goal is not only to increase performance, but also *energy efficiency*
- ❑ Specialization (i.e., fixed function hardware) can provide better energy efficiency

$$\textit{Power} = \frac{\textit{Op}}{\textit{second}} \times \frac{\textit{Joules}}{\textit{Op}}$$

❑ General-purpose processors are not energy-efficient


❑ For each instruction:

- Read instruction (Address translation, communicate with icache, access icache, etc.)
- Decode instruction (Translate op to uops, access uop cache, etc.)
- Check for dependencies/pipeline hazards
- Identify available execution resource
- Use decoded operands to control register file SRAM (retrieve data)
- Move data from register file to selected execution resource
- Perform arithmetic operation
- Move data from execution resource to register file
- Use decoded operands to control write to register file SRAM

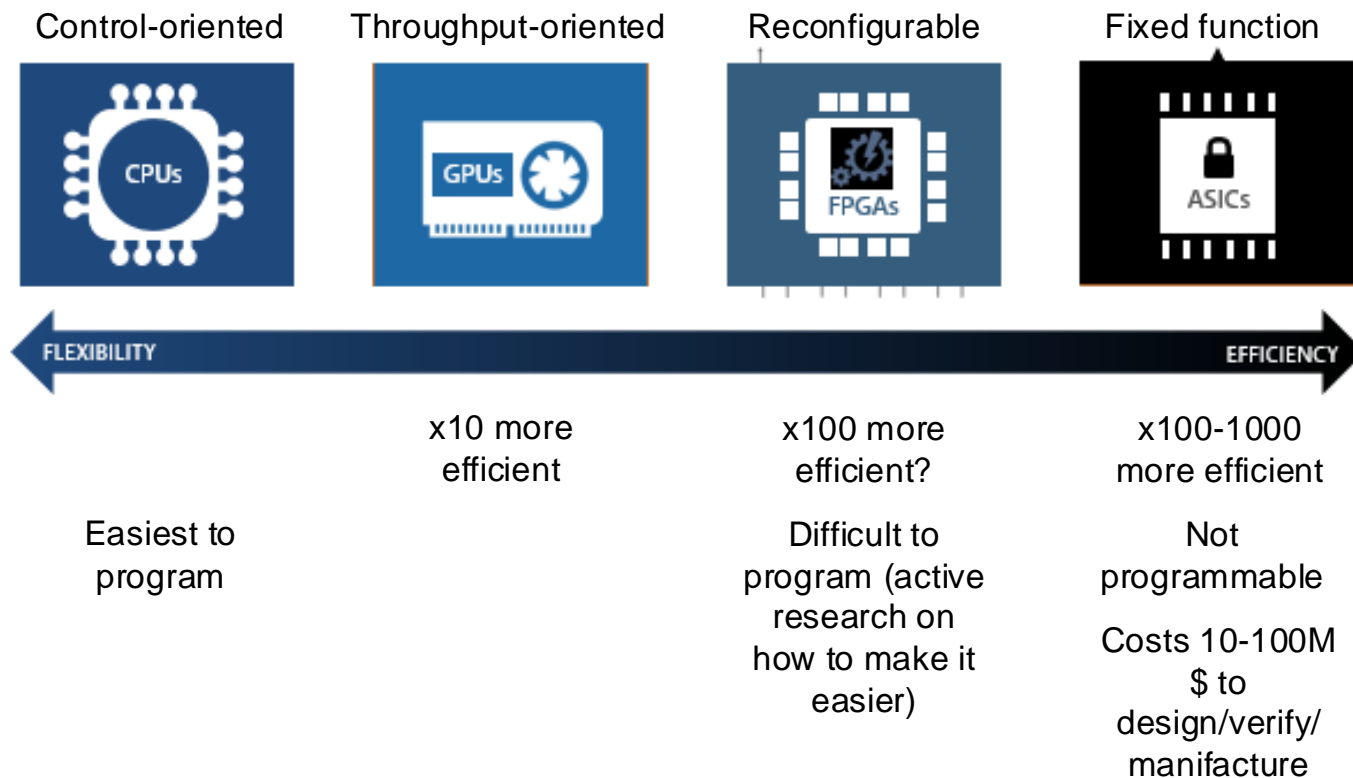


*Efficient Embedded Computing [Dally et al. 08]*

- ❑ On the other hand, the circuit that would be required to actually perform the computation can be very simple




## ❑ Choosing the right tool for the job:




- ❑ Two ideas to reduce energy consumption:
  - Use specialized processors (more operations per joule, higher energy efficiency)
  - Move less data!
- ❑ Minimizing communication overhead increases performance, and it also reduces energy consumption

Operation (32-bit operands)	Energy/Op (28 nm)	Cost (vs. ALU)
ALU op	1 pJ	-
Load from SRAM	1-5 pJ	5x
Move 10mm on-chip	32 pJ	32x
Send off-chip	500 pJ	500x
Send to DRAM	1 nJ	1,000x
Send over LTE	>10 $\mu$ J	10,000,000x

*data from John Brunhaver, Bill Dally, Mark Horowitz*




- How can we write efficient, portable parallel programs for emerging heterogeneous architectures?



Pthreads, OpenMP,  
Vectorized  
instructions, ...

CUDA,  
OpenCL,  
...



MPI, Spark, ...

Verilog, VHDL, ...






- ❑ The ideal parallel programming language provides
  - Performance
  - Productivity
  - Generality

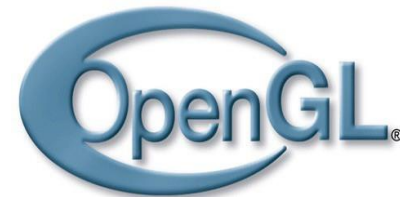


## ❑ Existing languages:


Domain-Specific Languages



- ❑ Programming languages with restricted expressiveness for a particular domain
- ❑ High-level, usually declarative, and deterministic
- ❑ External or embedded within another language

The logo for SQL, featuring the letters "SQL" in a bold, orange, sans-serif font, with a small blue icon of a cursor or arrow pointing upwards and to the right.The logo for PyTorch, featuring the word "PYTORCH" in a bold, black, sans-serif font, with a small orange flame icon above the letter "O".The logo for OpenGL, featuring the word "OpenGL" in a blue, sans-serif font, with a blue circular icon to the left of the text.The logo for MATLAB, featuring the word "MATLAB" in a black, sans-serif font, with a colorful 3D surface plot icon below the text.The logo for TeX, featuring the word "TEX" in a black, serif font.The logo for TVM, featuring a blue icon of a square with a smaller square inside, followed by the letters "tvm" in a black, sans-serif font.The logo for Halide, featuring the word "halide" in a black, sans-serif font.The logo for Rails, featuring a red square icon with a white circular arrow and the word "RAILS" in white, bold, sans-serif font.

- ❑ We are surrounded by data-intensive imaging applications
- ❑ Almost every imaging sensor is accessed by a powerful compute unit



Enormous opportunity, demands for  
**extreme optimization**  
(performance + **power consumption**)

Keywords: **parallelism, locality**

## ❑ Faster algorithms


- Advanced methods, sometimes approximate computing

## ❑ Faster hardware

- Faster CPUs (higher frequency, higher memory bandwidth, SIMD vector operations)
- GPUs

## ❑ More efficient use of hardware

- Parallelism
- Better cache management
- Memory usage



Require to *reorganize computation!*

- ❑ C++ with multithreading and SIMD operations
- ❑ CUDA/OpenCL for GPUs
- ❑ Optimized libraries (BLAS, IPP, MKL, OpenCV)
  
- ❑ Writing efficient image processing pipelines is *hard!*
  - Optimization requires **manually** transforming program and data structure to exploit parallelism and locality
  - Exploring of different options becomes difficult

- ❑ Halide proposes a radically different approach:

decouple the **algorithm** from the **schedule**

*what* is computed

fixed once and never  
modified by scheduling  
decisions

*where* and *when* it is computed

easily changed to  
explore different  
possibilities

- ❑ C/C++ implementation of a *box blur* algorithm with 3x3 window:

```
void box_filter_3x3(const Image &in, Image &blury){
    Image blurx(in.width(), in.height()); // allocate blurx array

    for (int y = 0; y < in.height(); y++)
        for (int x = 0; x < in.width(); x++)
            blurx(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y))/3;


    for (int y = 0; y < in.height(); y++)
        for (int x = 0; x < in.width(); x++)
            blury(x, y) = (blurx(x, y-1) + blurx(x, y) + blurx(x, y+1))/3;
}
```



```
void box_filter_3x3(const Image &in, Image &blurry){
    Image blurx(in.width(), in.height()); // allocate blurx array

    for (int y = 0; y < in.height(); y++)
        for (int x = 0; x < in.width(); x++)
            blurx(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y))/3;

    for (int y = 0; y < in.height(); y++)
        for (int x = 0; x < in.width(); x++)
            blurry(x, y) = (blurx(x, y-1) + blurx(x, y) + blurx(x, y+1))/3;
}
```



- ❑ C/C++ implementation of a *box blur* algorithm with 3x3 window:

```
void box_filter_3x3(const Image &in, Image &blury){
    Image blurx(in.width(), in.height()); // allocate blurx array

    for (int y = 0; y < in.height(); y++)
        for (int x = 0; x < in.width(); x++)
            blurx(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y))/3;

    for (int y = 0; y < in.height(); y++)
        for (int x = 0; x < in.width(); x++)
            blury(x, y) = (blurx(x, y-1) + blurx(x, y) + blurx(x, y+1))/3;
}
```

9.96 ms/megapixel on a quad core x86

## ❑ Optimized C/C++ implementation:

```
void box_filter_3x3(const Image &in, Image &blury) {  
    __m128i one_third = _mm_set1_epi16(21846);  
    #pragma omp parallel for  
    for (int yTile = 0; yTile < in.height(); yTile += 32) {  
        __m128i a, b, c, sum, avg;  
        __m128i blurx[(256/8)*(32+2)]; // allocate tile blurx array  
        for (int xTile = 0; xTile < in.width(); xTile += 256) {  
            __m128i *blurxPtr = blurx;  
            for (int y = -1; y < 32+1; y++) {  
                const uint16_t *inPtr = &(in[yTile+y][xTile]);  
                for (int x = 0; x < 256; x += 8) {  
                    a = _mm_loadu_si128((__m128i*)(inPtr-1));  
                    b = _mm_loadu_si128((__m128i*)(inPtr+1));  
                    c = _mm_load_si128((__m128i*)(inPtr));  
                    sum = _mm_add_epi16(_mm_add_epi16(a, b), c);  
                    avg = _mm_mulhi_epi16(sum, one_third);  
                    _mm_store_si128(blurxPtr++, avg);  
                    inPtr += 8;  
                }  
            }  
            blurxPtr = blurx;  
            for (int y = 0; y < 32; y++) {  
                __m128i *outPtr = (__m128i*)&(blury[yTile+y][xTile]);  
                for (int x = 0; x < 256; x += 8) {  
                    a = _mm_load_si128(blurxPtr+(2*256)/8);  
                    b = _mm_load_si128(blurxPtr+256/8);  
                    c = _mm_load_si128(blurxPtr++);  
                    sum = _mm_add_epi16(_mm_add_epi16(a, b), c);  
                    avg = _mm_mulhi_epi16(sum, one_third);  
                    _mm_store_si128(outPtr++, avg);  
                }  
            }  
        }  
    }  
}
```

**0.9 ms/megapixel on a quad core x86**  
**11x faster**

**parallelism**  
(work distributed among threads,  
SIMD parallel vectors)

## ❑ Optimized C/C++ implementation:

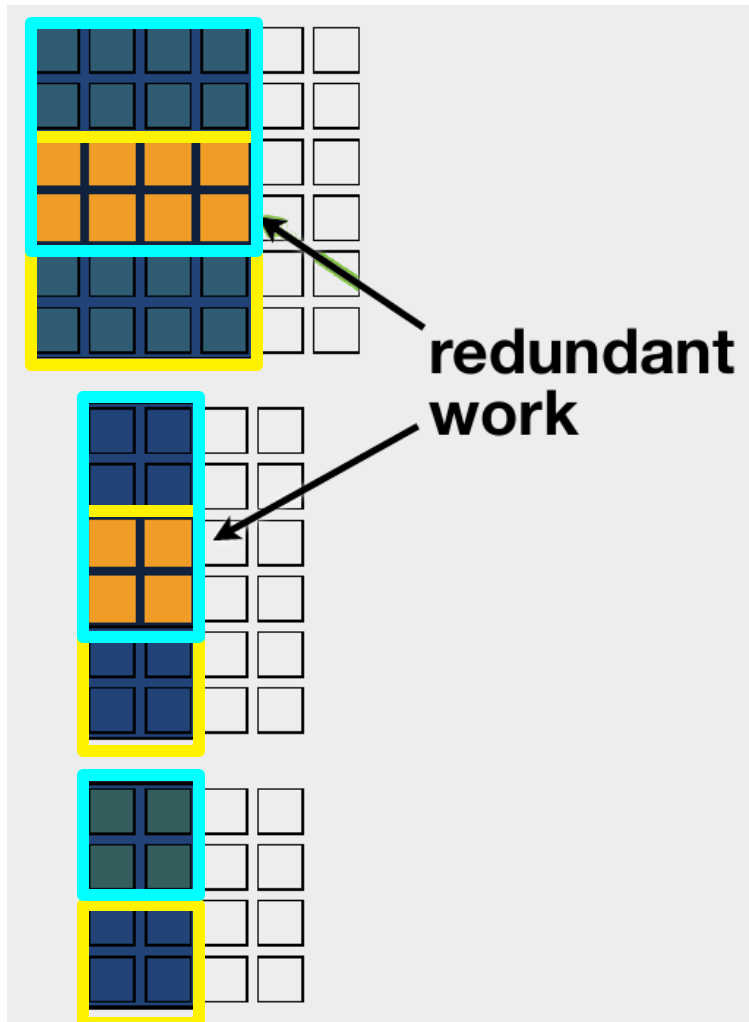
```
void box_filter_3x3(const Image &in, Image &blur) {  
    __m128i one_third = _mm_set1_epi16(21846);  
    #pragma omp parallel for  
    for (int yTile = 0; yTile < in.height(); yTile += 32) {  
        __m128i a, b, c, sum, avg;  
        __m128i blurx[(256/8)*(32+2)]; // allocate tile blurx array  
        for (int xTile = 0; xTile < in.width(); xTile += 256) {  
            __m128i *blurxPtr = blurx;  
            for (int y = -1; y < 32+1; y++) {  
                const uint16_t *inPtr = &(in[yTile+y][xTile]);  
                for (int x = 0; x < 256; x += 8) {  
                    a = _mm_loadu_si128((__m128i*)(inPtr-1));  
                    b = _mm_loadu_si128((__m128i*)(inPtr+1));  
                    c = _mm_load_si128((__m128i*)(inPtr));  
                    sum = _mm_add_epi16(_mm_add_epi16(a, b), c);  
                    avg = _mm_mulhi_epi16(sum, one_third);  
                    _mm_store_si128(blurxPtr++, avg);  
                    inPtr += 8;  
                }  
            }  
            blurxPtr = blurx;  
            for (int y = 0; y < 32; y++) {  
                __m128i *outPtr = (__m128i*)&(blur[yTile+y][xTile]);  
                for (int x = 0; x < 256; x += 8) {  
                    a = _mm_load_si128(blurxPtr+(2*256)/8);  
                    b = _mm_load_si128(blurxPtr+256/8);  
                    c = _mm_load_si128(blurxPtr++);  
                    sum = _mm_add_epi16(_mm_add_epi16(a, b), c);  
                    avg = _mm_mulhi_epi16(sum, one_third);  
                    _mm_store_si128(outPtr++, avg);  
                }  
            }  
        }  
    }  
}
```

**0.9 ms/megapixel on a quad core x86**  
**11x faster**

### locality


(compute in tiles,  
interleave tiles of blurx, blur,  
store blurx in local cache)

- ❑ So the **unoptimized** C/C++ version implements the pipeline like this:
  - Compute `blurx` on the whole input image
  - Store the result to memory
  - Load it back and compute `blury` on it
- ❑ The **optimized** version interleaves the execution of the two stages:
  - Compute `blurx` on a chunk of the input
  - Immediately compute `blury` on the intermediate result
  - Throw away the intermediate data, store the result of `blury` and move to the next chunk of the input



- ❑ Since we immediately throw away the intermediate results, there will be regions of `blurx` that are computed *several times*, once for each neighboring chunk of `blury`
- ❑ The same happens between `blurx` and the input image

- ❑ We were able to move from 9.9 to 0.9ms/megapixel
- ❑ The process required knowledge of parallelism and memory hierarchy
- ❑ **Expensive** (time consuming) because it had to be done manually
- ❑ Difficult to predict the effect of optimization choices
- ❑ **Not easily portable** to different hardware architectures



Now let's see how Halide **solves these problems** and at the same time **maintains (or enhances!) the benefits** in terms of performance

## ❑ Halide algorithm:

```
Var x, y; Func blurx, blurry;  
blurx(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y))/3;  
blurry(x, y) = (blurx(x, y-1) + blurx(x, y) + blurx(x, y+1))/3;
```

## ❑ Halide schedule:

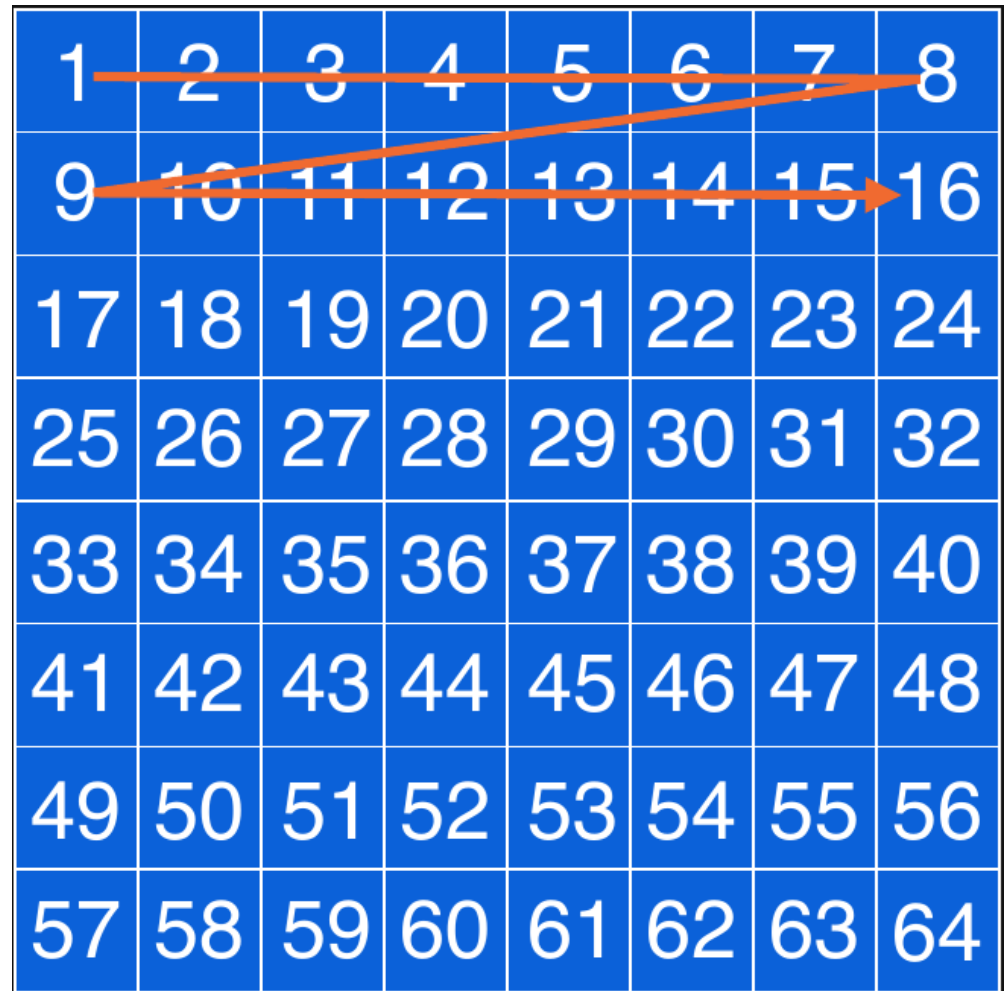
```
blurry.tile(x, y, xi, yi, 256, 32)  
    .vectorize(xi, 8).parallel(y);  
blurx.compute_at(blurry, x).store_at(blurry, x).vectorize(x.8);
```



- ❑ There is a choice space for optimizations in an image processing pipeline
  - In what order should *each stage* compute its values?
  - When should each stage compute its *inputs*?

- ❑ In what order should I compute my values?

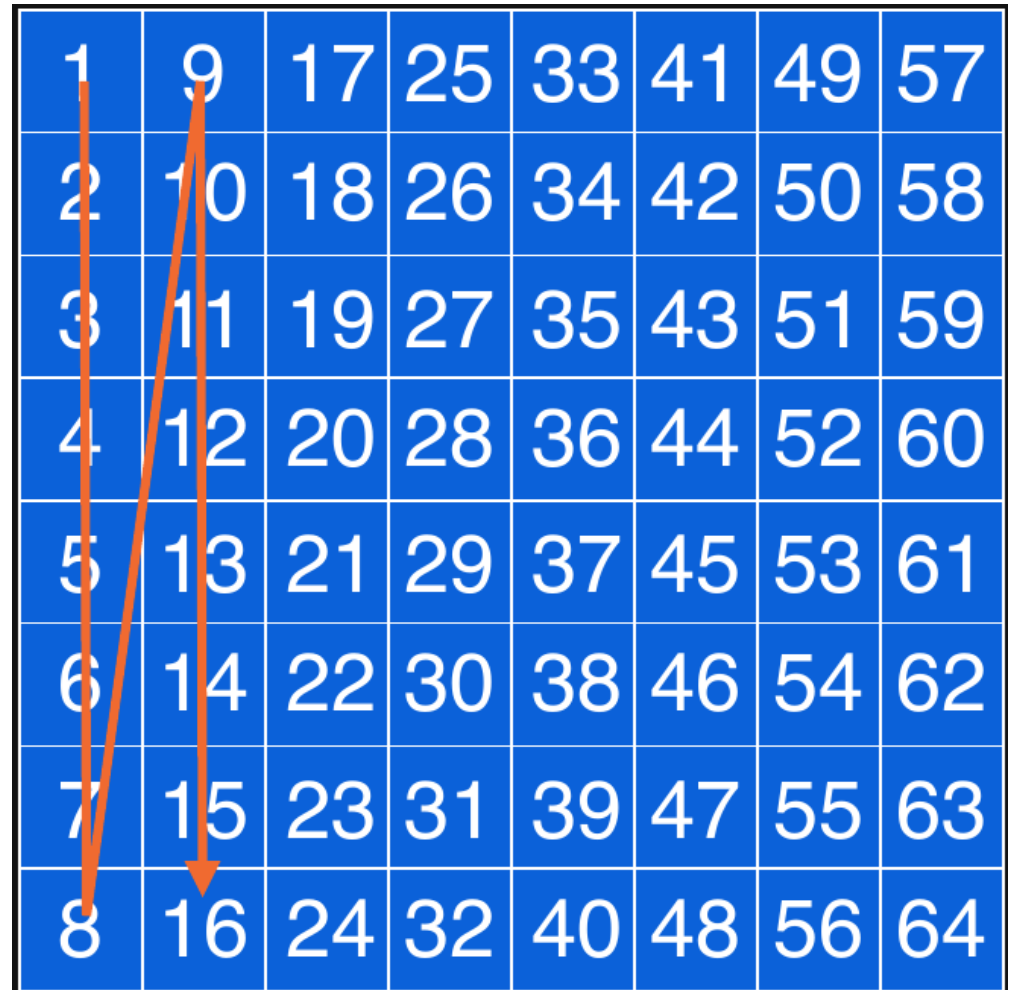
Serial y,  
serial x



1	2	3	4	5	6	7	8
9	10	11	12	13	14	15	16
17	18	19	20	21	22	23	24
25	26	27	28	29	30	31	32
33	34	35	36	37	38	39	40
41	42	43	44	45	46	47	48
49	50	51	52	53	54	55	56
57	58	59	60	61	62	63	64

- In what order should I compute my values?

Serial x,  
serial y



1	9	17	25	33	41	49	57
2	10	18	26	34	42	50	58
3	11	19	27	35	43	51	59
4	12	20	28	36	44	52	60
5	13	21	29	37	45	53	61
6	14	22	30	38	46	54	62
7	15	23	31	39	47	55	63
8	16	24	32	40	48	56	64

- ❑ In what order should I compute my values?

Serial y,  
Vectorize x by 4

	1			2	
	3			4	
	5			6	
	7			8	
	9			10	
	11			12	
	13			14	
	15			16	

- ❑ In what order should I compute my values?

Parallel y,  
Vectorize x by 4

	1			2
	1			2
	1			2
	1			2
	1			2
	1			2
	1			2
	1			2

❑ In what order should I compute my values?


Split x by 2,  
Split y by 2,  
Serial y<sub>outer</sub>,  
Serial x<sub>outer</sub>,  
Serial y<sub>inner</sub>,  
Serial x<sub>inner</sub>

1	2	5	6	9	10	13	14
3	4	7	8	11	12	15	16
17	18	21	22	25	26	29	30
19	20	23	24	27	28	31	32
33	34	37	38	41	42	45	46
35	36	39	40	43	44	47	48
49	50	53	54	57	58	61	62
51	52	55	56	59	60	63	64

## ❑ When should I compute my inputs?

```
compute blury:  
  for ...:  
    blury(...) = ...
```

```
compute blurx:  
  for ...:  
    blurx(...) = ...
```





Compute and store producer, then compute consumer

- + no recomputation
- poor locality

## ❑ When should I compute my inputs?

```
compute blurx:  
  for c:  
    for y:  
      for x:  
        blurx(...) =
```




Compute producer values when needed by consumer,  
then throw them away

- + locality
- redundant work




- When should I compute my inputs?



Compute producer values when needed by consumer,  
then keep the old values for reuse?

- + locality
- + no redundant work
- poor parallelism

- ❑ When should I compute my inputs?




Compute producer values when needed by consumer,  
then throw them away

+ locality

+ no redundant work (property of the consumer function)

## ❑ When should I compute my inputs?

```
compute blurx:  
  for c:  
    for y:  
      compute blury:  
        for x:  
          blury(...) = ...  
      for x:  
        blurx(...) = ...
```



Compute a subset of producer values all at once, then pass them to the consumer

+/- depend on chunk size, kernel size, parallelism...

Animated visualization of these and other scheduling options:

[https://www.youtube.com/watch?time\\_continue=14&v=3uiEyEKji0M&feature=emb\\_logo](https://www.youtube.com/watch?time_continue=14&v=3uiEyEKji0M&feature=emb_logo)

Halide Talk

blur\_x.compute\_root();      blur\_x.compute\_at(blur\_y, x);      blur\_x.store\_root().compute\_at(blur\_y, x);

blur\_x.compute\_at(blur\_y, x)  
  .vectorize(x, 4);  
blur\_y.tile(x, y, xi, yi, 8, 8)  
  .parallel(y)  
  .vectorize(xi, 4);


blur\_x.store\_root()  
  .compute\_at(blur\_y, y)  
  .split(x, x, xi, 8)  
  .vectorize(xi, 4).parallel(x);  
blur\_y.split(x, x, xi, 8)  
  .vectorize(xi, 4).parallel(x);

blur\_x.store\_at(blur\_y, y)  
  .compute\_at(blur\_y, yi)  
  .vectorize(x, 4);  
blur\_y.split(y, y, yi, 8)  
  .vectorize(x, 4)  
  .parallel(y);

36


▶ ⏮ 🔊 16:12 / 25:51      Scorri per i dettagli

- ❑ Performance requires complex trade-offs and frequent reorganization of computation




- ❑ Halide makes this iterative process easier

- How can we write efficient, portable parallel programs for emerging heterogeneous architectures?



Pthreads, OpenMP,  
Vectorized  
instructions, ...

CUDA,  
OpenCL,  
...

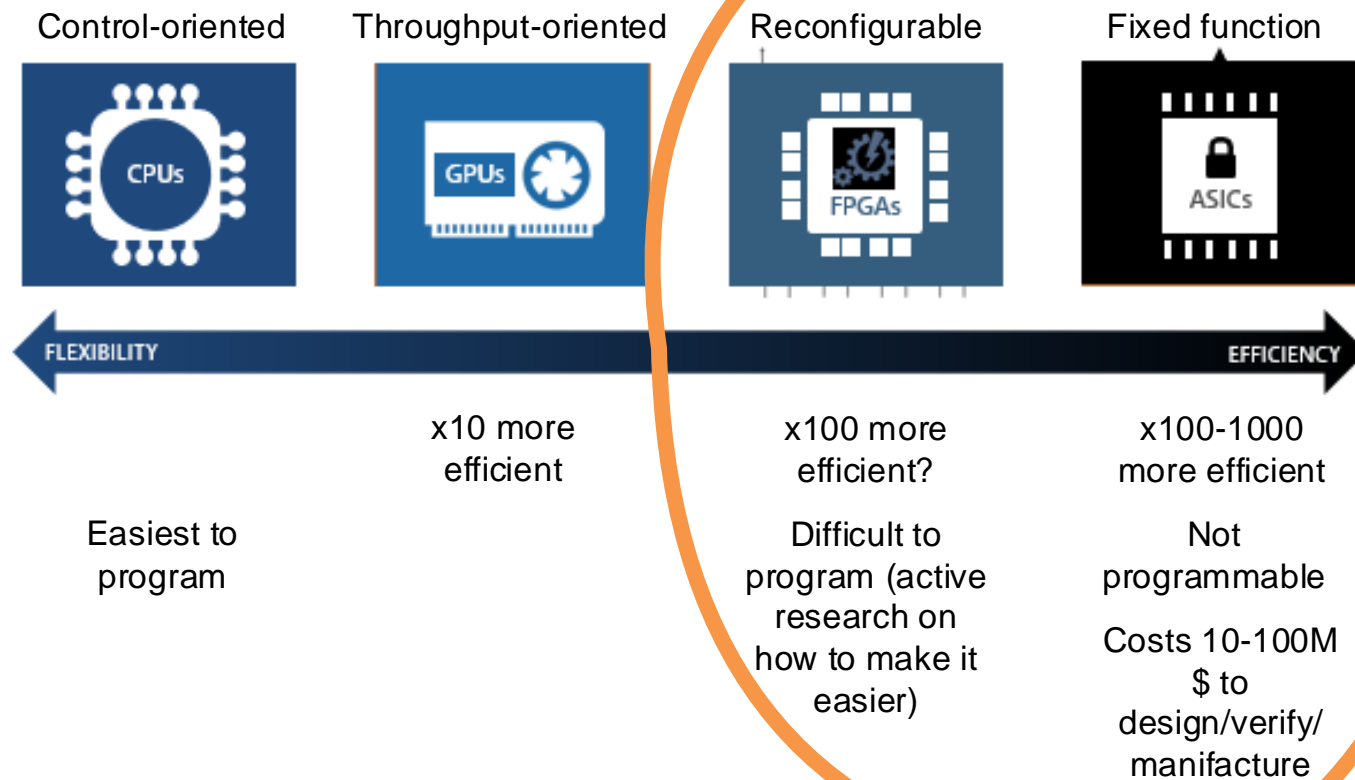


MPI, Spark, ...

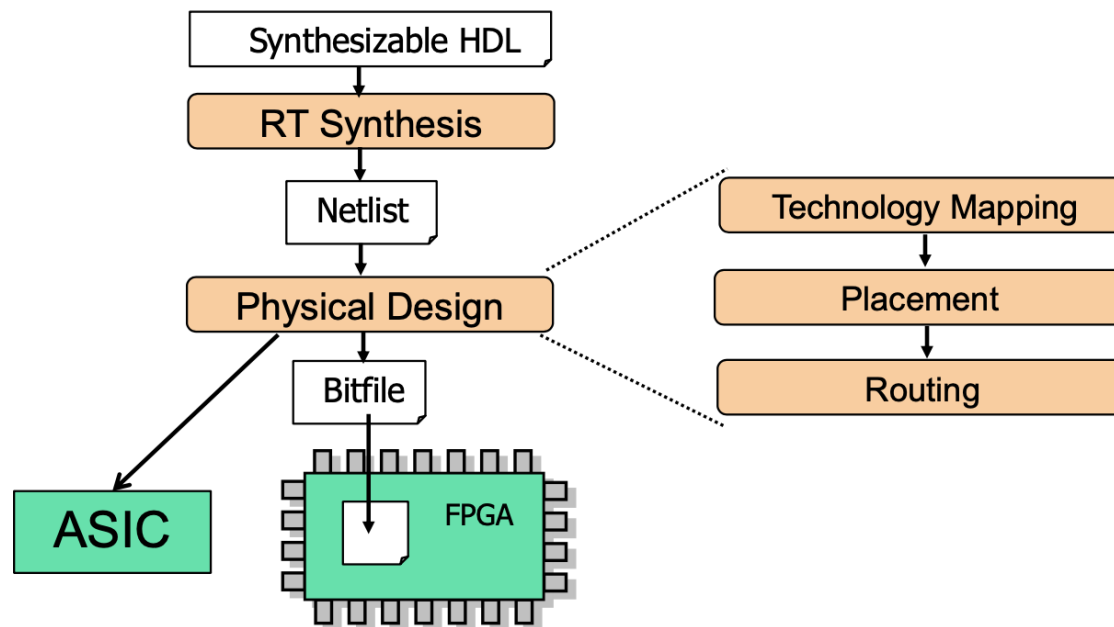
Verilog, VHDL, ...



## ❑ Choosing the right tool for the job:

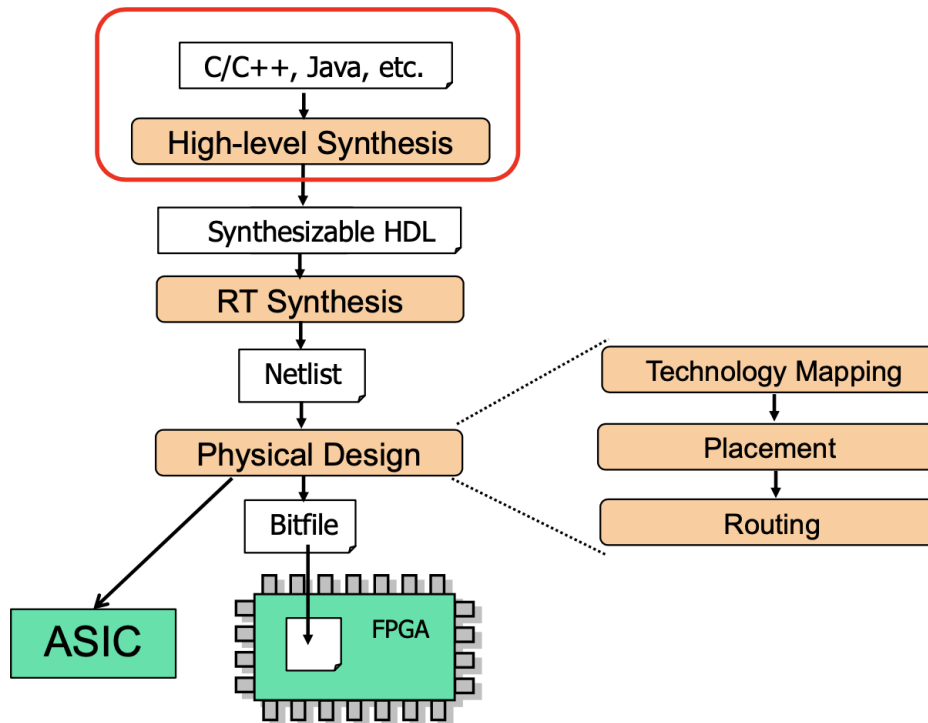


- ❑ Register-Transfer Level (RTL) design
  - Specify behavior and structure (muxes, registers...) in Verilog/VHDL
  - Precise, fine-grained control
  - Time-consuming and error-prone!





- High-level behavioral design through sw programming languages
  - Automated translation to RTL
  - Like compilers, but for hardware



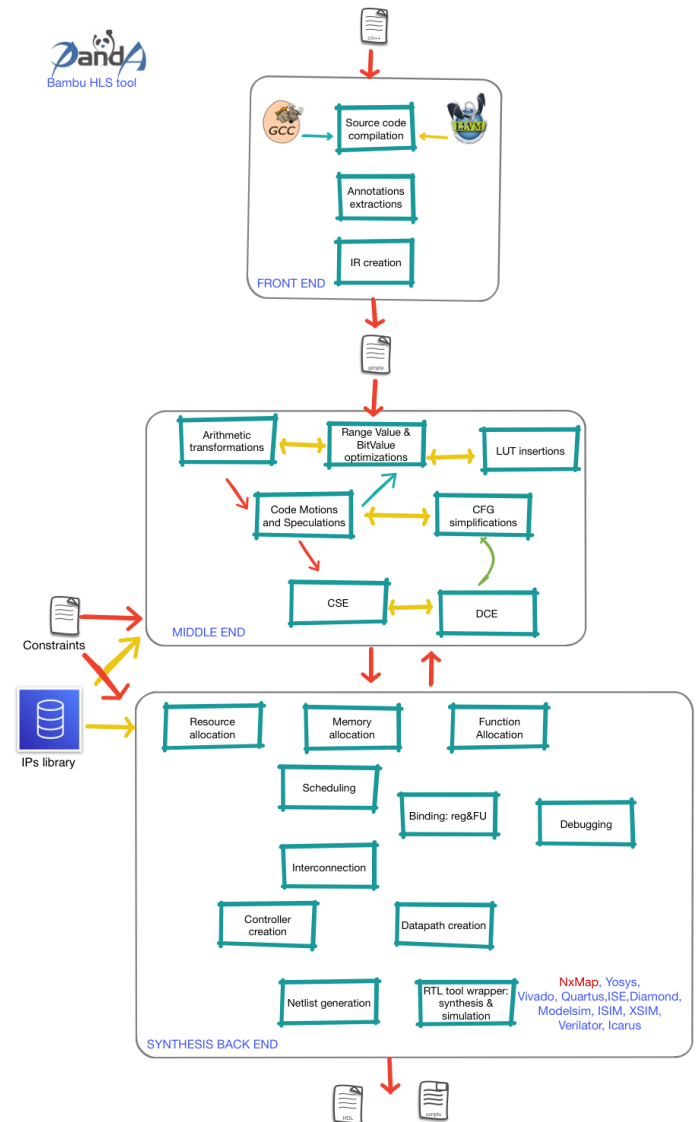
## Benefits

- ❑ Number of expert SW developers vs HW designers
- ❑ Easier to describe complex functionalities
- ❑ Technology-independent design
- ❑ Easier HW/SW partitioning
- ❑ Faster design cycle (including verification)

## Challenges

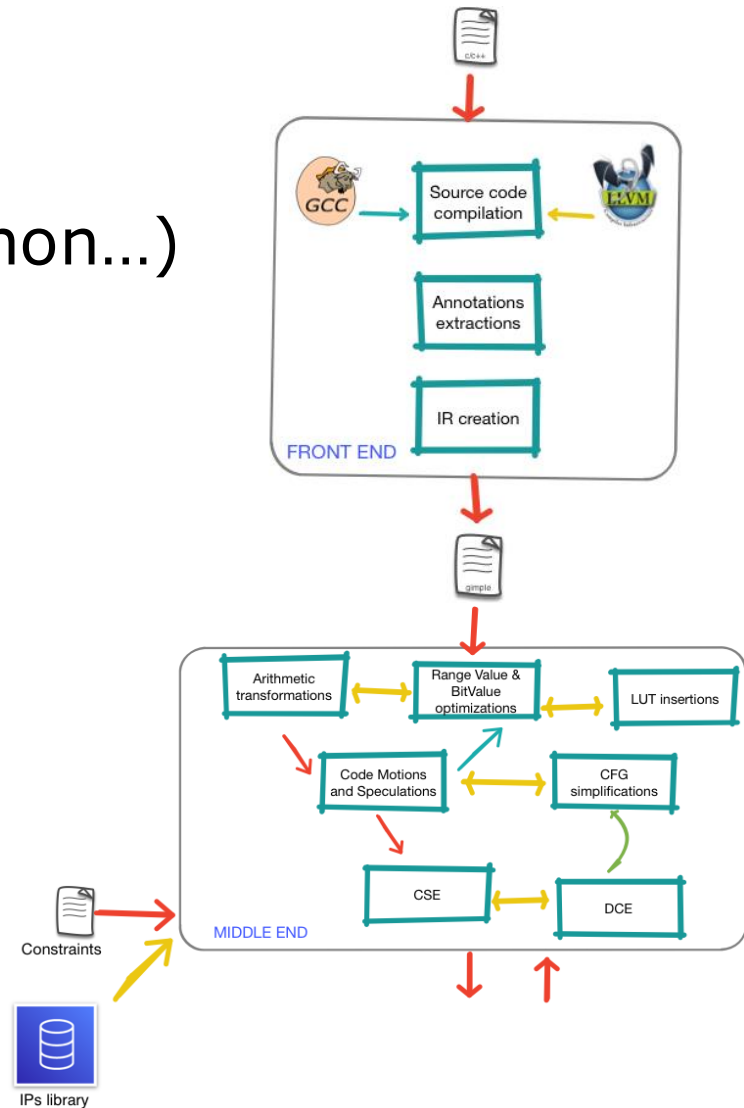
- ❑ While compilers map instructions to existing hardware, HLS tools must create a custom architecture
- ❑ Huge design space to explore
- ❑ Not all code is suitable to be translated into hardware

- ❑ Example HLS tool: Bambu, developed at DEIB
- ❑ (Probably) Most popular: Vitis HLS, from AMD/Xilinx
- ❑ Others: Intel HLS Compiler, SmartHLS (formerly LegUp) ...
- ❑ Compiler-like structure: front-end, middle-end, back-end
- ❑ Transformations applied on an Intermediate Representation (IR) of the input



## □ HLS inputs:

- High-level language specification (C/C++/Python...)
- Library of *characterized* modules
- Constraints (area/delay)
- (Optional) optimization directives



## □ HLS objectives:

- Minimum area (functional units, registers, memory, interconnect)
- Maximum speed (latency in clock cycles, clock period, throughput)

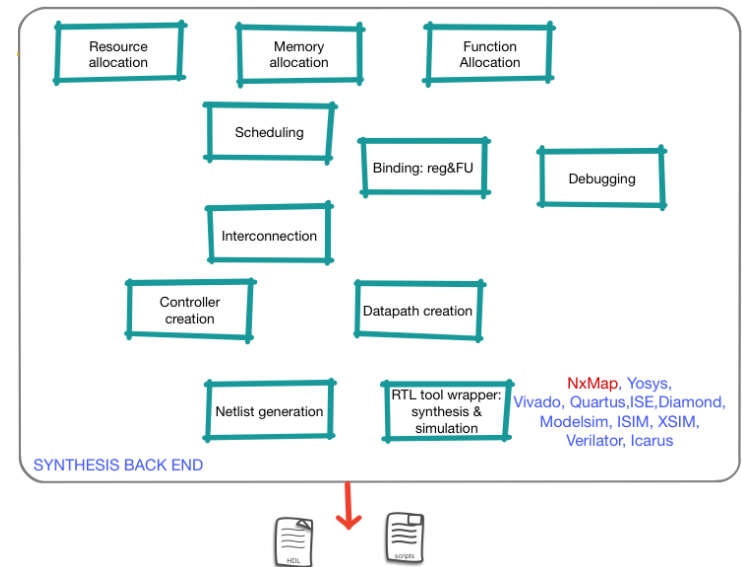
Generally, one parameter is set as a constraint,  
and the other one is optimized

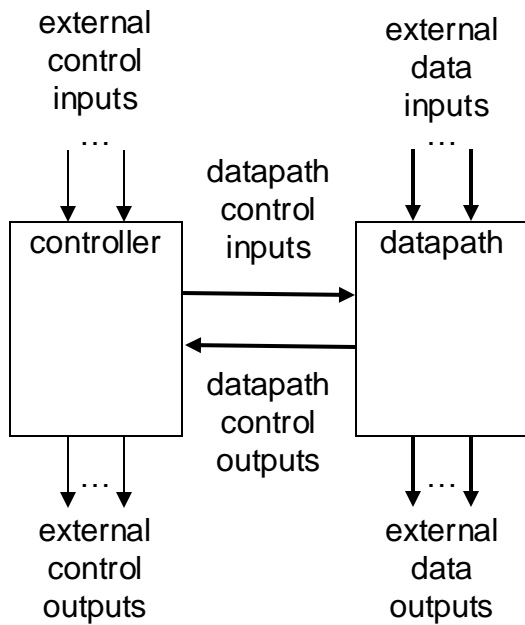
## □ Additional goals:

- Low power consumption, testability, more accurate estimation, fault tolerance...

## □ HLS output:

- Synthesizable Verilog/VHDL
- Controller + datapath (FSMD)
- Next steps: simulation, logic synthesis, implementation





controller and datapath

## Data-path:

- Functional resources: Perform operations on data (arithmetic and logic blocks).
- Memory resources: Store data (memory and registers).
- Interconnection resources: (muxes, busses and ports).

## Controller:

- Finite state machine (FSM)
- Controller + microprogram
- Synchronization scheme (e.g., global clock single phase with master-slave registers)

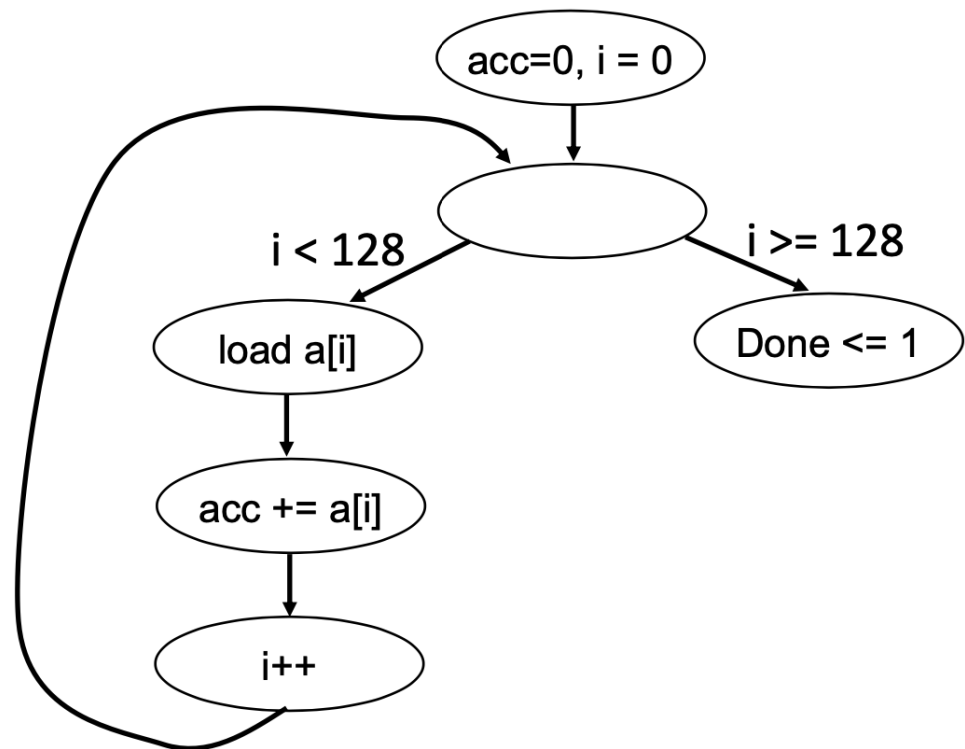


□ From high-level code to circuits:

1. Build FSM for the controller
2. Build datapath

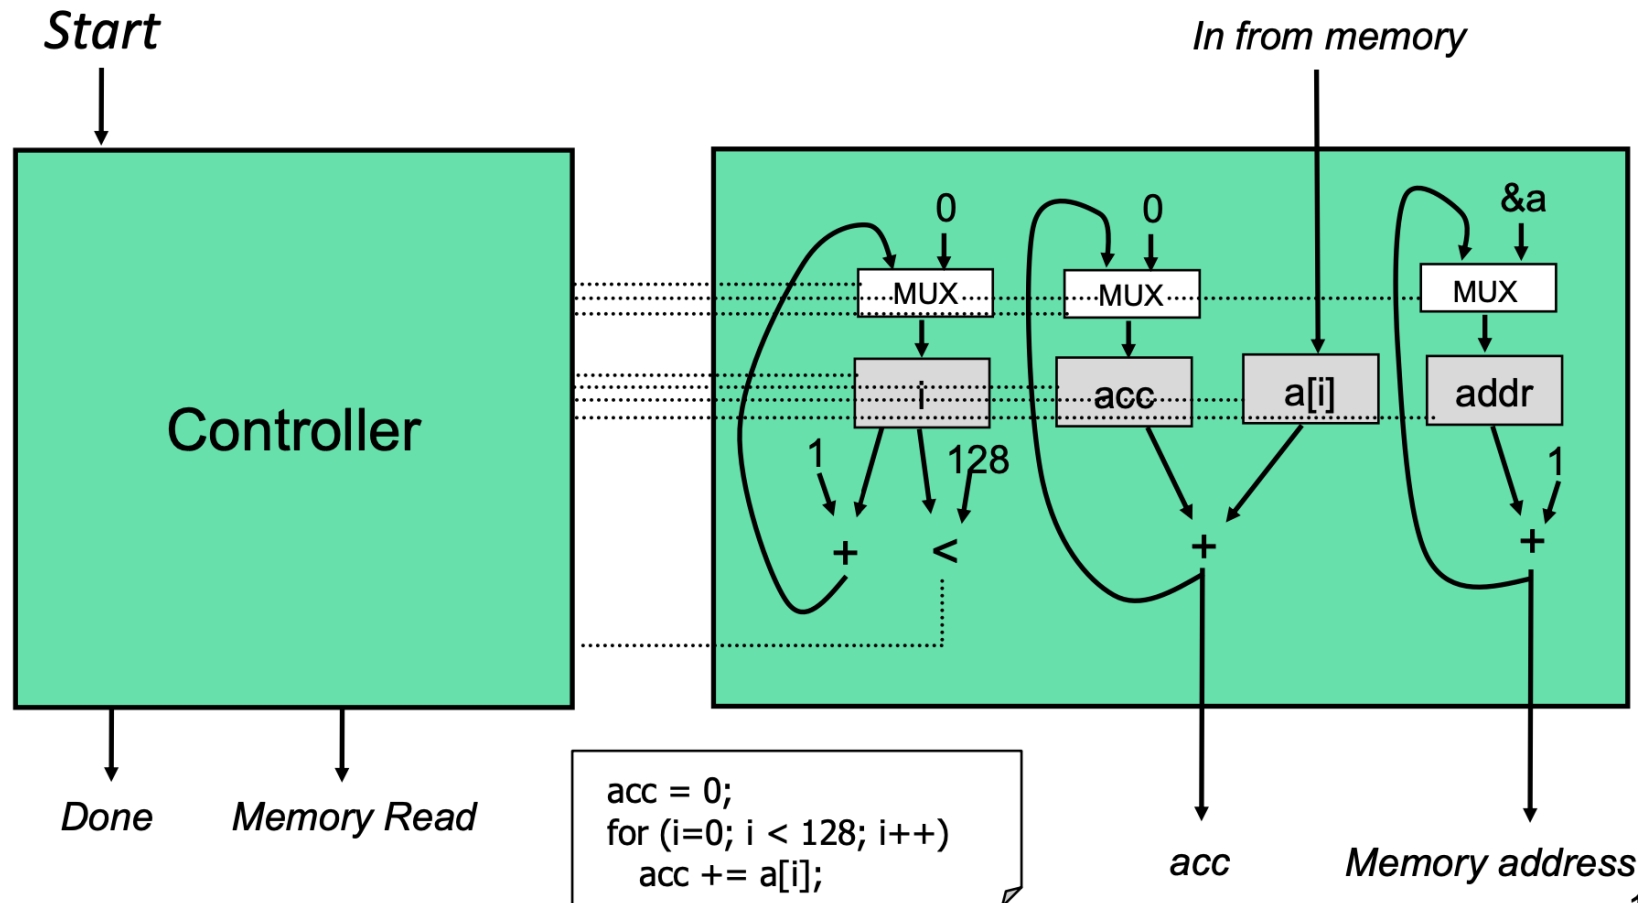
```
acc = 0;  
for (i=0; i < 128; i++)  
    acc += a[i];
```

Input code



FSM diagram


## □ Splitting controller and datapath



- ❑ HLS performs these steps automatically
  - **Scheduling** (determine when to perform each operation)
  - **Resource allocation** (allocate resources for each operation)
  - **Binding** (map operations to resources)

❑ Possible optimization: one less adder

❑ More optimization opportunities in the input program, in the FSM, in the datapath...



## ❑ Single Static Assignment form (SSA)

- A program is in SSA form if every variable is only assigned once

## ❑ 3-address form

- Single operator and at most three names

### Original

```
a := b + c
b := c + 1
d := b + c
a := a + 1
e := a + b
```

### SSA


```
a1 := b1 + c1
b2 := c1 + 1
d1 := b2 + c1
a2 := a1 + 1
e1 := a2 + b2
```

## Basic block

- Sequence of instructions with no labels (except at the first instruction) and no jumps (except at the last instruction)


## Control Flow Graph

- Nodes: basic blocks
- Edges: potential flow of control (loop, if/else, goto...)



## □ Data Flow Graph

- Nodes: operations
- Edges: data dependencies (read-after-write)
- One DFG for each basic block



- ❑ Definition: assignment of operations to time (control steps)
- ❑ Possibly within limits on hardware resources and timing
- ❑ Exploits potential parallelism
- ❑ Scheduling problems are *NP-hard*, so many techniques and heuristics can be used:
  - ASAP – As soon as possible
  - ALAP – As late as possible
  - List-based scheduling (resource constrained)
  - Integer Linear Programming
  - ...

## □ Inputs:

- IR (CDFG, SSA form)
- Operation delays in cycles
- Timing/resource constraints

## □ Output:


- Start time for every operation in the IR

## □ Goal:

- Best latency/area trade-off



## □ Example of different scheduling solutions



**1 adder & 1 mult.**

**2 adder & 1 mult.**

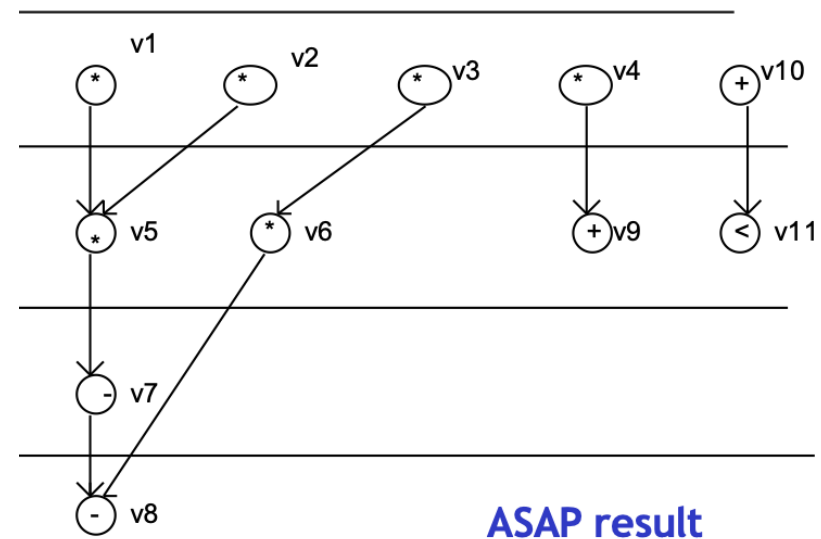

**step 1: +1**  
**step 2: +2**  
**step 3: X1**  
**step 4: +3, X2**

**step 1: +1, +2**  
**step 2: X1**  
**step 3: +3, X2**

## □ Simplest model: Minimum Latency Unconstrained Scheduling

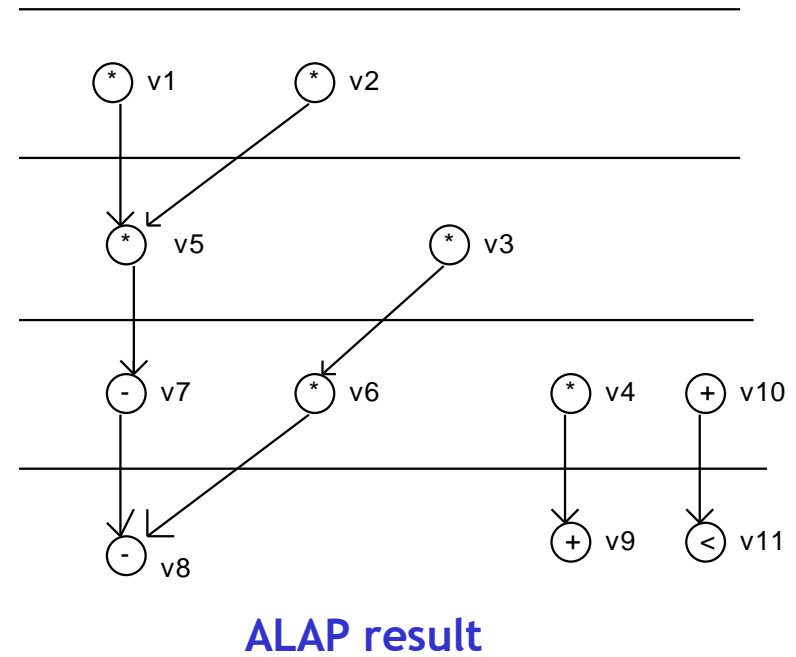

- All operations have bounded delays expressed as number of clock cycles, no area constraints

## □ ASAP scheduling




## □ Dual algorithm - **ALAP scheduling**

- Latency-constrained scheduling
- Maximum latency = ASAP latency



## □ Mobility:

- Difference between ASAP and ALAP start time
- Operations with mobility = 0 are on the critical path



- ❑ More realistic model: **Constrained Scheduling**
  - All operations have bounded delays expressed as number of clock cycles
  - Minimize latency given constraints on resources (ML-RCS)
  - Minimize resources given a bound on latency (MR-LCS)
- ❑ NP-hard
  - Exact methods (ILP)
  - Heuristics (list-based, force-directed)

## □ List-based Scheduling

- Heuristic method for both ML-RCS and MR-LCS
- Does not guarantee optimal solution
- Greedy strategy,  $O(n)$  time complexity

## □ Algorithm:


1. Construct a priority list (based e.g. on operation mobility)
2. While not all operations scheduled:
  1. For each available resource, select an operation in the ready list following the descending priority
  2. Assign these operations to the current clock cycle
  3. Update the ready list
  4. Clock cycle ++

- Exercise: schedule the given basic block from an HLS IR using list-based scheduling.

```
tmp_1 = a * b;  
tmp_2 = c * d;  
x_1 = tmp_1 * tmp_2;  
x_2 = a * d;  
x_3 = x_2 - x_1;  
x_4 = x_2 + x_3;  
x_5 = c * 3;  
x_6 = x_4 + x_4;  
x_7 = tmp_1 * x_5;  
out_1 = x_6 * x_7;
```


- Use mobility to create the priority list
- Available functional units:
  - 1 adder, latency 1 clock cycle
  - 1 multiplier, latency 2 clock cycles
  - 1 ALU, latency 2 clock cycles
- The ALU can perform any kind of arithmetic operations
- a, b, c, and d are input constants

```
tmp_1 = a * b;
tmp_2 = c * d;
x_1 = tmp_1 * tmp_2;
x_2 = a * d;
x_3 = x_2 - x_1;
x_4 = x_2 + x_3;
x_5 = c * 3;
x_6 = x_4 + x_4;
x_7 = tmp_1 * x_5;
out 1 = x_6 * x_7;
```



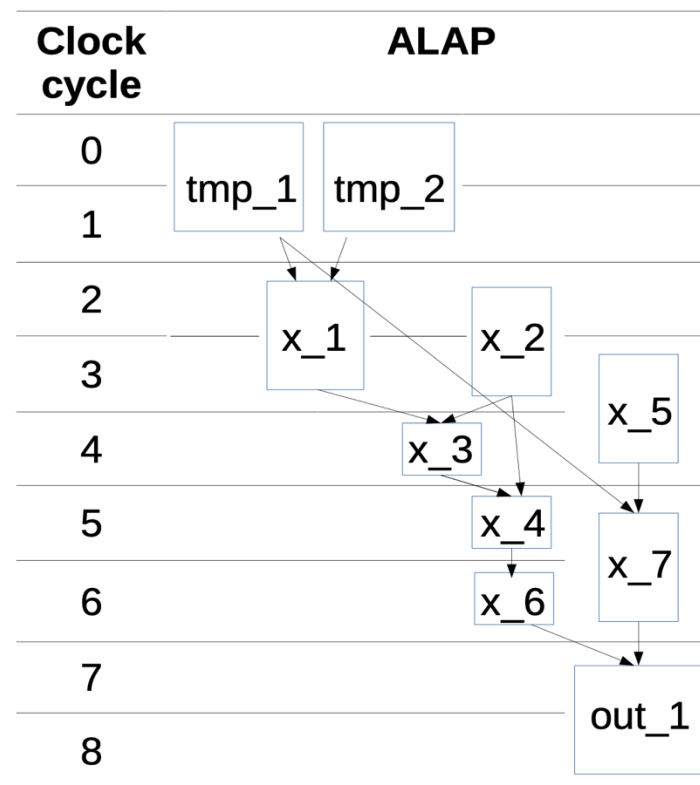
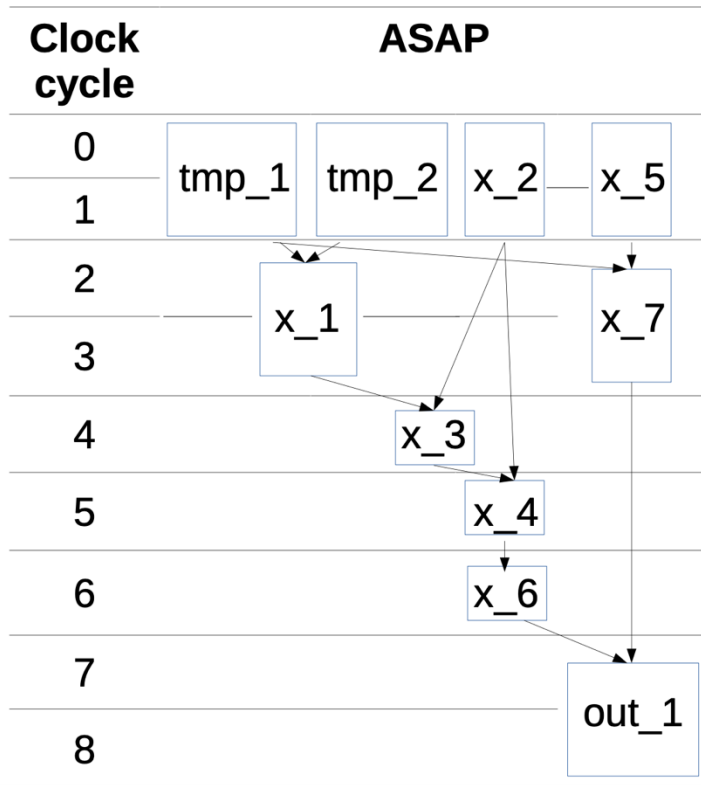


## Priority list = mobility



Clock cycle	ASAP			
0	tmp_1	tmp_2	x_2	x_5
1				
2				
3		x_1		x_7
4			x_3	
5				x_4
6			x_6	
7				out_1
8				

## □ Priority list = mobility



Mobility<sub>tmp\_1</sub> = 0  
 Mobility<sub>tmp\_2</sub> = 0  
 Mobility<sub>x\_2</sub> = 2  
 Mobility<sub>x\_5</sub> = 3  
 Mobility<sub>x\_1</sub> = 0  
 Mobility<sub>x\_7</sub> = 3  
 Mobility<sub>x\_3</sub> = 0  
 Mobility<sub>x\_4</sub> = 0  
 Mobility<sub>x\_6</sub> = 0  
 Mobility<sub>out\_1</sub> = 0

## □ List-based scheduling

Clock cycle	ADDER	MULTIPLIER	ALU
1		tmp_1 = a * b;	tmp_2 = c * d;
2		busy	busy

Ready list:

```
tmp_1 = a * b;  
tmp_2 = c * d;  
x_1 = tmp_1 * tmp_2;  
x_2 = a * d;  
x_3 = x_2 - x_1;  
x_4 = x_2 + x_3;  
x_5 = c * 3;  
x_6 = x_4 + x_4;  
x_7 = tmp_1 * x_5;  
out_1 = x_6 * x_7;
```

```
Mobilitytmp_1 = 0  
Mobilitytmp_2 = 0  
Mobilityx_2 = 2  
Mobilityx_5 = 3  
Mobilityx_1 = 0  
Mobilityx_7 = 3  
Mobilityx_3 = 0  
Mobilityx_4 = 0  
Mobilityx_6 = 0  
Mobilityout_1 = 0
```

## □ List-based scheduling

Clock cycle	ADDER	MULTIPLIER	ALU
1		tmp_1 = a * b;	tmp_2 = c * d;
2		busy	busy
3		x_1 = tmp_1 * tmp_2;	x_2 = a * d;
4		busy	busy

Ready list:

```
x_1 = tmp_1 * tmp_2;  
x_2 = a * d;  
x_3 = x_2 - x_1;  
x_4 = x_2 + x_3;  
x_5 = c * 3;  
x_6 = x_4 + x_4;  
x_7 = tmp_1 * x_5;  
out_1 = x_6 * x_7;
```

```
Mobilityx_2 = 2  
Mobilityx_5 = 3  
Mobilityx_1 = 0  
Mobilityx_7 = 3  
Mobilityx_3 = 0  
Mobilityx_4 = 0  
Mobilityx_6 = 0  
Mobilityout_1 = 0
```

## □ List-based scheduling

Clock cycle	ADDER	MULTIPLIER	ALU
1		tmp_1 = a * b;	tmp_2 = c * d;
2		busy	busy
3		x_1 = tmp_1 * tmp_2;	x_2 = a * d;
4		busy	busy
5	x_3 = x_2 - x_1;	x_5 = c * 3;	

Ready list:

```
x_3 = x_2 - x_1;  
x_4 = x_2 + x_3;  
x_5 = c * 3;  
x_6 = x_4 + x_4;  
x_7 = tmp_1 * x_5;  
out_1 = x_6 * x_7;
```

Mobility<sub>x\_5</sub> = 3

Mobility<sub>x\_7</sub> = 3

Mobility<sub>x\_3</sub> = 0

Mobility<sub>x\_4</sub> = 0

Mobility<sub>x\_6</sub> = 0

Mobility<sub>out\_1</sub> = 0

## List-based scheduling

Clock cycle	ADDER	MULTIPLIER	ALU
1		tmp_1 = a * b;	tmp_2 = c * d;
2		busy	busy
3		x_1 = tmp_1 * tmp_2;	x_2 = a * d;
4		busy	busy
5	x_3 = x_2 - x_1;	x_5 = c * 3;	
6	x_4 = x_2 + x_3;	busy	

Ready list:

```

x_4 = x_2 + x_3;
x_6 = x_4 + x_4;
x_7 = tmp_1 * x_5;
out_1 = x_6 * x_7;

```

Mobility<sub>x\_7</sub> = 3

Mobility<sub>x\_4</sub> = 0

Mobility<sub>x\_6</sub> = 0

Mobility<sub>out\_1</sub> = 0

## List-based scheduling

Clock cycle	ADDER	MULTIPLIER	ALU
1		tmp_1 = a * b;	tmp_2 = c * d;
2		busy	busy
3		x_1 = tmp_1 * tmp_2;	x_2 = a * d;
4		busy	busy
5	x_3 = x_2 - x_1;	x_5 = c * 3;	
6	x_4 = x_2 + x_3;	busy	
7	x_6 = x_4 + x_4;	x_7 = tmp_1 * x_5;	
8		busy	

Ready list:

```
x_6 = x_4 + x_4;
x_7 = tmp_1 * x_5;
out_1 = x_6 * x_7;
```

Mobility<sub>x\_7</sub> = 3

Mobility<sub>x\_6</sub> = 0

Mobility<sub>out\_1</sub> = 0

## List-based scheduling

Clock cycle	ADDER	MULTIPLIER	ALU
1		tmp_1 = a * b;	tmp_2 = c * d;
2		busy	busy
3		x_1 = tmp_1 * tmp_2;	x_2 = a * d;
4		busy	busy
5	x_3 = x_2 - x_1;	x_5 = c * 3;	
6	x_4 = x_2 + x_3;	busy	
7	x_6 = x_4 + x_4;	x_7 = tmp_1 * x_5;	
8		busy	
9		out_1 = x_6 * x_7;	
10		busy	

Ready list: `out_1 = x_6 * x_7;`

Mobility<sub>out\_1</sub> = 0



[https://colab.research.google.com/drive/1muPx2tSure0GmyljW4q5G0J6QdHclXZE?usp=share\\_link](https://colab.research.google.com/drive/1muPx2tSure0GmyljW4q5G0J6QdHclXZE?usp=share_link)

- ❑ Stanford course CS149, Parallel Computing
- ❑ <https://halide-lang.org/>
  - Links to source code, publications, tutorials, all sorts of useful material
- ❑ “Digital & Computational Photography”, MIT course by Frédo Durand (lectures 14 to 17)  
<https://stellar.mit.edu/S/course/6/sp15/6.815/materials.html>

- ❑ M. Fingeroff, “High-Level Synthesis Blue Book”
  - [https://cse.usf.edu/~haozheng/teach/cda4253/doc/hls/hls\\_bluebook\\_uv.pdf](https://cse.usf.edu/~haozheng/teach/cda4253/doc/hls/hls_bluebook_uv.pdf)
- ❑ Bambu tutorial from [DATE 2022](#) (somewhat outdated, but still useful)
- ❑ In-depth Vitis HLS [user guide and tutorials](#)
- ❑ J. Cong et al., “FPGA HLS Today: Successes, Challenges, and Opportunities”  
<https://doi.org/10.1145/3530775>
- ❑ G. Gozzi et al., “SPARTA: High-Level Synthesis of Parallel Multi-Threaded Accelerators”  
<https://doi.org/10.1145/3677035>