

A decorative graphic on the left side of the slide consists of a network of lines and circles. The lines are primarily dark blue, with some transitioning to a lighter blue and then to a light beige color towards the bottom. The circles are small and white with dark outlines. The overall pattern resembles a stylized circuit board or a network diagram.

# 058165 - PARALLEL COMPUTING

Fabrizio Ferrandi

a.a. 2022-2023

# Acknowledge

## Material from:

Parallel Computing lectures from prof. Kayvon and prof. Olukotun Stanford University

**An example:**  
**Programming with ISPC**

# ISPC

- Intel SPMD Program Compiler (ISPC)
- SPMD: single program multiple data
- <http://ispc.github.com/>
- A great read: “The Story of ISPC” (by Matt Pharr)
  - <https://pharr.org/matt/blog/2018/04/30/ispc-all.html>

# Recall: example program from last class

Compute  $\sin(x)$  using Taylor expansion:  $\sin(x) = x - x^3/3! + x^5/5! - x^7/7! + \dots$

for each element of an array of N floating-point numbers

```
void sinx(int N, int terms, float* x, float* result)
{
    for (int i=0; i<N; i++)
    {
        float value = x[i];
        float numer = x[i] * x[i] * x[i];
        int denom = 6; // 3!
        int sign = -1;

        for (int j=1; j<=terms; j++)
        {
            value += sign * numer / denom;
            numer *= x[i] * x[i];
            denom *= (2*j+2) * (2*j+3);
            sign *= -1;
        }

        result[i] = value;
    }
}
```

# Invoking sinx()

C++ code: main.cpp


```
#include "sinx.h"

int main(int argc, void** argv) {
    int N = 1024;
    int terms = 5;
    float* x = new float[N];
    float* result = new float[N];

    // initialize x here

    sinx(N, terms, x, result);

    return 0;
}
```



Call to sinx()  
Control transferred to sinx() func

Return from sinx()  
Control transferred back to main()

C++ code: sinx.cpp

```
void sinx(int N, int terms, float* x, float* result)
{
    for (int i=0; i<N; i++)
    {
        float value = x[i];
        float numer = x[i] * x[i] * x[i];
        int denom = 6; // 3!
        int sign = -1;

        for (int j=1; j<=terms; j++)
        {
            value += sign * numer / denom;
            numer *= x[i] * x[i];
            denom *= (2*j+2) * (2*j+3);
            sign *= -1;
        }

        result[i] = value;
    }
}
```

# sinx() in ISPC

## C++ code: main.cpp

```
#include "sinx_ispc.h"

int main(int argc, void** argv) {
    int N = 1024;
    int terms = 5;
    float* x = new float[N];
    float* result = new float[N];

    // initialize x here

    // execute ISPC code
    ispc_sinx(N, terms, x, result);
    return 0;
}
```

## SPMD programming abstraction:

Call to ISPC function spawns “gang” of ISPC  
“program instances”

All instances run ISPC code concurrently

Each instance has its own copy of local variables  
(blue variables in code, we’ll talk about “uniform” later)

Upon return, all instances have completed

## ISPC code: sinx.ispc

```
export void ispc_sinx(
    uniform int N,
    uniform int terms,
    uniform float* x,
    uniform float* result)
{
    // assume N % programCount = 0
    for (uniform int i=0; i<N; i+=programCount)
    {
        int idx = i + programIndex;
        float value = x[idx];
        float numer = x[idx] * x[idx] * x[idx];
        uniform int denom = 6; // 3!
        uniform int sign = -1;

        for (uniform int j=1; j<=terms; j++)
        {
            value += sign * numer / denom
            numer *= x[idx] * x[idx];
            denom *= (2*j+2) * (2*j+3);
            sign *= -1;
        }
        result[idx] = value;
    }
}
```

# Invoking sinx() in ISPC

C++ code: main.cpp

```
#include "sinx_ispc.h"

int main(int argc, void** argv) {
    int N = 1024;
    int terms = 5;
    float* x = new float[N];
    float* result = new float[N];

    // initialize x here

    // execute ISPC code
    ispc_sinx(N, terms, x, result);
    return 0;
}
```

## SPMD programming abstraction:


Call to ISPC function spawns “gang” of ISPC “program instances”

All instances run ISPC code concurrently

Each instance has its own copy of local variables

Upon return, all instances have completed

In this illustration programCount = 8





# sinx() in ISPC

“Interleaved” assignment of array elements to program instances

## C++ code: main.cpp

```
#include "sinx_ispc.h"

int main(int argc, void** argv) {
    int N = 1024;
    int terms = 5;
    float* x = new float[N];
    float* result = new float[N];

    // initialize x here

    // execute ISPC code
    ispc_sinx(N, terms, x, result);
    return 0;
}
```

## ISPC language keywords:

**programCount**: number of simultaneously executing instances in the gang (uniform value)

**programIndex**: id of the current instance in the gang. (a non-uniform value: “varying”)

**uniform**: A type modifier. All instances have the same value for this variable. Its use is purely an optimization. Not needed for correctness.


## ISPC code: sinx.ispc

```
export void ispc_sinx(
    uniform int N,
    uniform int terms,
    uniform float* x,
    uniform float* result)
{
    // assumes N % programCount == 0
    for (uniform int i=0; i<N; i+=programCount)
    {
        int idx = i + programIndex;
        float value = x[idx];
        float numer = x[idx] * x[idx] * x[idx];
        uniform int denom = 6; // 3!
        uniform int sign = -1;

        for (uniform int j=1; j<=terms; j++)
        {
            value += sign * numer / denom;
            numer *= x[idx] * x[idx];
            denom *= (2*j+2) * (2*j+3);
            sign *= -1;
        }
        result[idx] = value;
    }
}
```

# Interleaved assignment of program instances to loop iterations

Elements of output array (results)



“Gang” of ISPC program instances

In this illustration: gang contains eight instances:

programCount = 8

# ISPC implements the gang abstraction using SIMD instructions

C++ code: main.cpp

```
#include "sinx_ispc.h"

int main(int argc, void** argv) {
    int N = 1024;
    int terms = 5;
    float* x = new float[N];
    float* result = new float[N];

    // initialize x here

    // execute ISPC code
    ispc_sinx(N, terms, x, result);
    return 0;
}
```

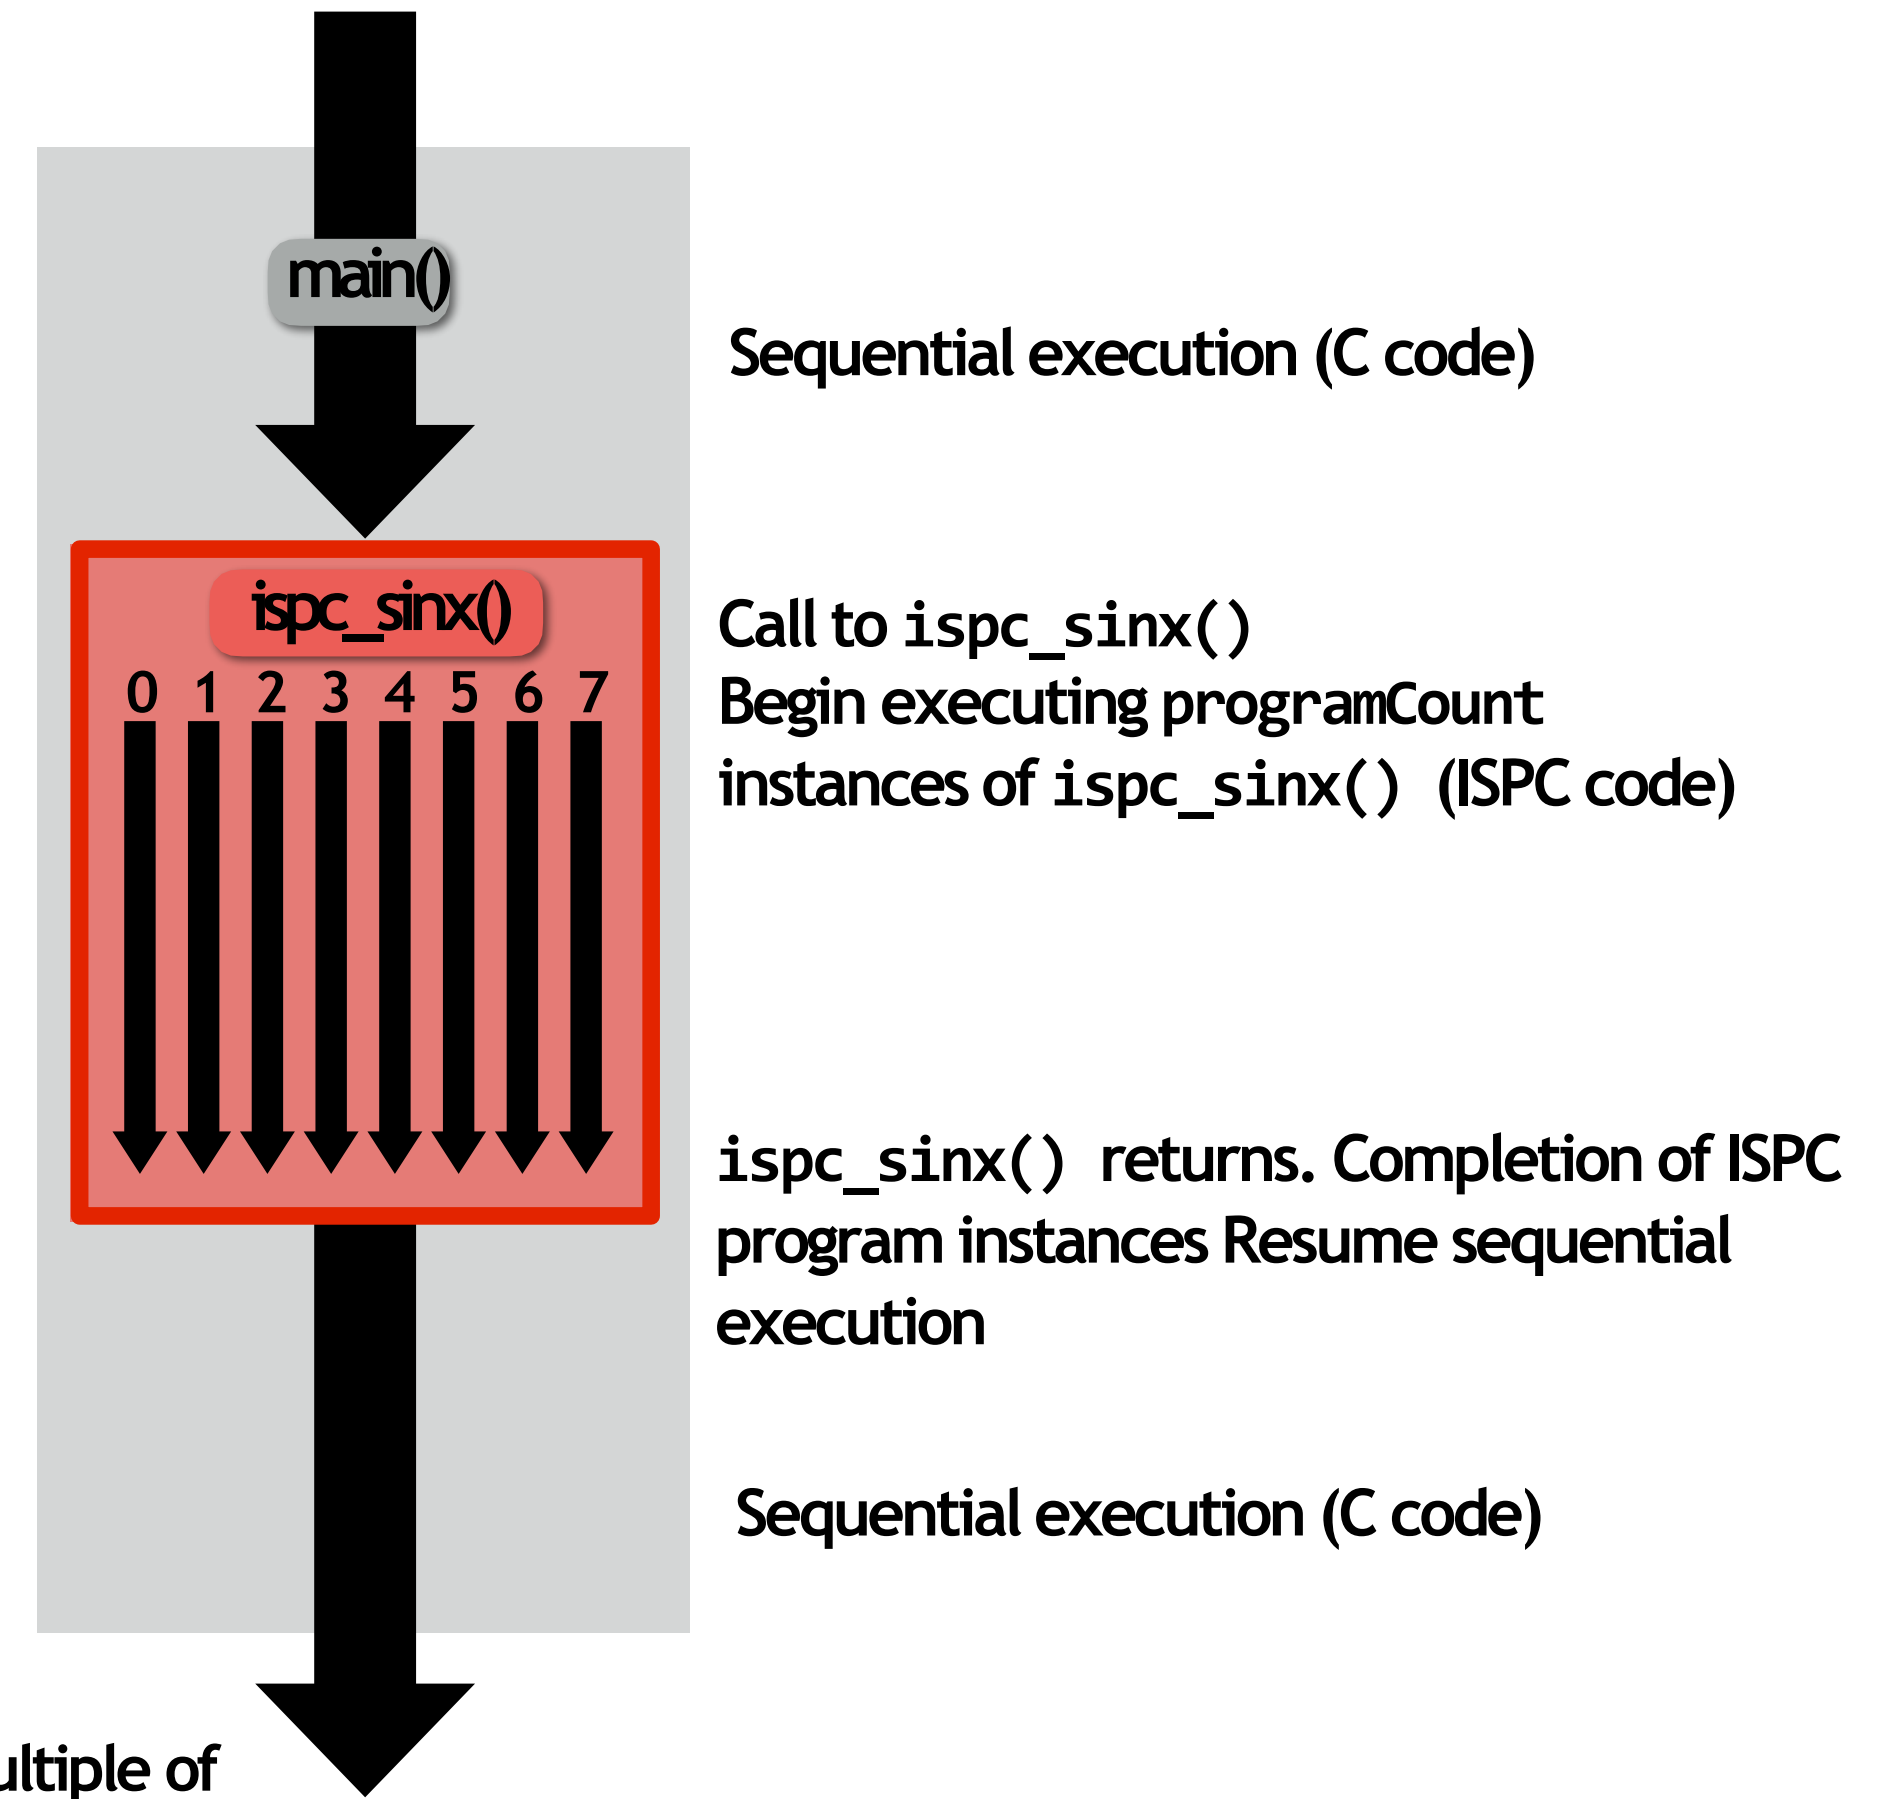
## SPMD programming abstraction:

Call to ISPC function spawns “gang” of ISPC “program instances” All instances run ISPC code simultaneously  
Upon return, all instances have completed

## ISPC compiler generates SIMD implementation:

Number of instances in a gang is the SIMD width of the hardware (or a small multiple of SIMD width) ISPC compiler generates a C++ function binary (.o) whose body contains SIMD instructions

C++ code links against generated object file as usual



# sinx() in ISPC: version 2

“Blocked” assignment of array elements to program instances

## C++ code: main.cpp

```
#include "sinx_ispc.h"

int main(int argc, void** argv) {
    int N = 1024;
    int terms = 5;
    float* x = new float[N];
    float* result = new float[N];

    // initialize x here

    // execute ISPC code
    ispc_sinx_v2(N, terms, x, result);
    return 0;
}
```

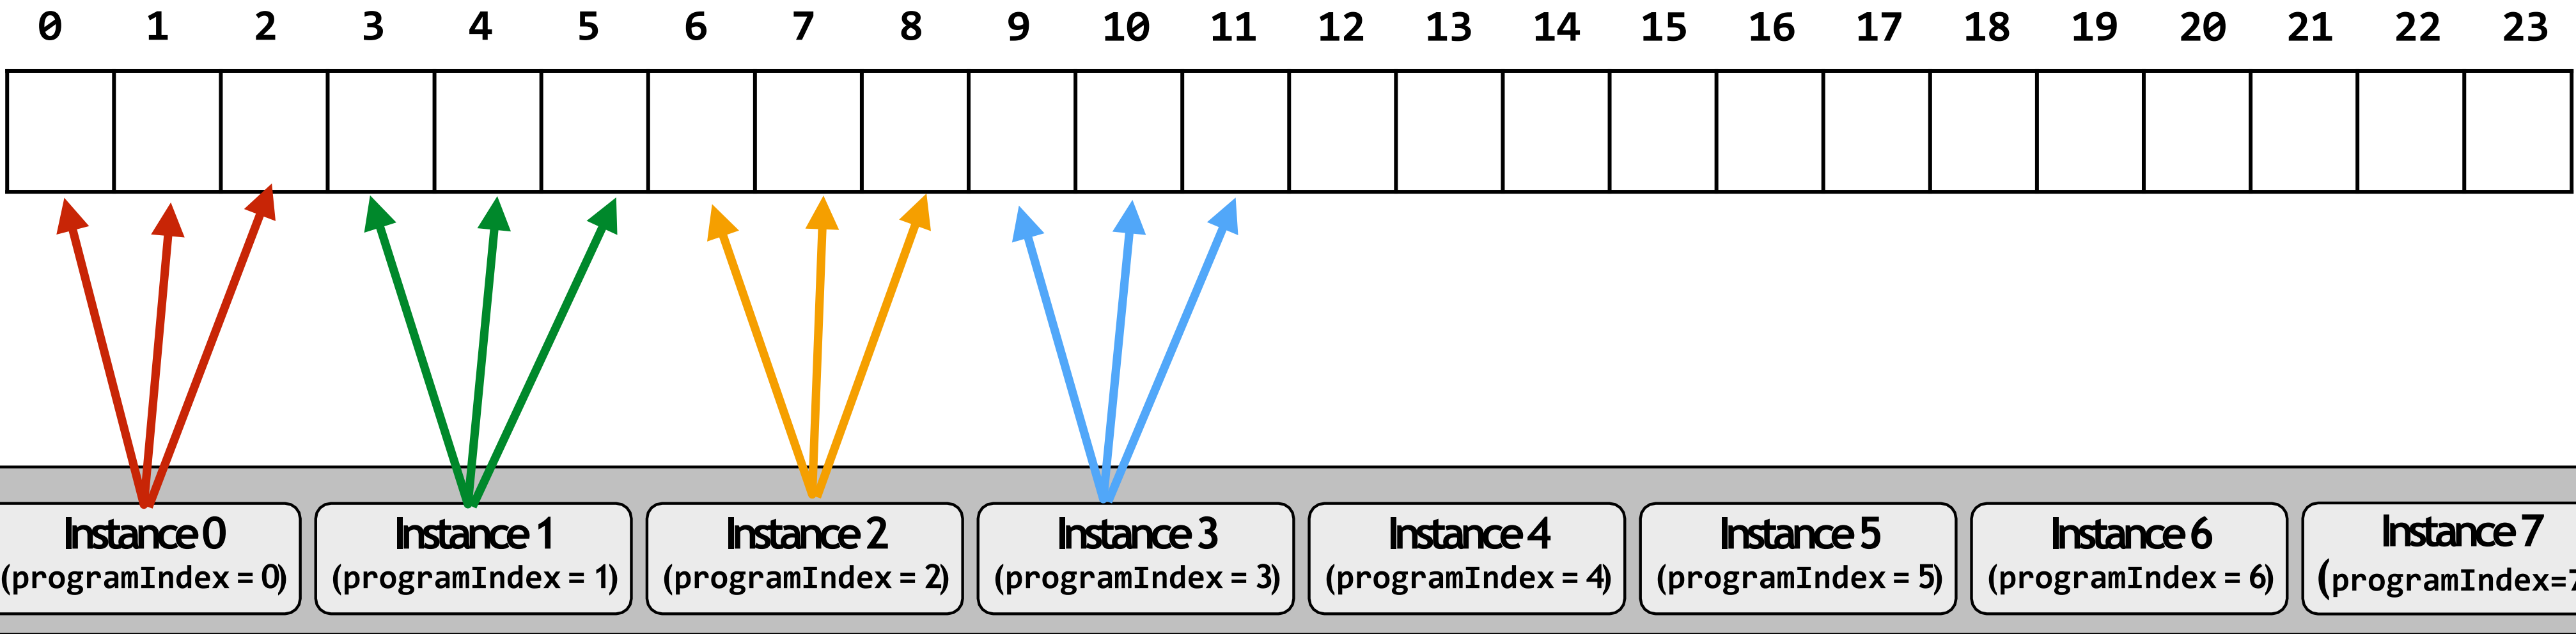
## ISPC code: sinx.ispc

```
export void ispc_sinx_v2(
    uniform int N,
    uniform int terms,
    uniform float* x,
    uniform float* result)
{
    // assume N % programCount = 0
    uniform int count = N / programCount;
    int start = programIndex * count;
    for (uniform int i=0; i<count; i++)
    {
        int idx = start + i;
        float value = x[idx];
        float numer = x[idx] * x[idx] * x[idx];
        uniform int denom = 6; // 3!
        uniform int sign = -1;

        for (uniform int j=1; j<=terms; j++)
        {
            value += sign * numer / denom
            numer *= x[idx] * x[idx];
            denom *= (j+3) * (j+4);
            sign *= -1;
        }
        result[idx] = value;
    }
}
```

# Blocked assignment of program instances to loop iterations

Elements of output array (results)



“Gang” of ISPC program instances


In this illustration: gang contains eight instances:

programCount = 8

# Schedule: interleaved assignment

“Gang” of ISPC program instances

Gang contains four instances: programCount = 8



A single “packed vector load” instruction (`vmovaps *`) efficiently implements:  
`float value = x[idx];`  
for all program instances, since the eight values are contiguous in memory


```
...  
// assumes N % programCount = 0  
for (uniform int i=0; i<N; i+=programCount)  
{  
    int idx = i + programIndex;  
    float value = x[idx];  
}  
...
```

\* see `_mm256_load_ps()` intrinsic function

# Schedule: blocked assignment

“Gang” of ISPC program instances

Gang contains four instances: `programCount = 8`



`float value = x[idx];`  
For all program instances now touches eight non-contiguous values in memory. Need “gather” instruction (`vgatherdps *`) to implement (gather is a more complex, and more costly SIMD instruction...)

```
uniform int count = N / programCount;  
int start = programIndex * count;  
for (uniform int i=0; i<count; i++) {  
    int idx = start + i;  
    float value = x[idx];  
    ...  
}
```

\* see `_mm256_i32gather_ps()` intrinsic function



# Raising level of abstraction with foreach

## C++ code: main.cpp

```
#include "sinx_ispc.h"

int N = 1024;
int terms = 5;
float* x = new float[N];
float* result = new float[N];

// initialize x here

// execute ISPC code
sinx(N, terms, x, result);
```

## ISPC code: sinx.ispc

```
export void ispc_sinx(
    uniform int N,
    uniform int terms,
    uniform float* x,
    uniform float* result)
{
    foreach (i = 0 ... N)
    {
        float value = x[i];
        float number = x[i] * x[i] * x[i];
        uniform int denom = 6; // 3!
        uniform int sign = -1;

        for (uniform int j=1; j<=terms; j++)
        {
            value += sign * number / denom;
            number *= x[i] * x[i];
            denom *= (2*j+2) * (2*j+3);
            sign *= -1;
        }
        result[i] = value;
    }
}
```

### foreach: key ISPC language construct

- foreach declares parallel loop iterations
  - Programmer says: these are the iterations the entire gang (not each instance) must perform
- ISPC implementation assigns iterations to program instances in the gang
  - Current ISPC implementation will perform a static interleaved assignment (but the abstraction permits a different assignment)

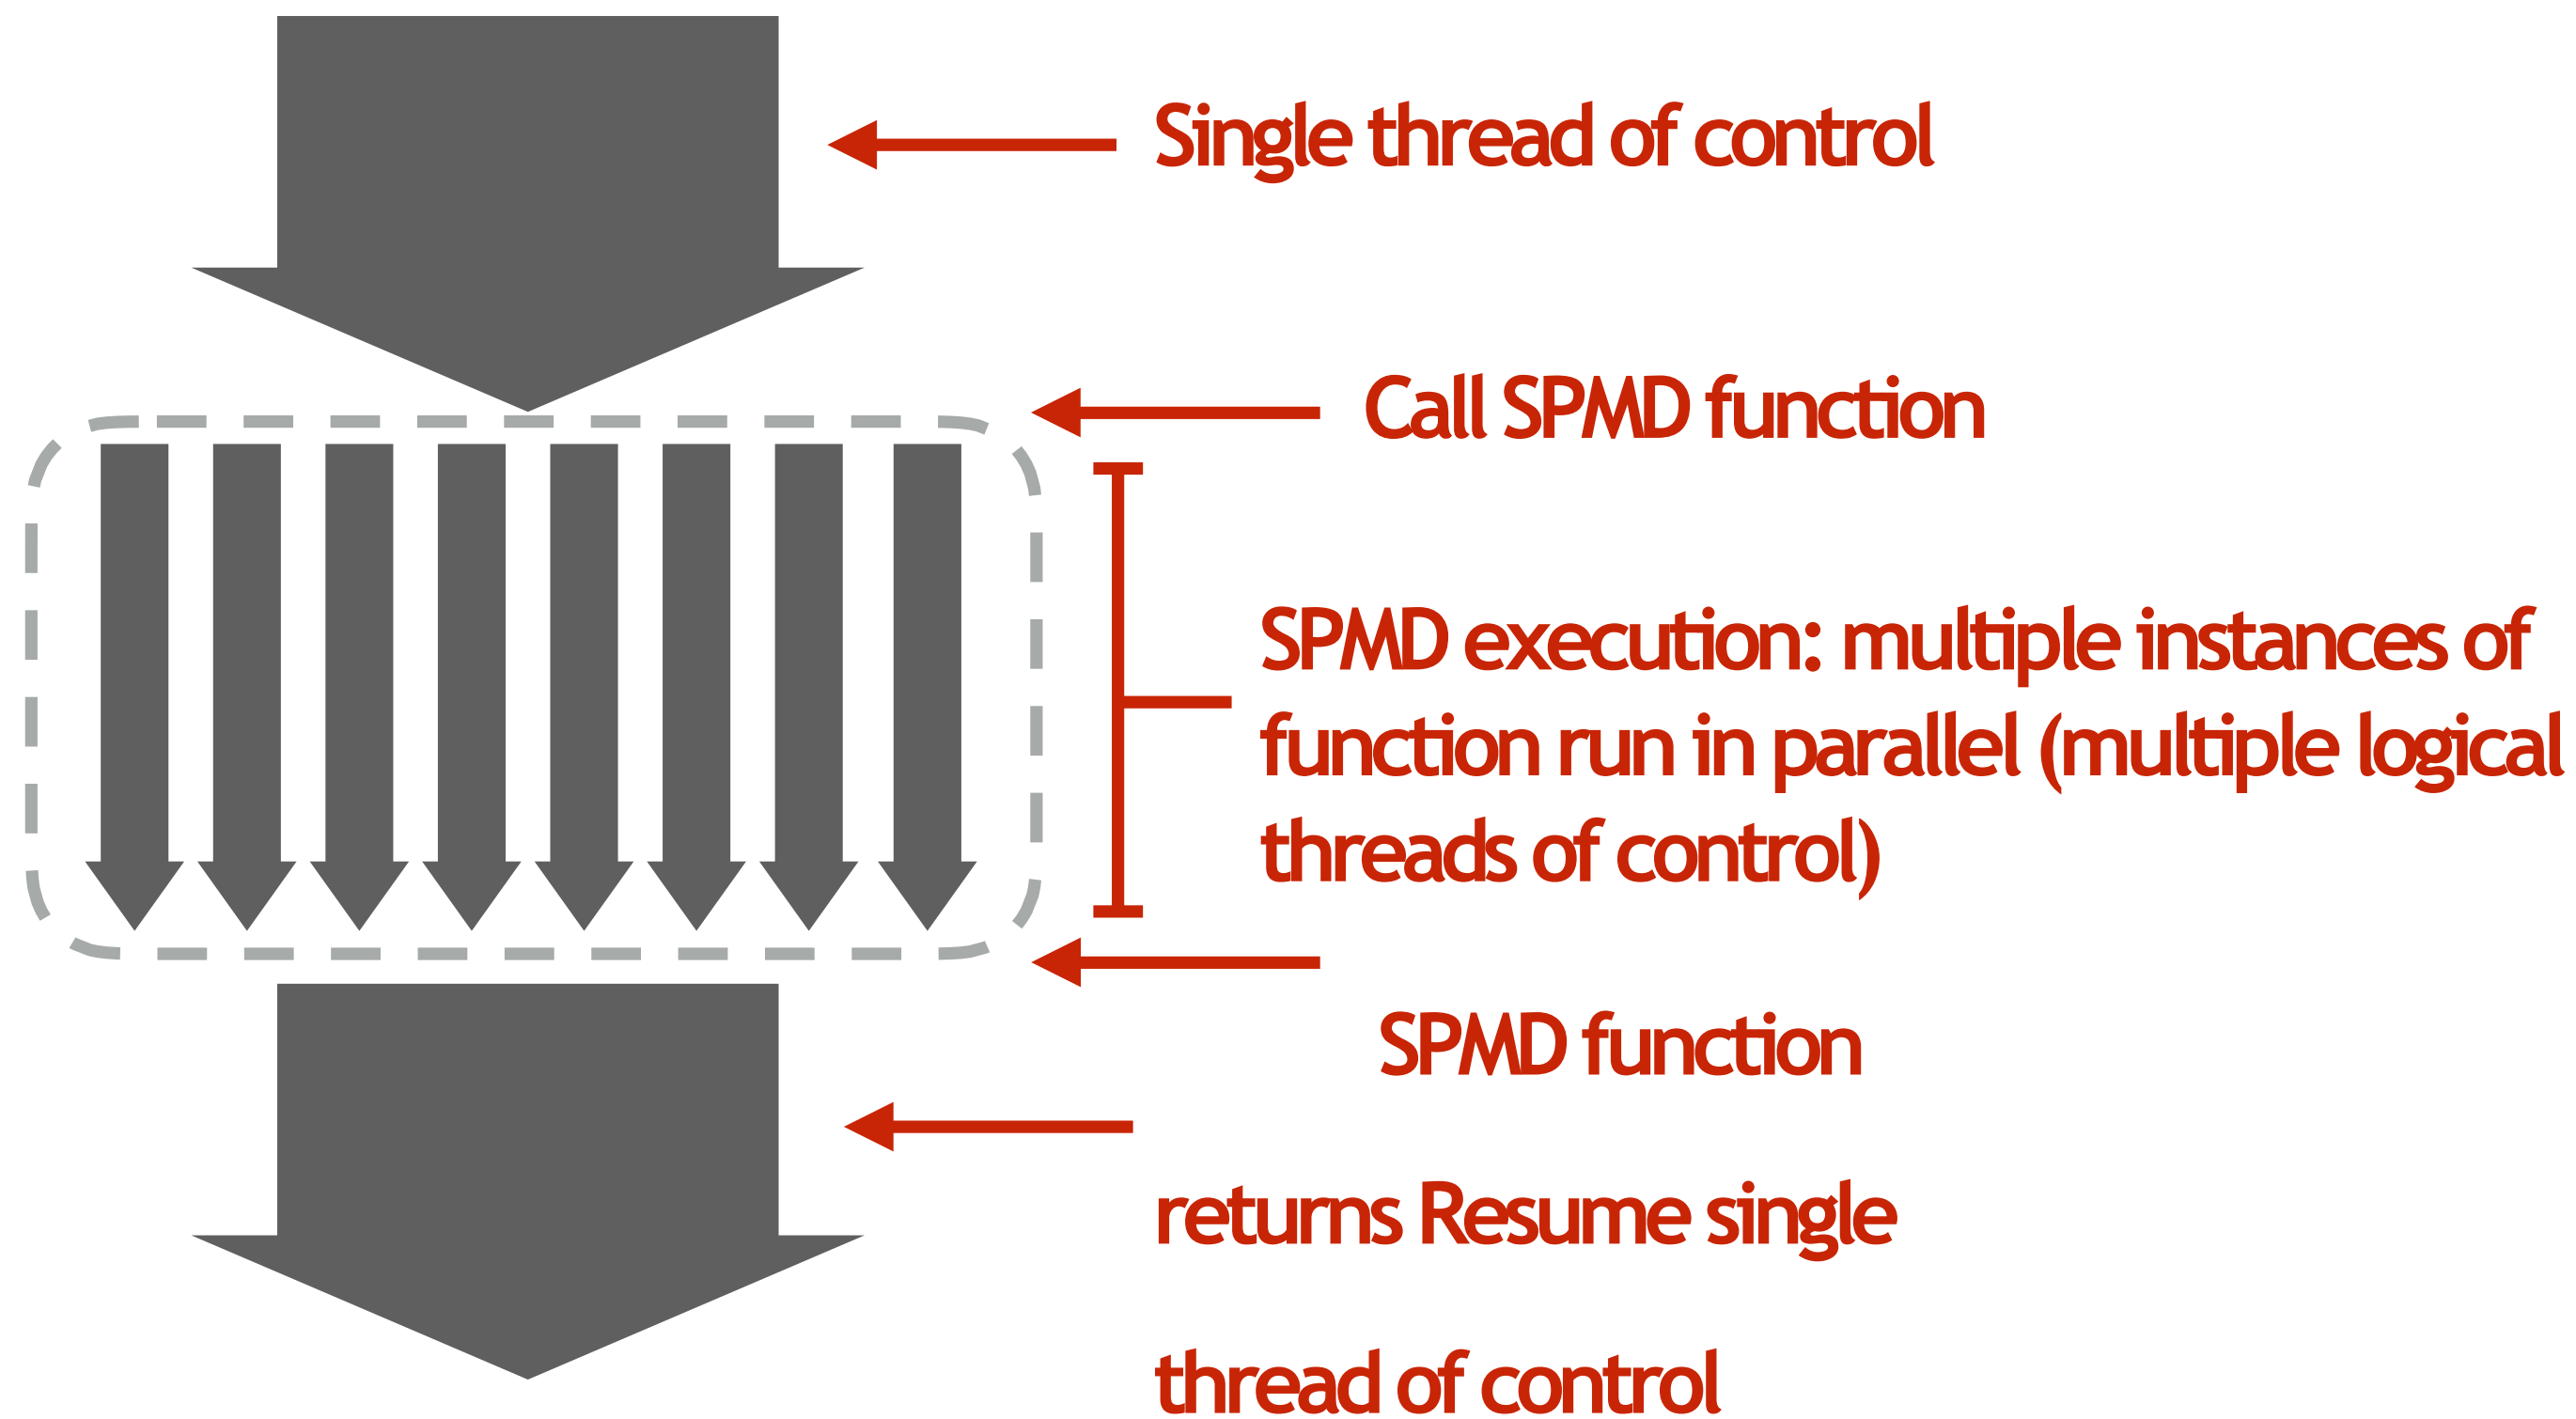


# ISPC: abstraction vs. implementation

- **Single program, multiple data (SPMD) programming model**
  - Programmer “thinks”: running a gang is spawning `programCount` logical instruction streams (each with a different value of `programIndex`)
  - This is the programming abstraction
  - Program is written in terms of this abstraction
- **Single instruction, multiple data (SIMD) implementation**
  - ISPC compiler emits vector instructions (e.g., AVX2, ARM NEON) that carry out the logic performed by a ISPC gang
  - ISPC compiler handles mapping of conditional control flow to vector instructions (by masking vector lanes, etc.)
- **Semantics of ISPC can be tricky**
  - SPMD abstraction + uniform values  
(allows implementation details to peek through abstraction a bit)

# SPMD programming model summary

- SPMD = “single program, multiple data”
- Define one function, run multiple instances of that function in parallel on different input arguments



# ISPC tasks

- The ISPC gang abstraction is implemented by SIMD instructions that execute within on thread running on one x86 core of a CPU.
- So all the code I've shown you in the previous slides would have executed on only one of the four cores of the myth machines.
- ISPC contains another abstraction: a “task” that is used to achieve multi-core execution. I'll let you read up about that.

# Part 2 of today's lecture

- **Three parallel programming models**
  - That differ in what communication abstractions they present to the programmer
  - Programming models are important because they (1) influence how programmers think when writing programs and (2) influence the design of parallel hardware platforms designed to execute them efficiently
- **Corresponding machine architectures**
  - Abstraction presented by the hardware to low-level software
- **We'll focus on differences in communication/synchronization**

# Three programming models (abstractions)

**1. Shared address space**

**2. Message passing**

**3. Data parallel**

# Shared address space model

# Review: a program's memory address space

- A computer's memory is organized as a array of bytes
- Each byte is identified by its “address” in memory (its position in this array)  
(in this class we assume memory is byte-addressable)

*“The byte stored at address 0x8 has the value 32.”*


*“The byte stored at address 0x10 (16) has the value 128.”*

In the illustration on the right, the program's memory address space is 32 bytes in size  
(so valid addresses range from 0x0 to 0x1F)

Address	Value
0x0	16
0x1	255
0x2	14
0x3	0
0x4	0
0x5	0
0x6	6
0x7	0
0x8	32
0x9	48
0xA	255
0xB	255
0xC	255
0xD	0
0xE	0
0xF	0
0x10	128
⋮	⋮
0x1F	0

# The implementation of the linear memory address space abstraction on a modern computer is complex

The instruction “load the value stored at address X into register R0” might involve a complex sequence of operations by multiple data caches and access to DRAM





# Shared address space model (abstraction)

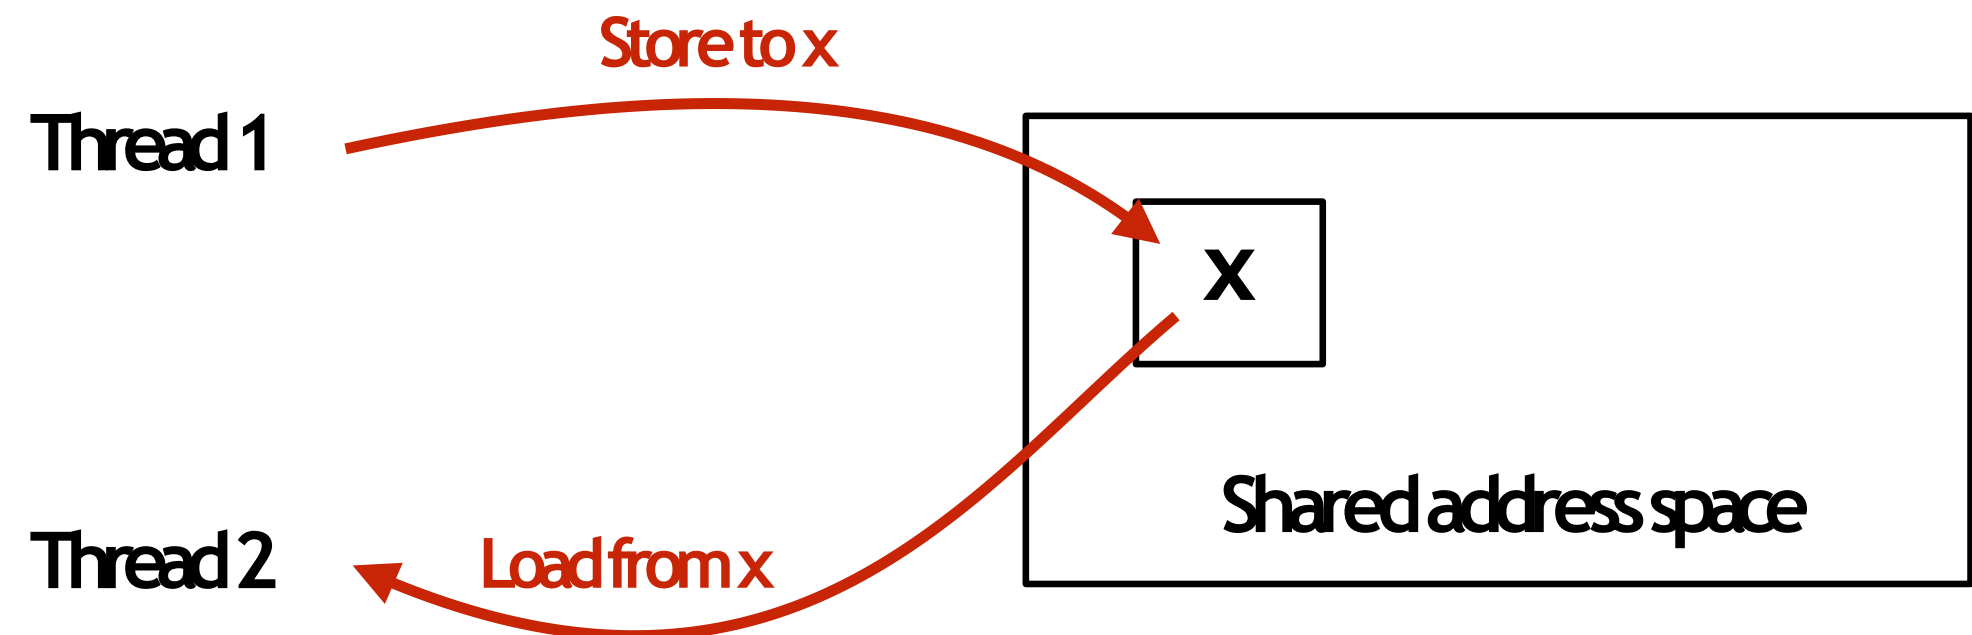
Threads communicate by reading/writing to locations in a shared address space (shared variables)

Thread 1:

```
int x = 0;  
spawn_thread(foo, &x);  
  
// write to address holding  
// contents of variable x  
x = 1;
```

Thread 2:

```
void foo(int* x) {  
  
// read from addr storing  
// contents of variable x  
while (x == 0) {}  
print x;  
}
```



**A common metaphor:  
A shared address space is like a bulletin board**

**(Everyone can read/write)**



# Coordinating access to shared variables with synchronization

Thread 1:

```
int x = 0;
Lock my_lock;

spawn_thread(foo, &x, &my_lock);

mylock.lock();
x++;
mylock.unlock();
```

Thread 2:

```
void foo(int* x, Lock* my_lock) {
    my_lock->lock();
    x++;
    my_lock->unlock();

    print(x);
}
```

# Review: why do we need mutual exclusion?

- Each thread executes
  - Load the value of variable x from a location in memory into register r1  
(this stores a copy of the value in memory in the register)
  - Add the contents of register r2 to register r1
  - Store the value of register r1 into the address storing the program variable x
- One possible interleaving: (let starting value of x=0, r2=1)

T1	T2
r1 ← x	T1 reads value 0
	T2 reads value 0
r1 ← r1 + r2	T1 sets value of its r1 to 1
	T2 sets value of its r1 to 1
X ← r1	T1 stores 1 to address of x
	T2 stores 1 to address of x

- Need this set of three instructions must be “atomic”

# Mechanisms for preserving atomicity

- Lock/unlock mutex around a critical section

```
mylock.lock();  
// critical section  
mylock.unlock();
```

- Some languages have first-class support for atomicity of code blocks

```
atomic {  
    // critical section  
}
```

- Intrinsic for hardware-supported atomic read-modify-write operations


```
atomicAdd(x, 10);
```

# Review: shared address space model




- **Threads communicate by:**
  - Reading/writing to shared variables in a shared address space
    - Inter-thread communication is implicit in memory loads/stores
  - Manipulating synchronization primitives
    - e.g., ensuring mutual exclusion via use of locks
- **This is a natural extension of sequential programming**
  - In fact, all our discussions in class have assumed a shared address space so far!

# Hardware implementation of a shared address space

Key idea: any processor can directly reference contents of any memory location



## Examples of interconnects

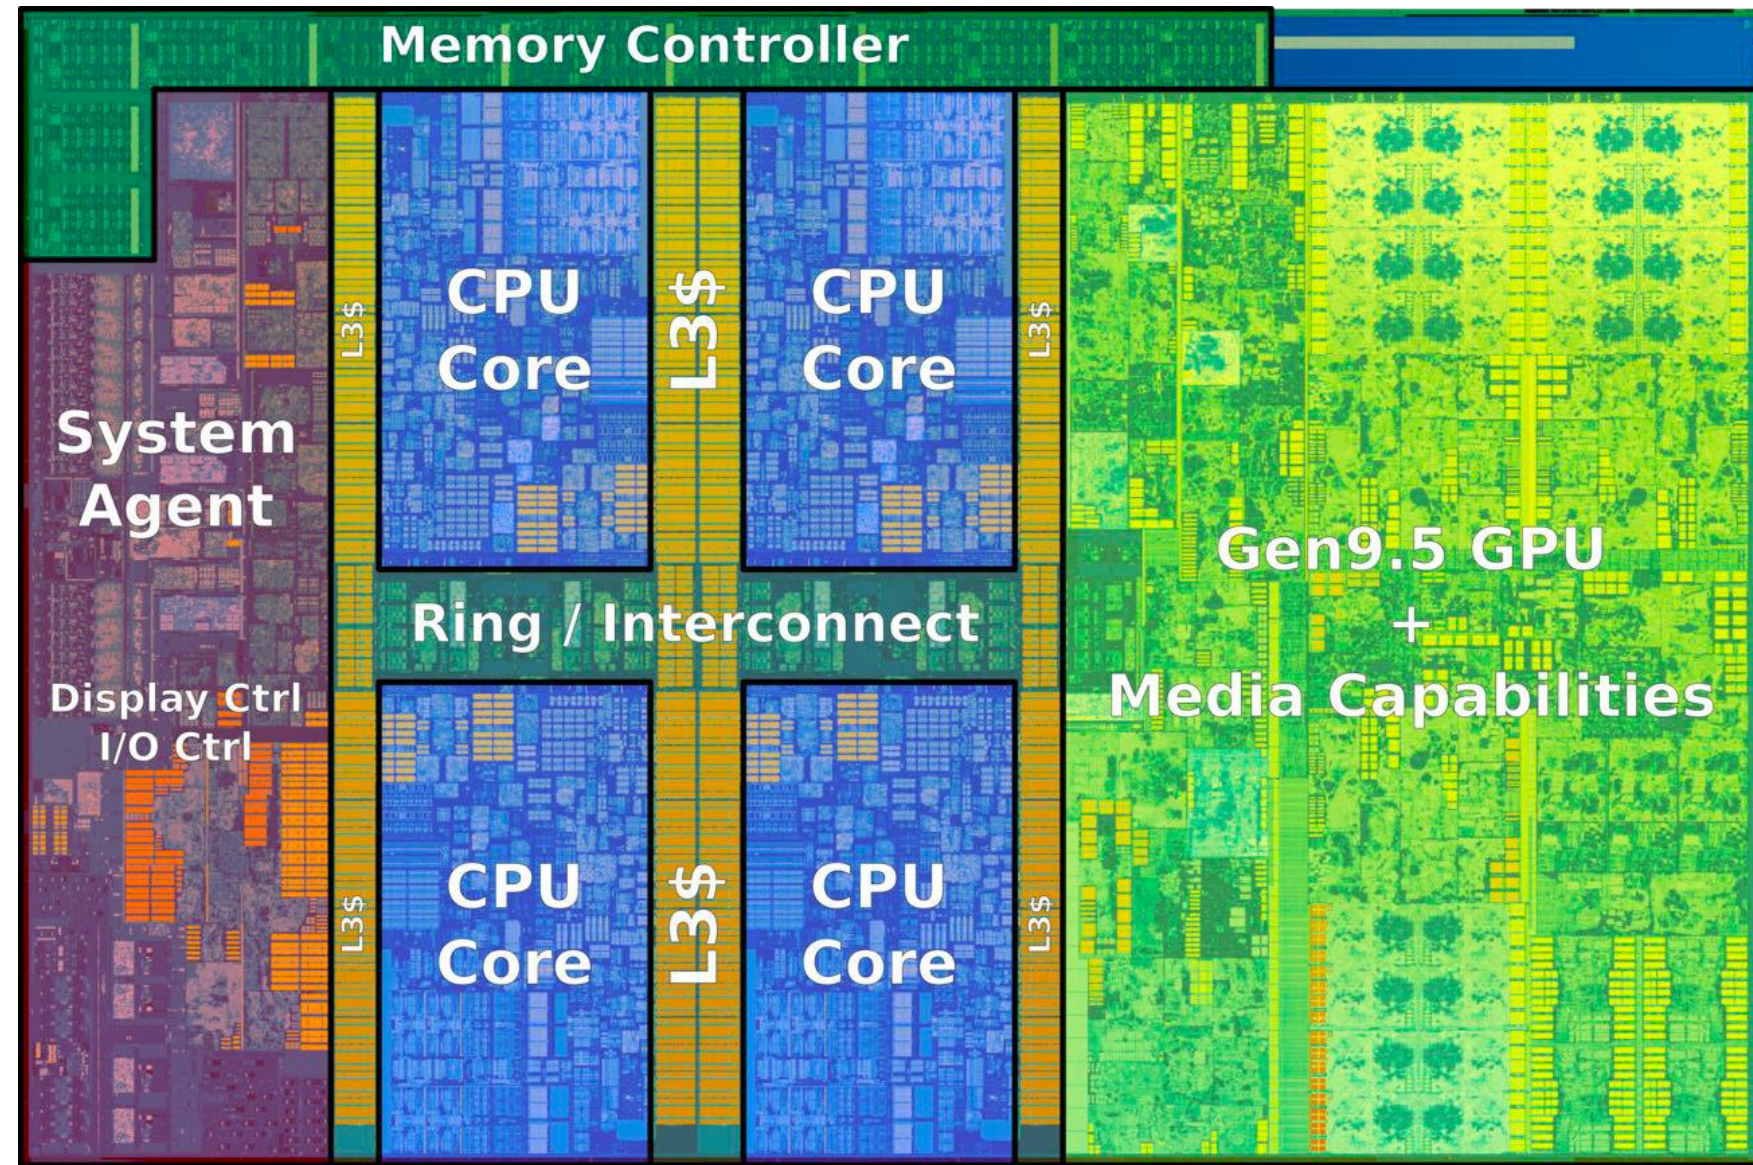


\* Caches (not shown) are another implementation of a shared address space




# Shared address space hardware architecture

Any processor can directly reference any memory location



Example: Intel Core i7 processor (Kaby Lake)




Intel Core i7 (quad core)  
(interconnect is a ring)



# Intel's ring interconnect

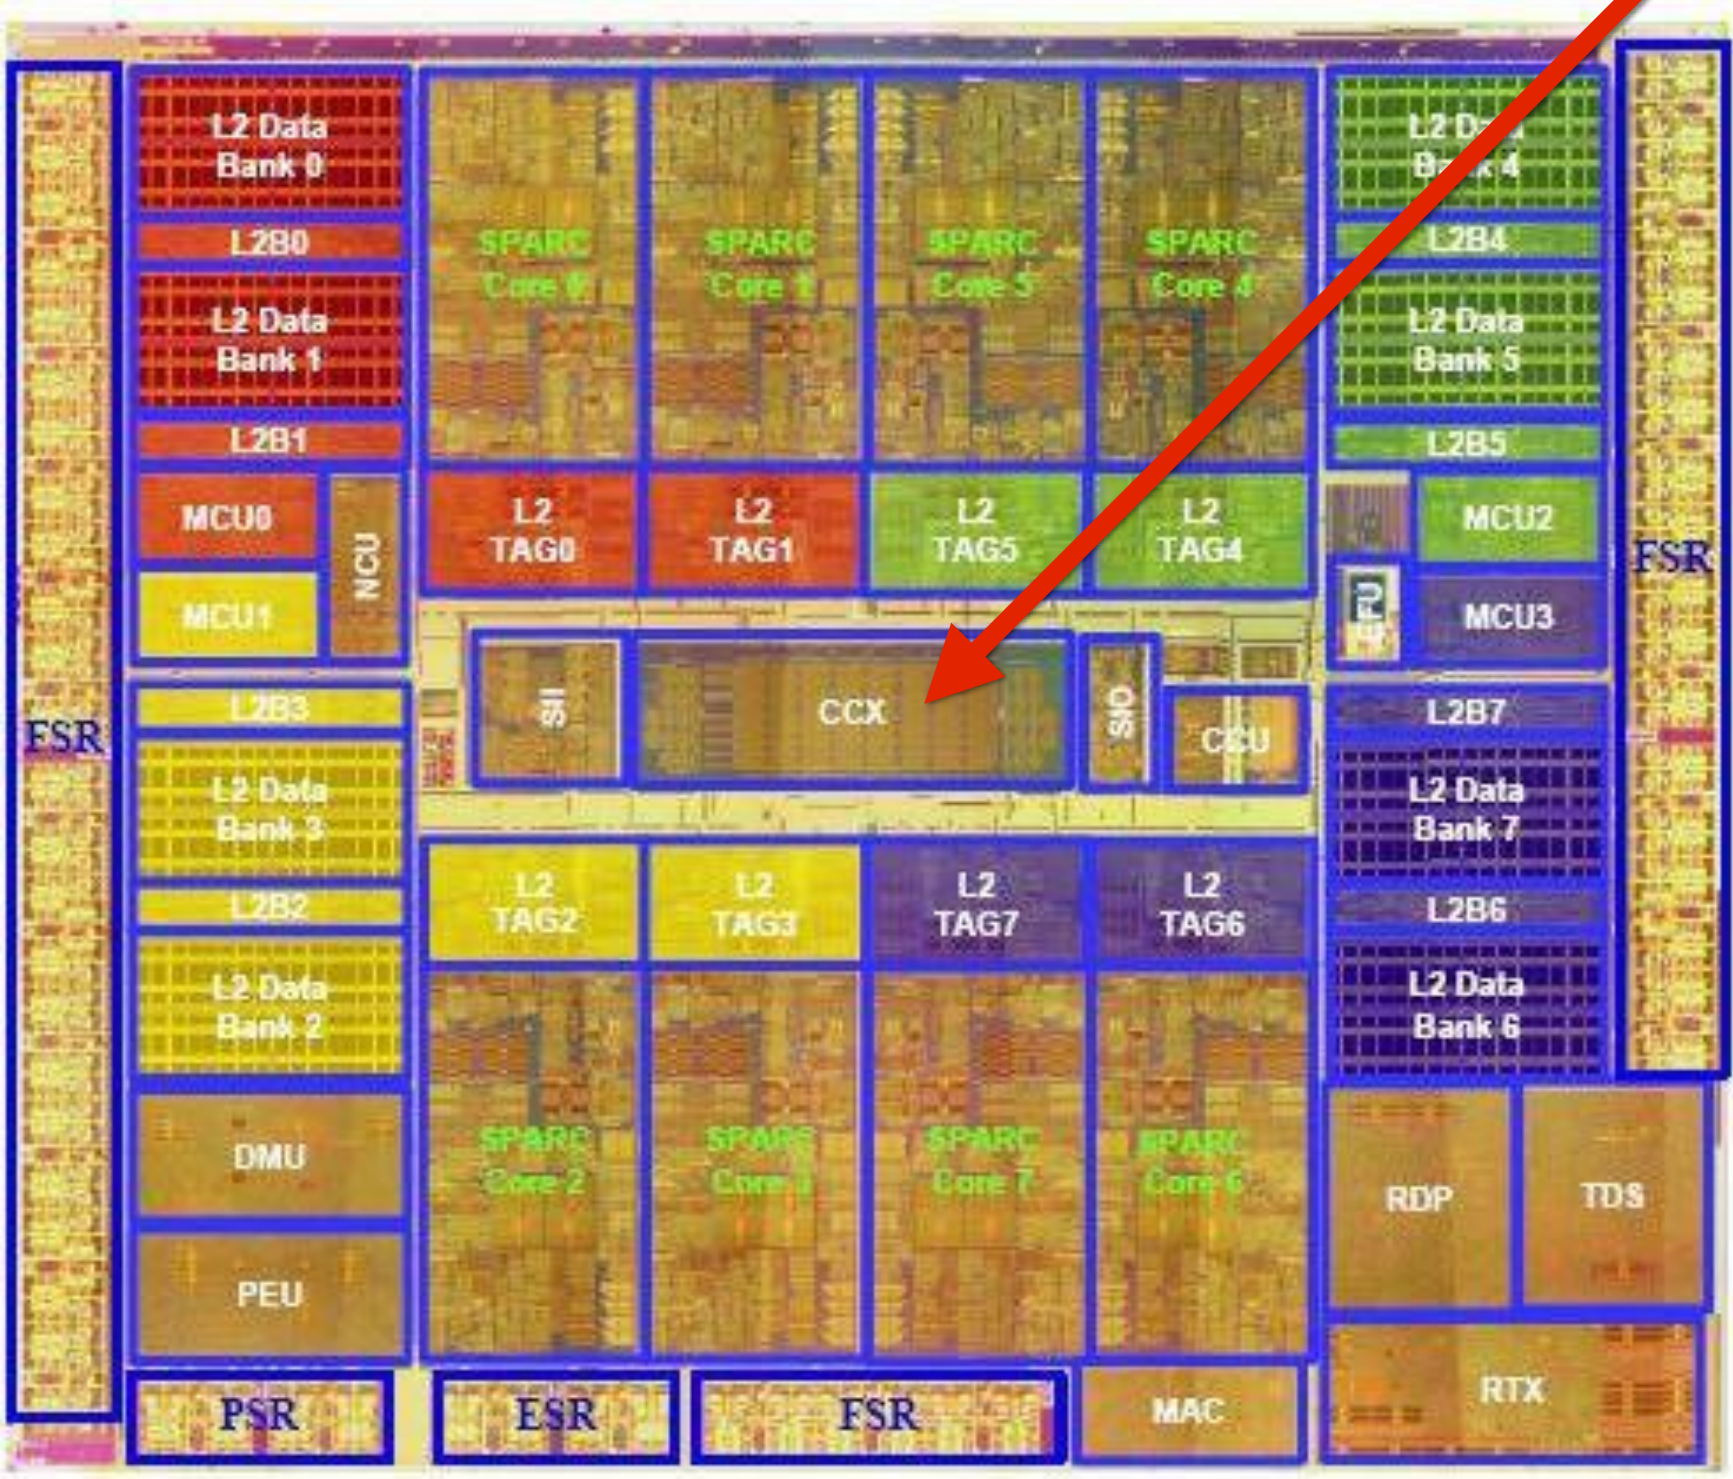
Introduced in Sandy Bridge microarchitecture



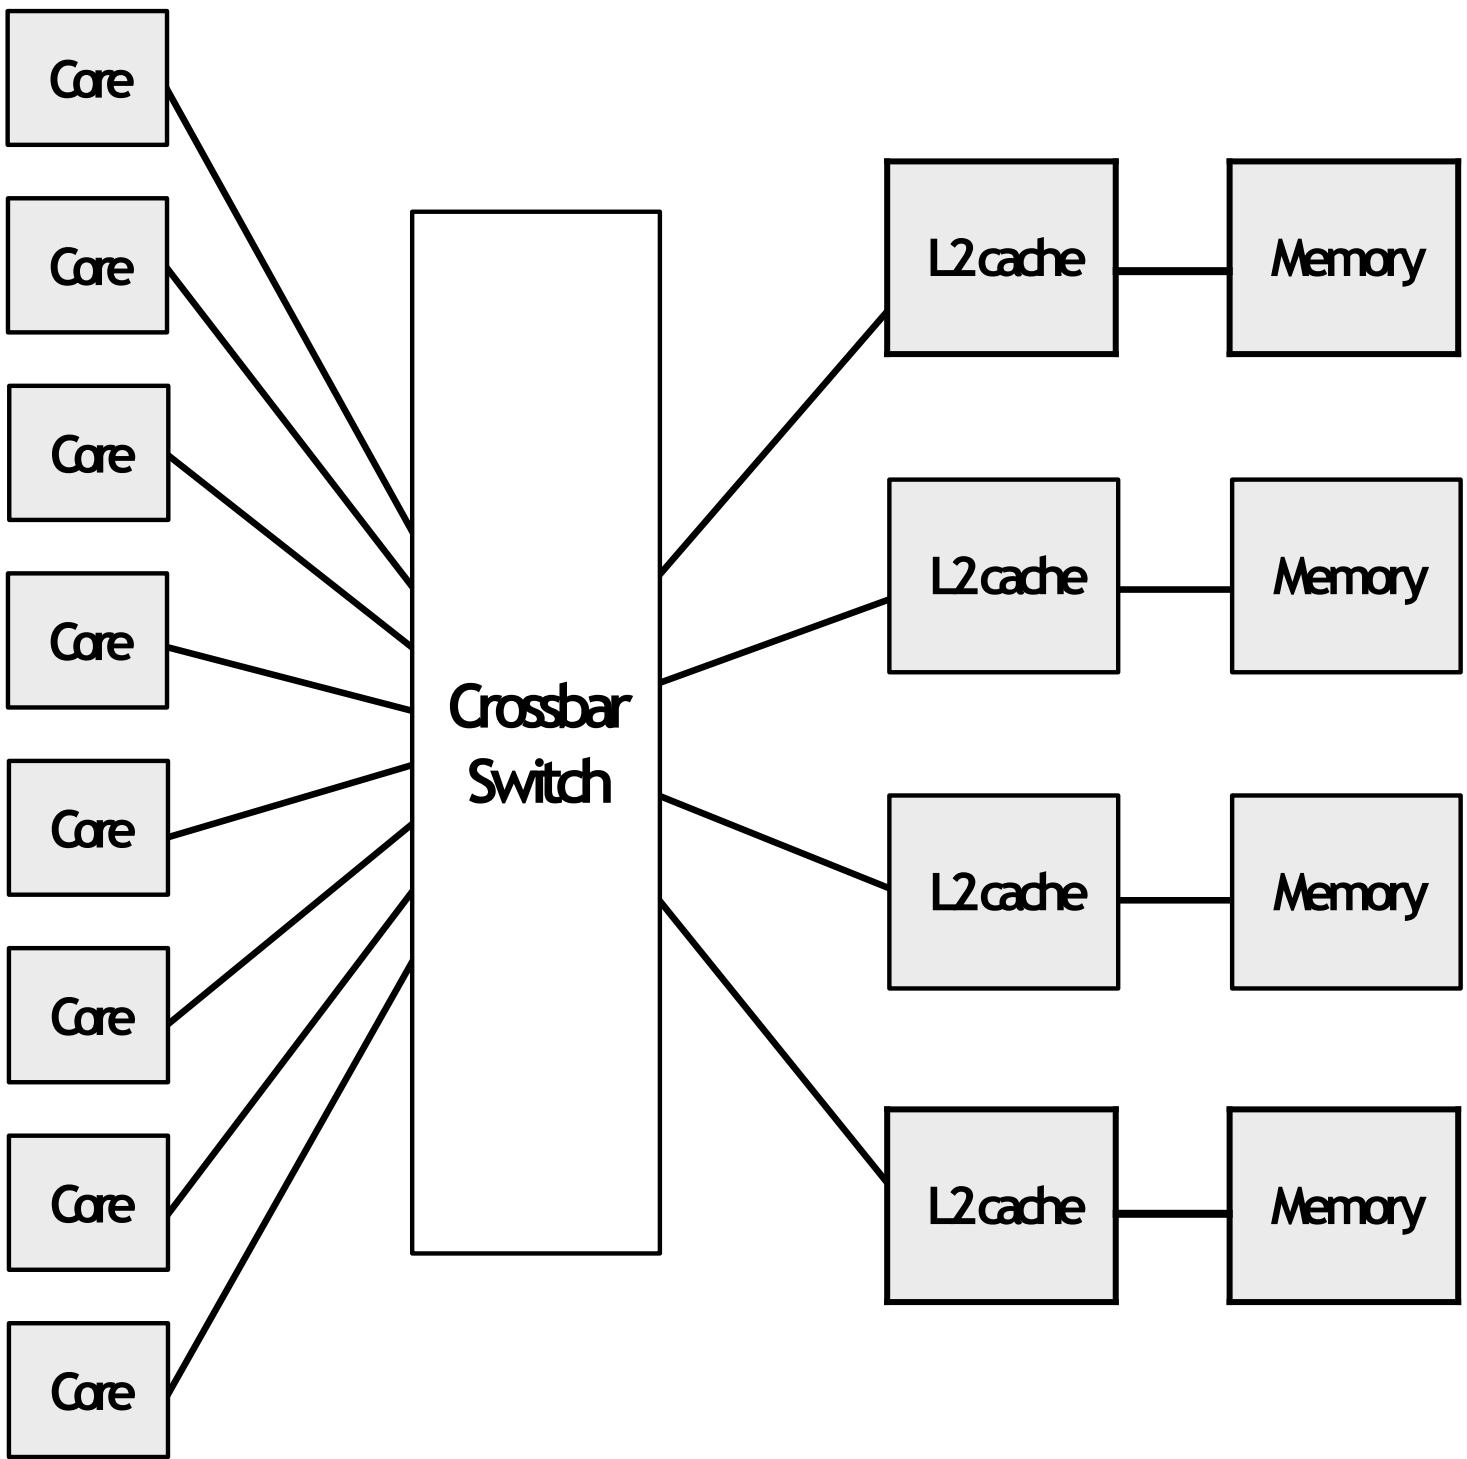
- Four rings
  - request
  - snoop
  - ack
  - data (32 bytes)
- Six interconnect nodes: four “slices” of L3 cache + system agent + graphics
- Each bank of L3 connected to ring bus twice
- Theoretical peak BW from cores to L3 at 3.4 GHz ~ 435 GB/sec
  - When each core is accessing its local slice

# SUN Niagara 2 (UltraSPARC T2): crossbar interconnect

Note area of crossbar (CCX):  
about same area as one core on chip




Eight core processor






# Intel Xeon Phi (Knights Landing)



- 72 cores, arranged as 6x6 mesh of tiles (2 cores/tile) YX routing of messages:
- - Message travels in Y direction
  - “Turn”
  - Message travels in X direction



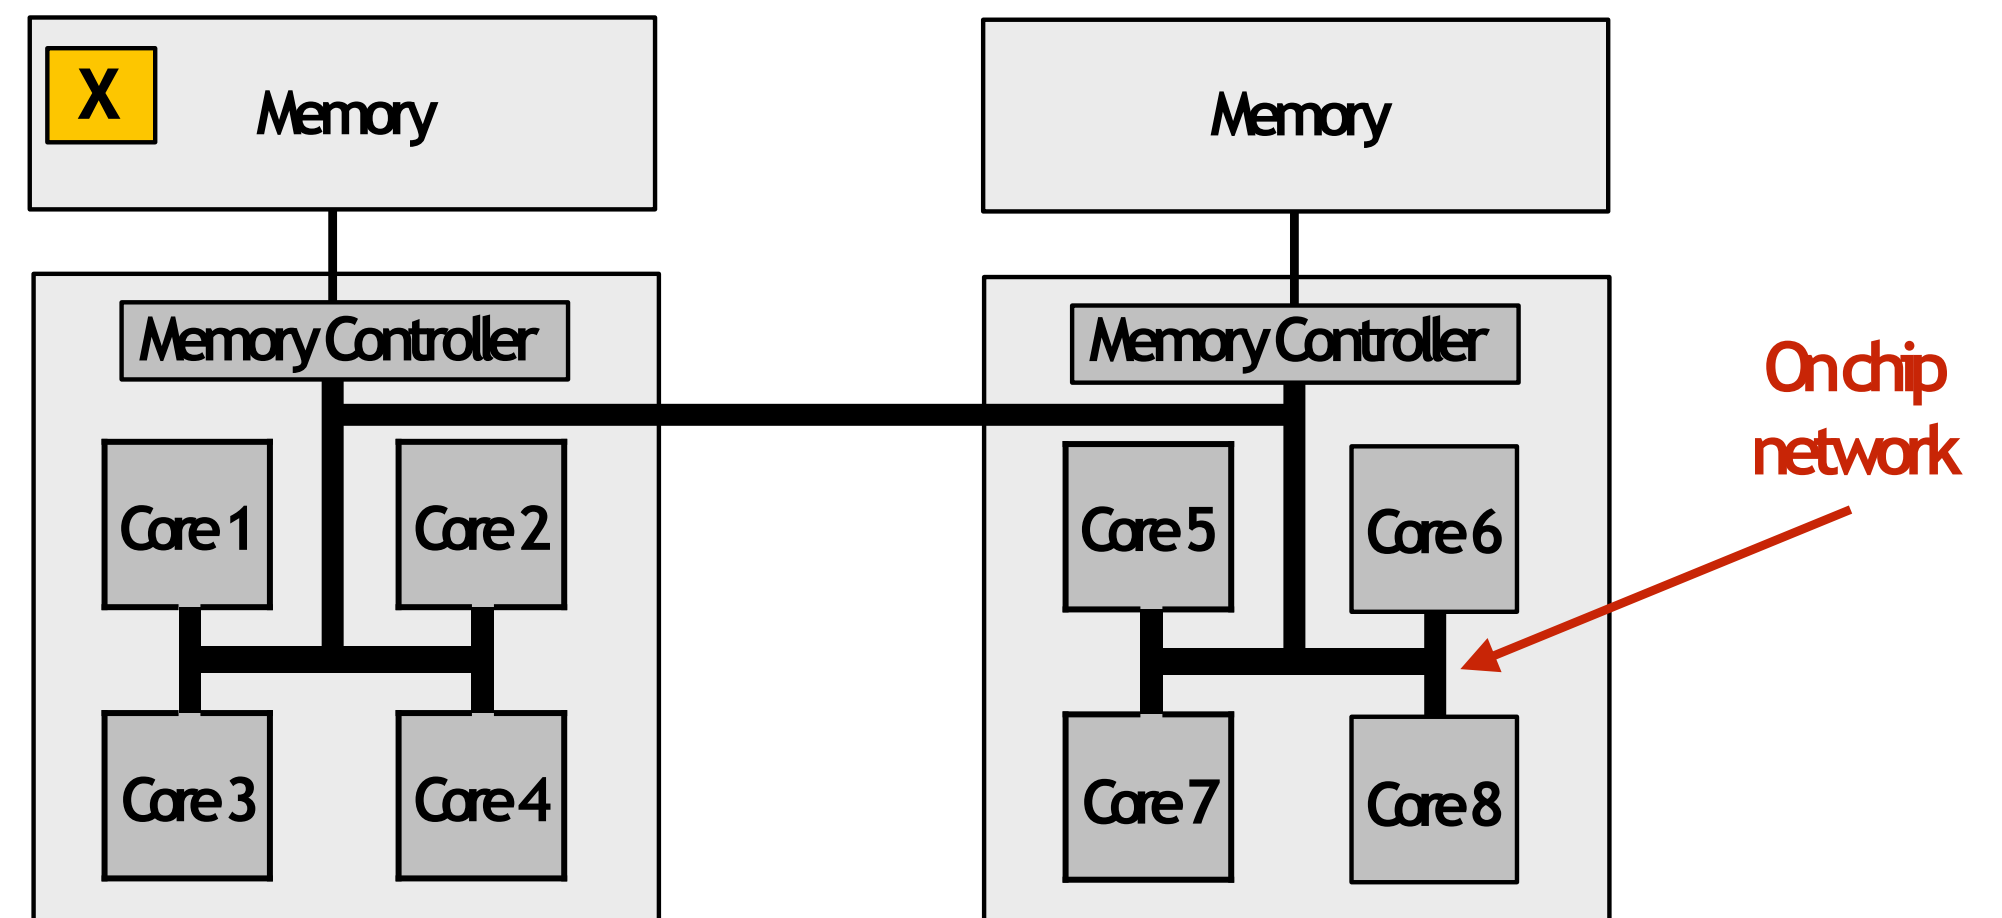


# Non-uniform memory access (NUMA)

The latency of accessing a memory location may be different from different processing cores in the system Bandwidth from any one location may also be different to different CPU cores \*



Example: modern multi-socket configuration



\* In practice, you'll find NUMA behavior on a single-socket system as well (recall: different cache slices are a different distance from each core)

# Summary: shared address space model

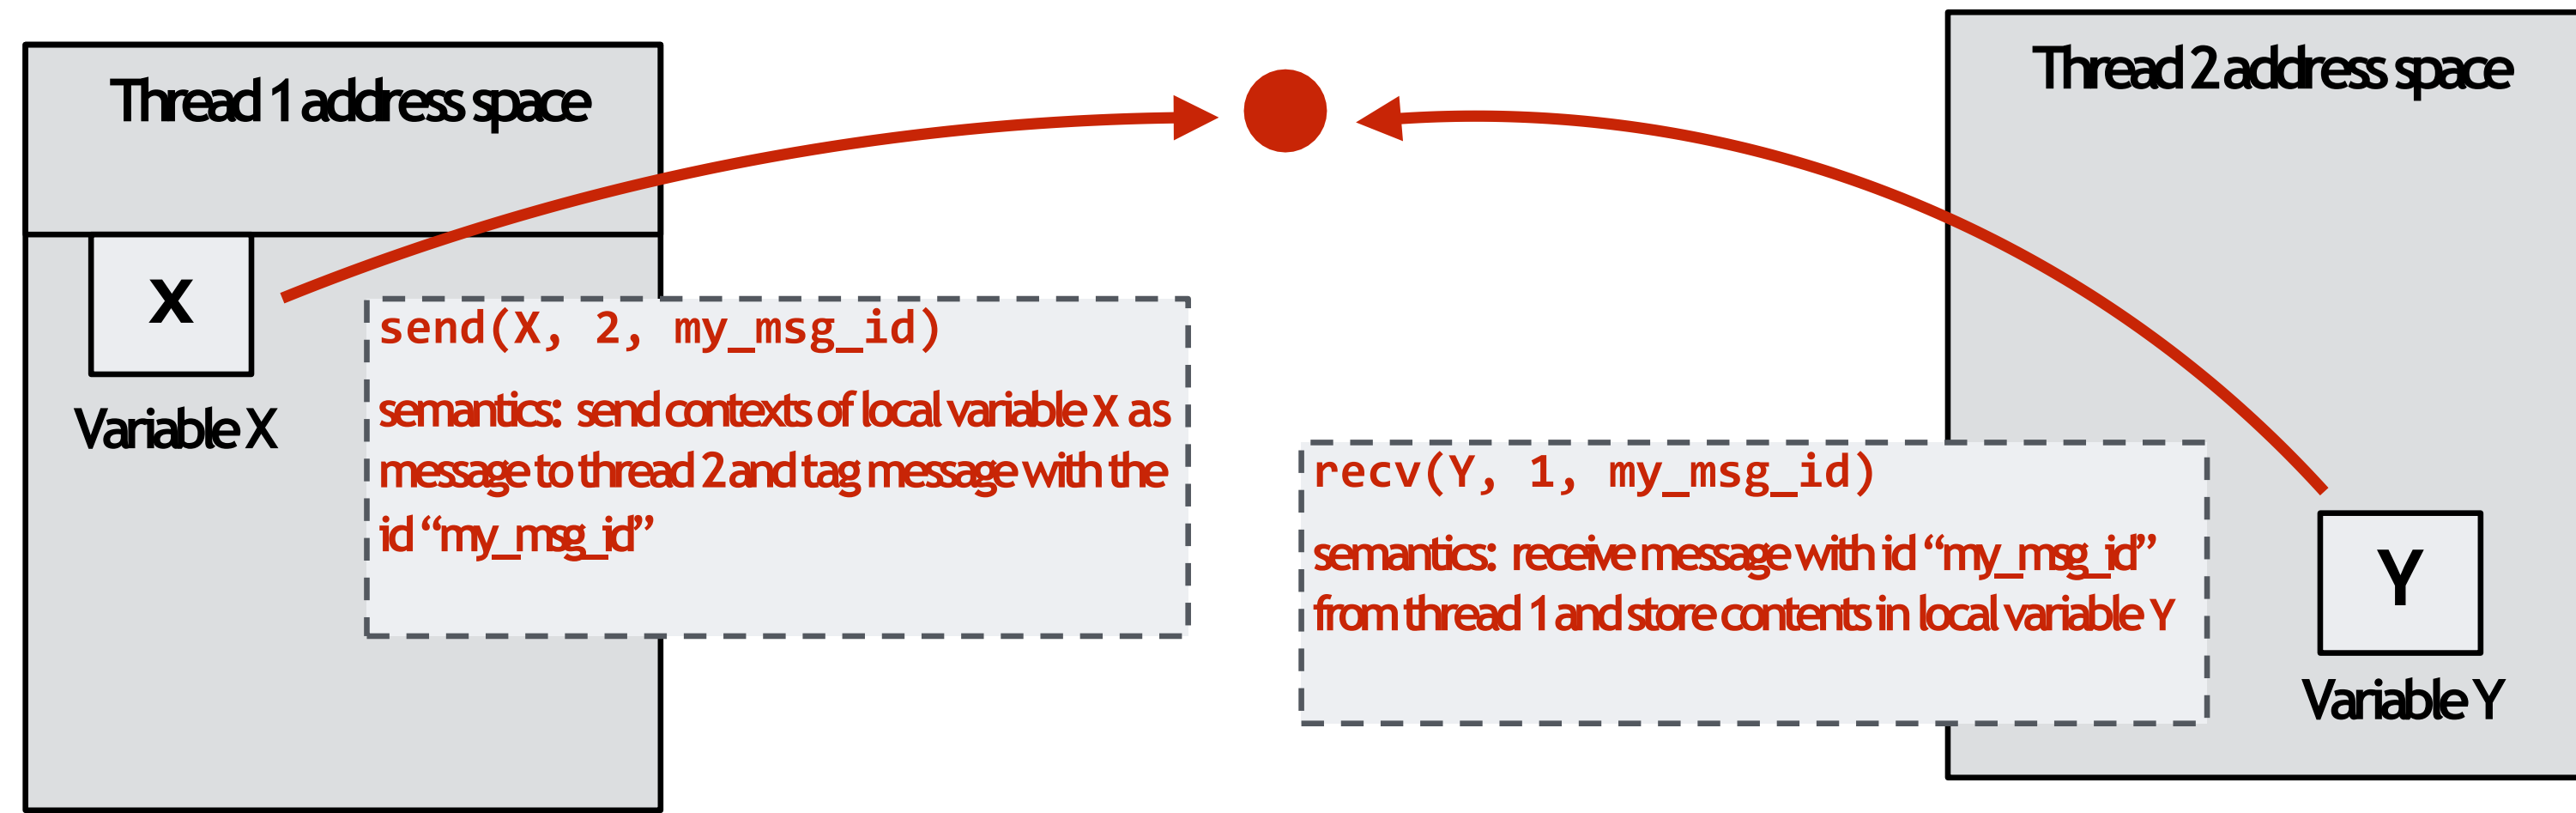
- **Communication abstraction**
  - Threads read/write variables in shared address space
  - Threads manipulate synchronization primitives: locks, atomic ops, etc.
  - Logical extension of uniprocessor programming \*
- **Requires hardware support to implement efficiently**
  - Any processor can load and store from any address
  - Can be costly to scale to large numbers of processors  
(one of the reasons why high-core count processors are expensive)

\* But NUMA implementations requires reasoning about locality for performance optimization

# **Message passing model of communication**

# Message passing model (abstraction)

- Threads operate within their own private address spaces
- Threads communicate by sending/receiving messages
  - send: specifies recipient, buffer to be transmitted, and optional message identifier (“tag”)
  - receive: sender, specifies buffer to store data, and optional message identifier
  - Sending messages is the only way to exchange data between threads 1 and 2
    - Why?



(Communication operations shown in red)




# A common metaphor: snail mail





# Message passing (implementation)

- Hardware need not implement system-wide loads and stores to execute message passing programs (it need only communicate messages between nodes)
  - Can connect commodity systems together to form a large parallel machine (message passing is a programming model for clusters and supercomputers)



Cluster of workstations  
(Infiniband network)



# The data-parallel model

# Programming models provide a way to think about the organization of parallel programs (by imposing structure)

- **Shared address space: very little structure to communication**
  - All threads can read and write to all shared variables
- **Message passing: communication is structured in the form of messages**
  - All communication occurs in the form of messages
  - Communication is explicit in source code—the sends and receives)
- **Data parallel structure: more rigid structure to computation**
  - Perform same function on elements of large collections

# Data-parallel model \*

- Organize computation as operations on sequences of elements
  - e.g., perform same function on all elements of a sequence
- A well-known modern example: NumPy:  $C = A + B$  (A, B, and C are vectors of same length)

Something you've seen early in the lecture...

\* We'll have multiple lectures in the course about data-parallel programming and data-parallel thinking: this is just a taste


# Key data type: sequences

- Ordered collection of elements
- For example, in a C++ like language: `Sequence<T>`
- Scala lists: `List[T]`
- In a functional language (like Haskell): `seq T`
- Program can only access elements of sequence through sequence operators:
  - map, reduce, scan, shift, etc.

# Map

- Higher order function (function that takes a function as an argument) that operates on sequences
- Applies side-effect-free unary function  $f :: a \rightarrow b$  to all elements of input sequence, to produce output sequence of the same length
- In a functional language (e.g., Haskell)
  - `map :: (a -> b) -> seq a -> seq b`
- In C++:

```
template<class InputIt, class OutputIt, class UnaryOperation> OutputIt
transform(InputIt first1, InputIt last1,
          OutputIt d_first, UnaryOperation
          unary_op);
```



# Parallelizing map

- Since  $f :: a \rightarrow b$  is a function (side-effect free), then applying  $f$  to all elements of the sequence can be done **in any order** without changing the output of the program
- The implementation of map has flexibility to reorder/parallelize processing of elements of sequence however it sees fit

# Data parallelism in ISPC

```
// main C++ code:
const int N = 1024;
float* x = new float[N];
float* y = new float[N];

// initialize N elements of x here

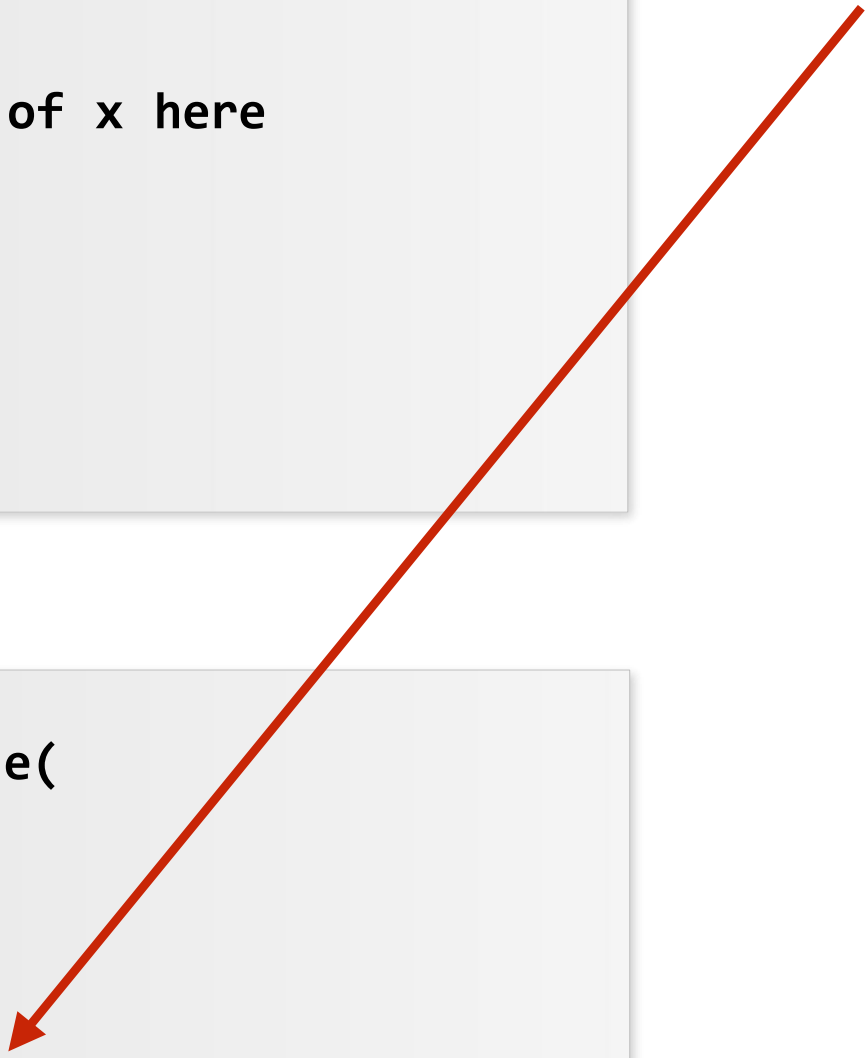
absolute_value(N, x, y);
```

foreach construct

Think of loop body as a function

Given this program, it is reasonable to think of the program as using foreach to “map the loop body onto each element” of the arrays X and Y.

```
// ISPC code:
export void absolute_value(
    uniform int N,
    uniform float* x,
    uniform float* y)
{
    foreach (i = 0 ... N)
    {
        if (x[i] < 0)
            y[i] = -x[i];
        else
            y[i] = x[i];
    }
}
```



But if we want to be more precise: a sequence is not a first-class ISPC concept. It is implicitly defined by how the program has implemented array indexing logic in the foreach loop.

(There is no operation in ISPC with the semantic: “map this code over all elements of this sequence”)



# Data parallelism in ISPC

```
// main C++ code:
const int N = 1024;
float* x = new float[N/2];
float* y = new float[N];

// initialize N/2 elements of x here

absolute_repeat(N/2, x, y);
```

Think of loop body as a function

The input/output sequences being mapped over are implicitly defined by array indexing logic

```
// ISPC code:
export void absolute_repeat(
    uniform int N,
    uniform float* x,
    uniform float* y)
{
    foreach (i = 0 ... N)
    {
        if (x[i] < 0)
            y[2*i] = -x[i];
        else
            y[2*i] = x[i];
        y[2*i+1] = y[2*i];
    }
}
```

This is also a valid ISPC program!

It takes the absolute value of elements of x, then repeats it twice in the output array y

(Less obvious how to think of this code as mapping the loop body onto existing sequences.)

# Data parallelism in ISPC

```
// main C++ code:  
const int N = 1024;  
float* x = new float[N];  
float* y = new float[N];  
  
// initialize N elements of x  
  
shift_negative(N, x, y);
```

Think of loop body as a function

The input/output sequences being mapped over are implicitly defined by array indexing logic

```
// ISPC code:  
export void shift_negative(  
    uniform int N,  
    uniform float* x,  
    uniform float* y)  
{  
    foreach (i = 0 ... N)  
    {  
        if (i >= 1 && x[i] < 0)  
            y[i-1] = x[i];  
        else  
            y[i] = x[i];  
    }  
}
```

The output of this program is undefined!

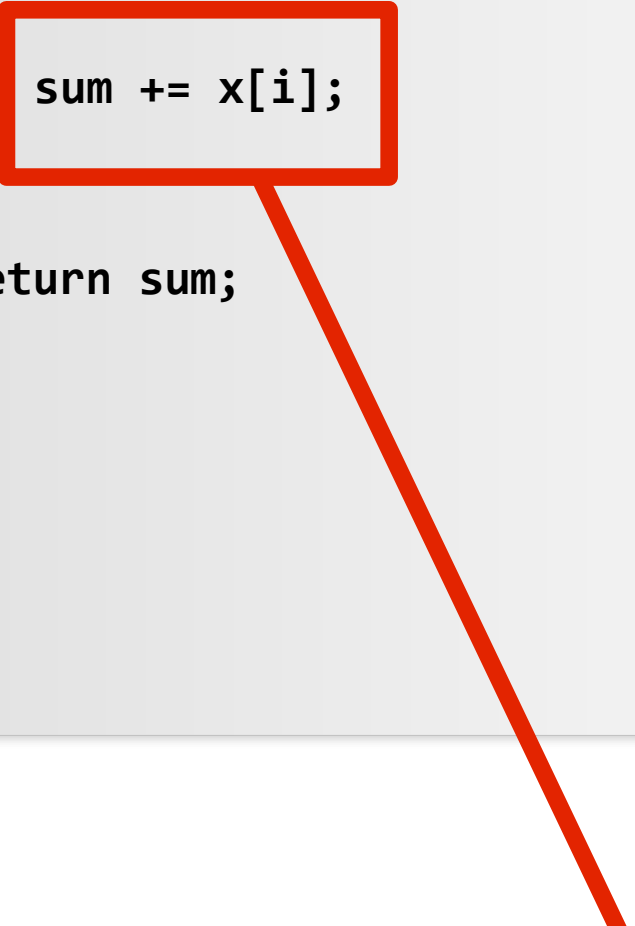
Possible for multiple iterations of the loop body to write to same memory location

Data-parallel model (foreach) provides no specification of order in which iterations occur

# ISPC discussion: sum “reduction”

Compute the sum of all array elements in parallel

```
export uniform float sumall11(  
    uniform int N,  
    uniform float* x)  
{  
    uniform float sum = 0.0f;  
    foreach (i = 0 ... N)  
    {  
        sum += x[i];  
    }  
    return sum;  
}
```



sum is of type uniform float (one copy of variable for all program instances) x[i] is not a uniform expression (different value for each program instance) Result: compile-time type error

```
export uniform float sumall12(  
    uniform int N,  
    uniform float* x)  
{  
    uniform float sum;  
    float partial = 0.0f;  
    foreach (i = 0 ... N)  
    {  
        partial += x[i];  
    }  
  
    // from ISPC math library  
    sum = reduce_add(partial);  
  
    return sum;  
}
```

Correct ISPC solution

# ISPC discussion: sum “reduction”

Each instance accumulates a private partial sum  
(no communication)

Partial sums are added together using the `reduce_add()` cross-instance communication primitive. The result is the same total sum for all program instances (`reduce_add()` returns a uniform float)

The ISPC code at right will execute in a manner similar to handwritten C + AVX intrinsics implementation below. \*

```
float sumall2(int N, float* x) {  
  
    float tmp[8]; // assume 16-byte alignment  
    __mm256 partial = _mm256_broadcast_ss(0.0f);  
  
    for (int i=0; i<N; i+=8)  
        partial = _mm256_add_ps(partial, _mm256_load_ps(&x[i]));  
  
    _mm256_store_ps(tmp, partial);  
  
    float sum = 0.f;  
    for (int i=0; i<8; i++)  
        sum += tmp[i];  
  
    return sum;  
}
```

```
export uniform float sumall2(  
    uniform int N,  
    uniform float* x)  
{  
    uniform float sum;  
    float partial = 0.0f;  
    foreach (i = 0 ... N)  
    {  
        partial += x[i];  
    }  
  
    // from ISPC math library  
    sum = reduce_add(partial);  
  
    return sum;  
}
```

**\* Self-test:** If you understand why this implementation complies with the semantics of the ISPC gang abstraction, then you’ve got a good command of ISPC

# Summary: data-parallel model

- Data-parallelism is about imposing rigid program structure to facilitate simple programming and advanced optimizations
- Basic structure: map a function onto a large collection of data
  - Functional: side-effect free execution
  - No communication among distinct function invocations  
(allow invocations to be scheduled in any order, including in parallel)
- Other data parallel operators express more complex patterns on sequences: gather, scatter, reduce, scan, shift, etc.
  - This will be a topic of a later lecture
- You will think in terms of data-parallel primitives often in this class, but many modern performance-oriented data-parallel languages do not enforce this structure in the language
  - Many languages (like ISPC, CUDA, etc.) choose flexibility/familiarity of imperative C-style syntax over the safety of a more functional form

# Summary

# Summary

- Programming models provide a way to think about the organization of parallel programs.
- They provide abstractions that permit multiple valid implementations.
- *I want you to always be thinking about abstraction vs. implementation for the remainder of this course.*