



# 058165 - PARALLEL COMPUTING

Fabrizio Ferrandi

a.a. 2022-2023

- ❑ “Structured Parallel Programming: Patterns for Efficient Computation,” Michael McCool, Arch Robinson, James Reinders, 1<sup>st</sup> edition, Morgan Kaufmann, ISBN: 978-0-12-415993-8, 2012

- ❑ What are Collectives?
- ❑ Reduce Pattern
- ❑ Scan Pattern
- ❑ Sorting


- ❑ Collective operations deal with a *collection* of data as a whole, rather than as separate elements
- ❑ Collective patterns include:
  - ▶ Reduce
  - ▶ Scan
  - ▶ Partition
  - ▶ Scatter
  - ▶ Gather

- ❑ Collective operations deal with a *collection* of data as a whole, rather than as separate elements
- ❑ Collective patterns include:
  - ▶ Reduce
  - ▶ Scan
  - ▶ Partition
  - ▶ Scatter
  - ▶ Gather


Reduce and Scan  
will be covered in  
this lecture

- ❑ **Reduce** is used to combine a collection of elements into one summary value
- ❑ A combiner function combines elements pairwise
- ❑ A combiner function only needs to be *associative* to be parallelizable
- ❑ Example combiner functions:
  - ▶ Addition
  - ▶ Multiplication
  - ▶ Maximum / Minimum

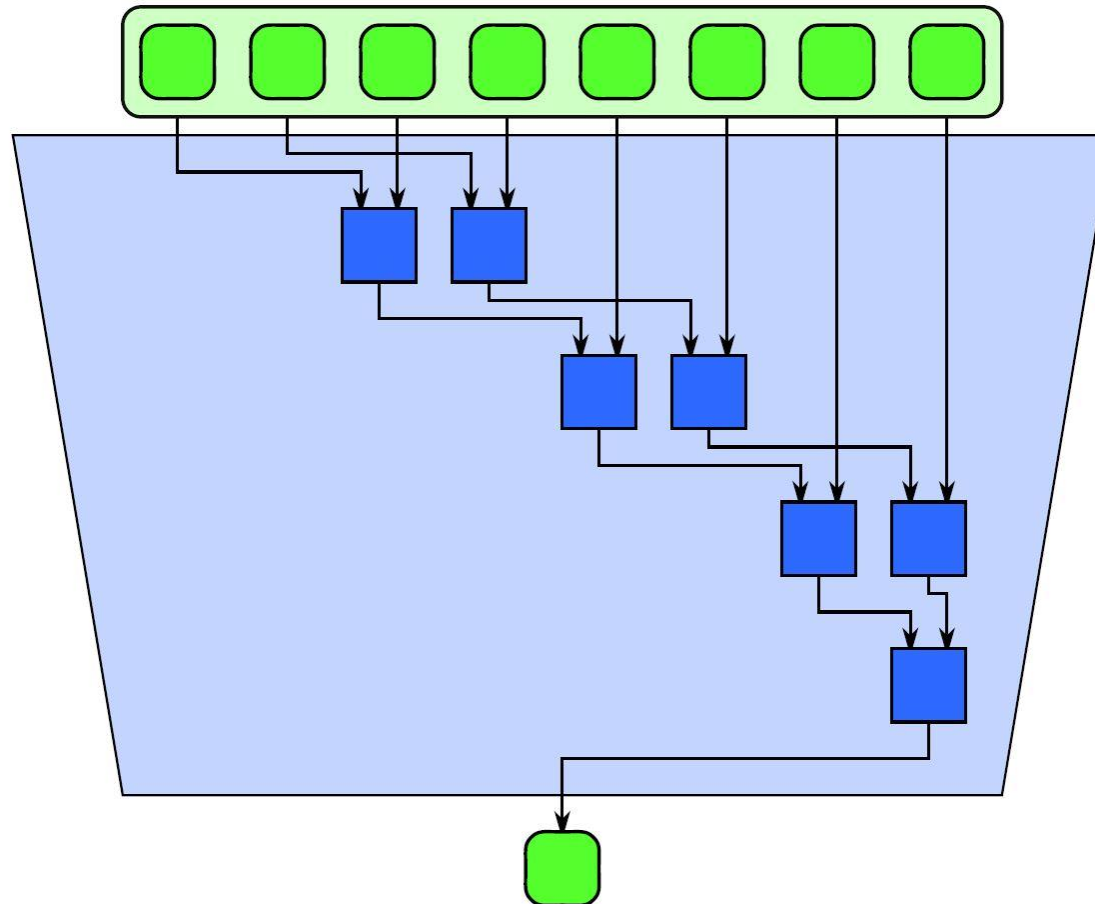
## Serial Reduction



## Parallel Reduction

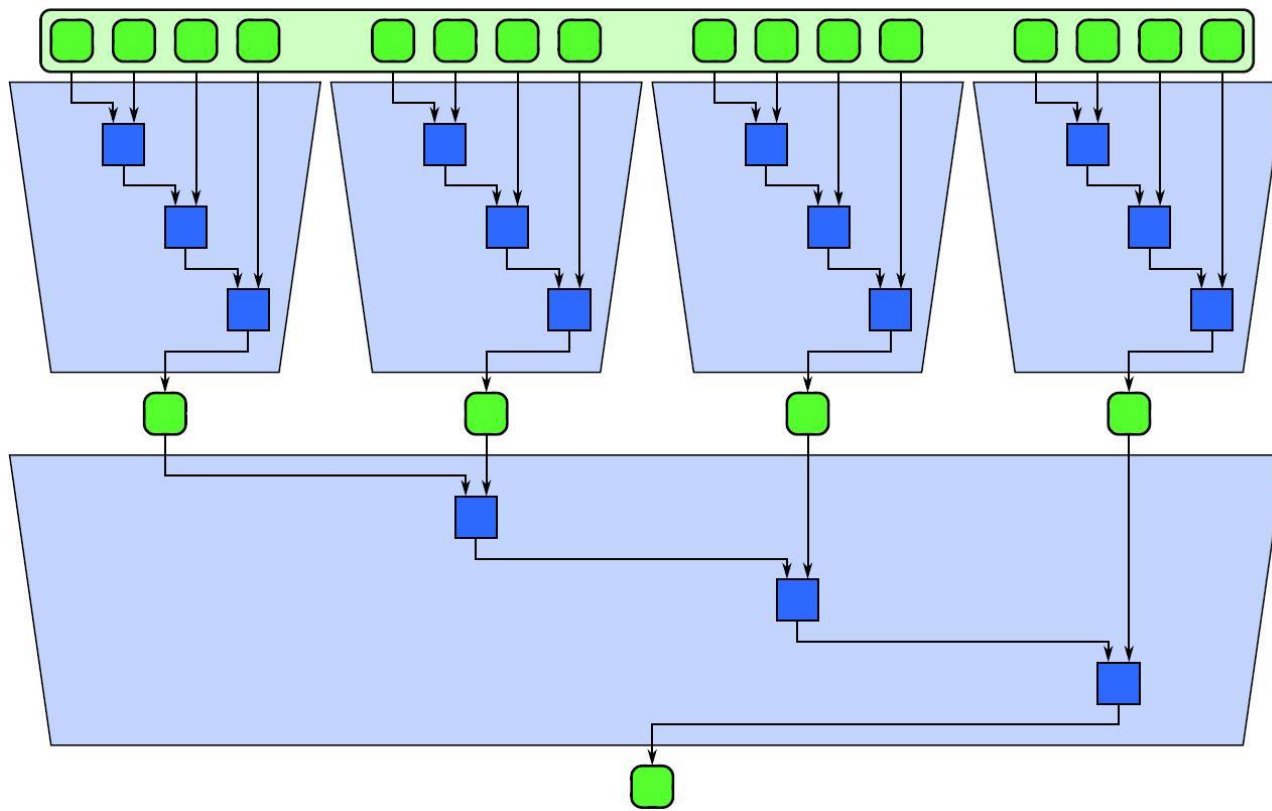


## ❑ Vectorization






- ❑ **Tiling** is used to break chunks of work up for workers to reduce serially



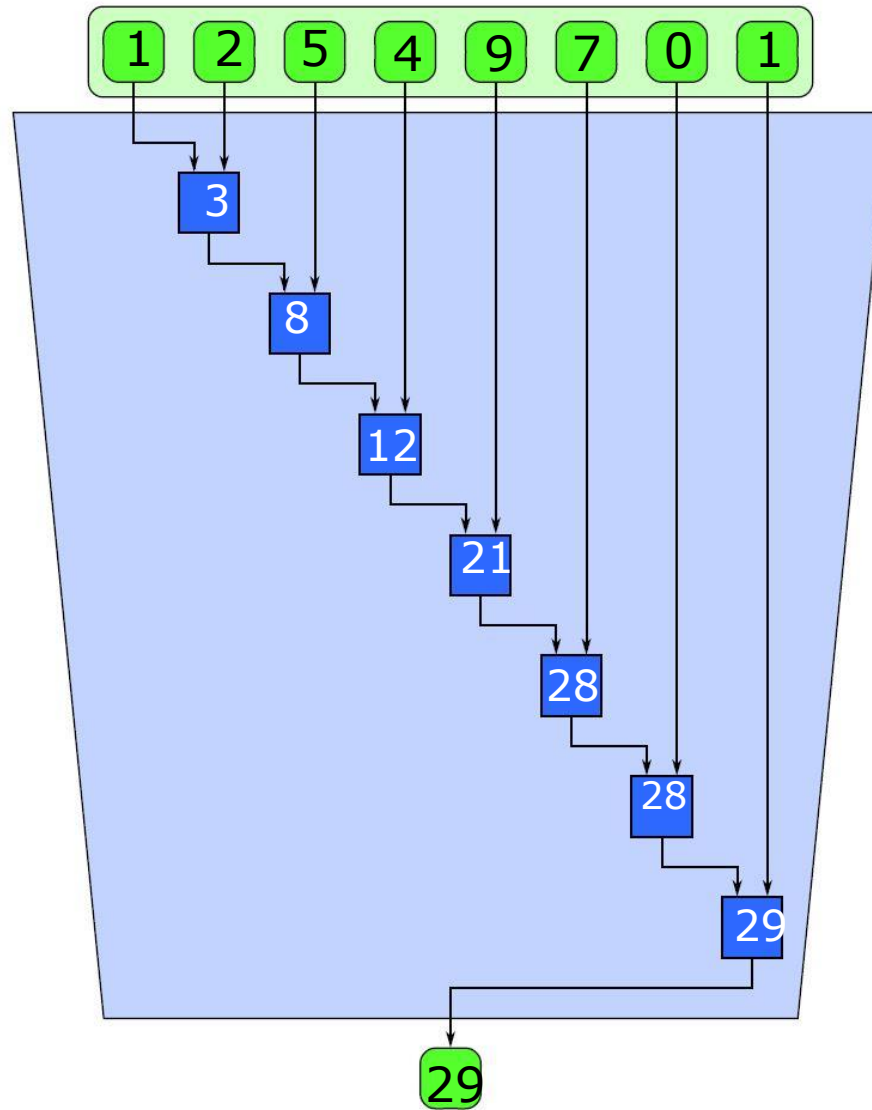
# Reduce – Add Example

10



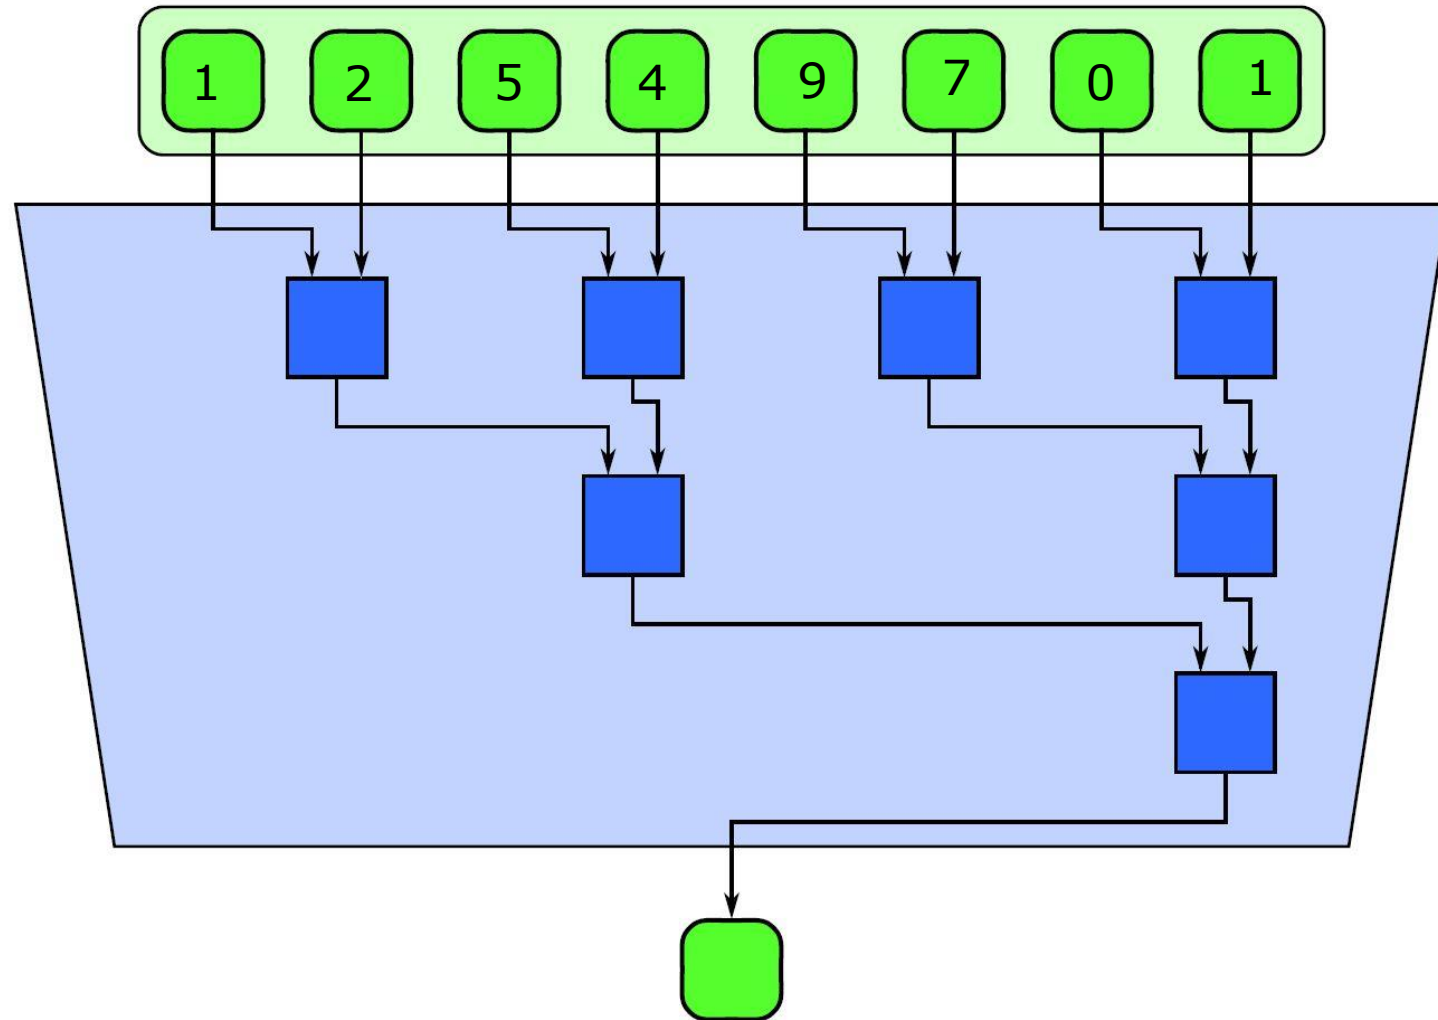
# Reduce – Add Example

11



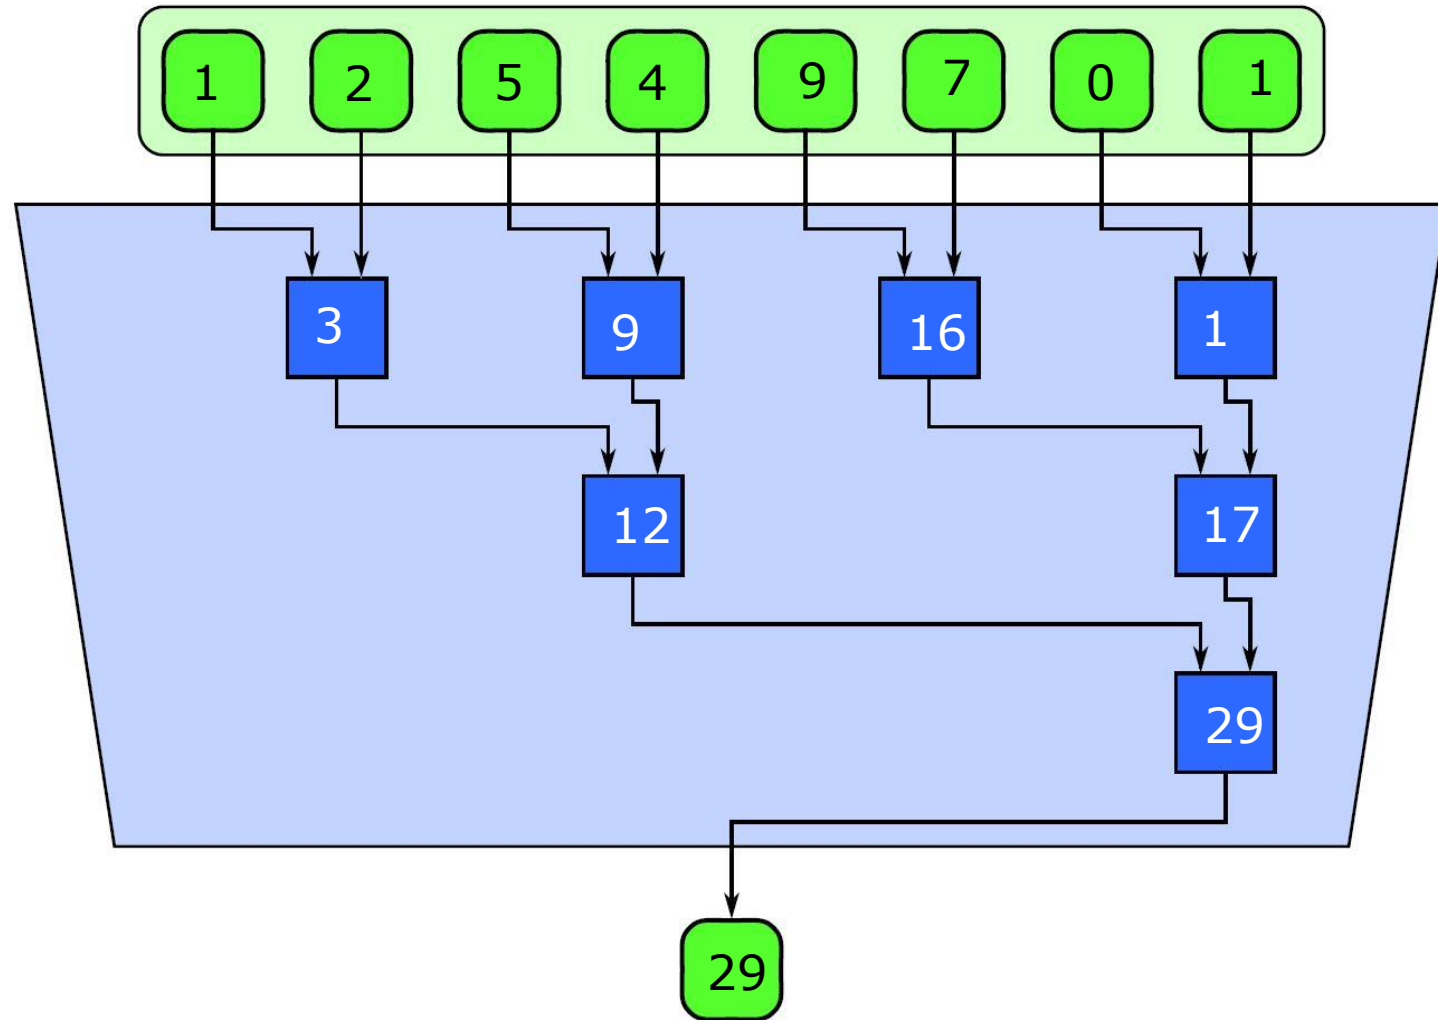
# Reduce – Add Example

12




# Reduce – Add Example

13



- We can “fuse” the map and reduce patterns



- ❑ Precision can become a problem with reductions on floating point data
- ❑ Different orderings of floating-point data can change the reduction value

- ❑ 2 vectors of same length
- ❑ Map (\*) to multiply the components
- ❑ Then reduce with (+) to get the final answer

$$\mathbf{a} \cdot \mathbf{b} = \sum_{i=0}^{n-1} a_i b_i.$$

Also:  $\vec{a} \cdot \vec{b} = |\vec{a}| \cos(\theta) |\vec{b}|$



# Dot Product – Example Uses

17

- ❑ Essential operation in physics, graphics, video games,...
- ❑ Gaming analogy: in Mario Kart, there are “boost pads” on the ground that increase your speed
  - ▶ red vector is your speed (x and y direction)
  - ▶ blue vector is the orientation of the boost pad (x and y direction). Larger numbers are more power.

How much boost will you get? For the analogy, imagine the pad multiplies your speed:

- If you come in going 0, you'll get nothing
- If you cross the pad perpendicularly, you'll get 0 [just like the banana obliteration, it will give you 0x boost in the perpendicular direction]

$$\text{Total} = \text{speed}_x \cdot \text{boost}_x + \text{speed}_y \cdot \text{boost}_y$$




Ref: <http://betterexplained.com/articles/vector-calculus-understanding-the-dot-product/>


- ❑ The **scan** pattern produces partial reductions of input sequence, generates new sequence
- ❑ Trickier to parallelize than reduce
- ❑ Inclusive scan vs. exclusive scan
  - ▶ Inclusive scan: includes current element in partial reduction
  - ▶ Exclusive scan: excludes current element in partial reduction, partial reduction is of all prior elements prior to current element

- ❑ Lexical comparison of strings – e.g., determine that “strategy” should appear before “stratification” in a dictionary
- ❑ Add multi-precision numbers (those that cannot be represented in a single machine *word*)
- ❑ Evaluate polynomials
- ❑ Implement radix sort or quicksort
- ❑ Delete marked elements in an array
- ❑ Dynamically allocate processors
- ❑ Lexical analysis – parsing programs into tokens
- ❑ Searching for regular expressions
- ❑ Labeling components in 2-D images
- ❑ Some tree algorithms – e.g., finding the depth of every vertex in a tree

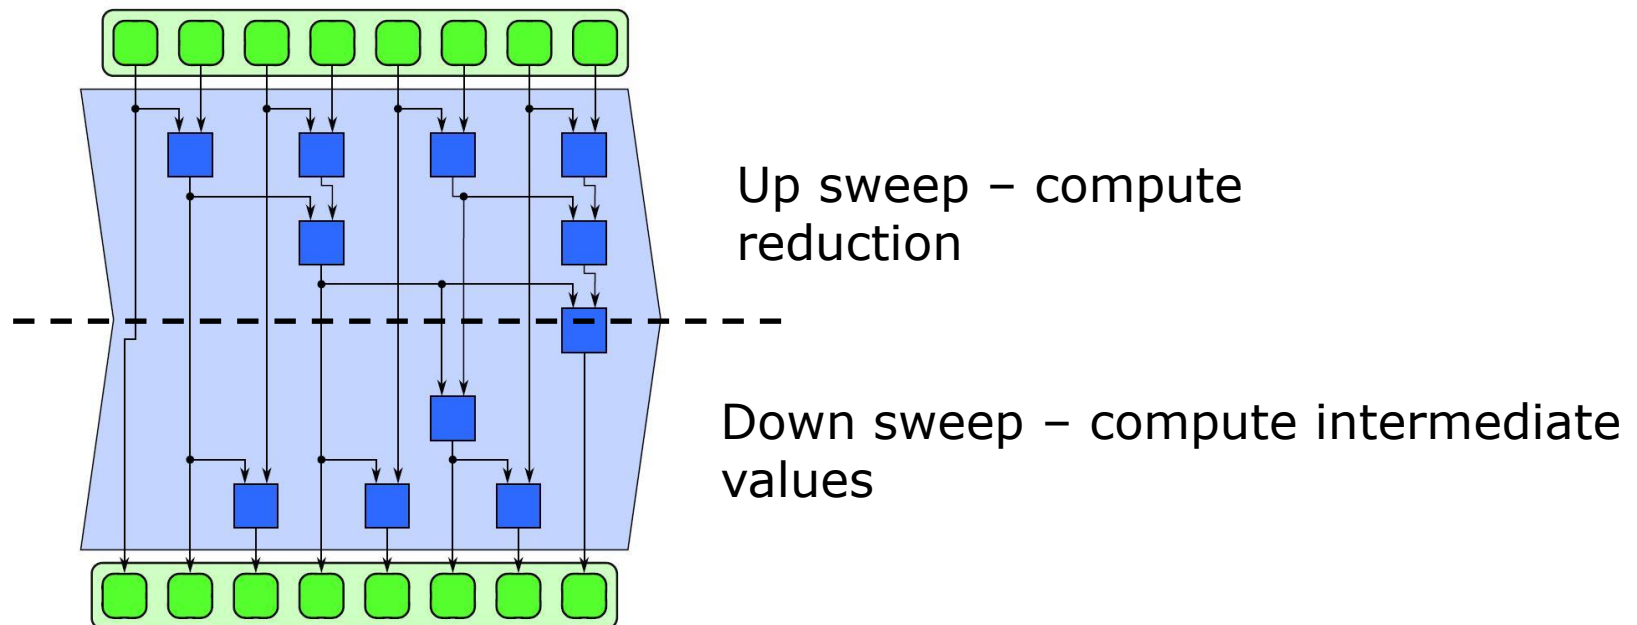
## Serial Scan



## Parallel Scan





- ❑ One algorithm for parallelizing scan is to perform an “up sweep” and a “down sweep”
- ❑ Reduce the input on the up sweep
- ❑ The down sweep produces the intermediate results



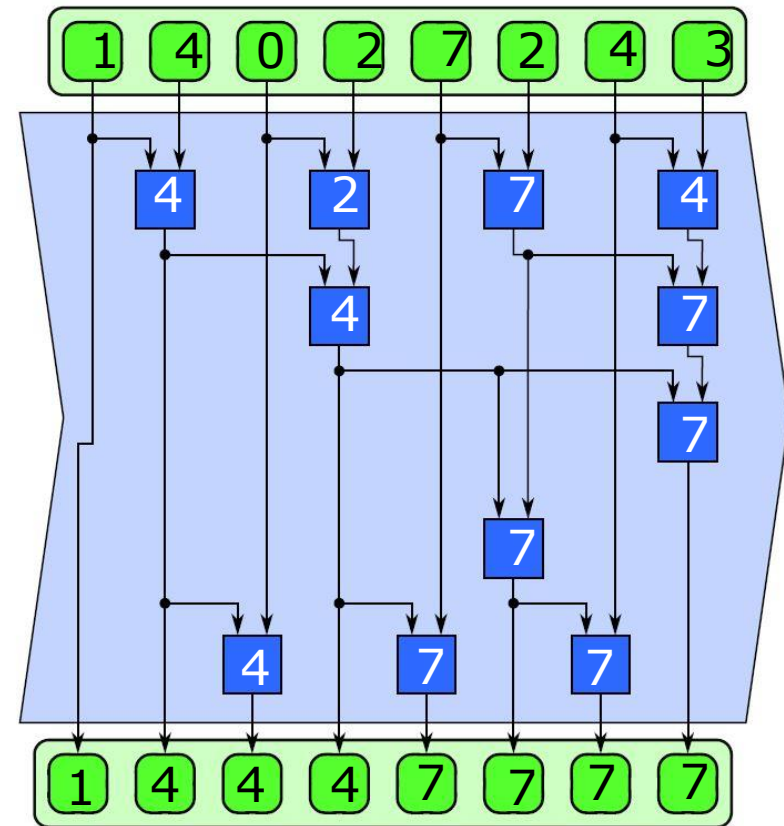
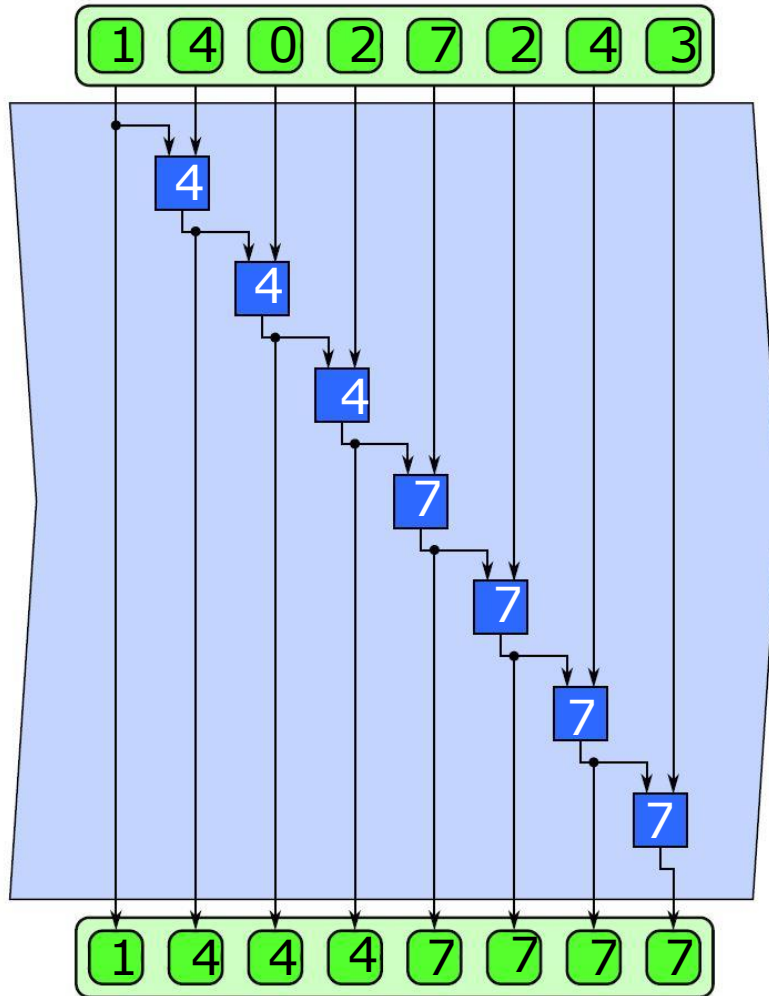
# Scan – Maximum Example

22




# Scan – Maximum Example

23

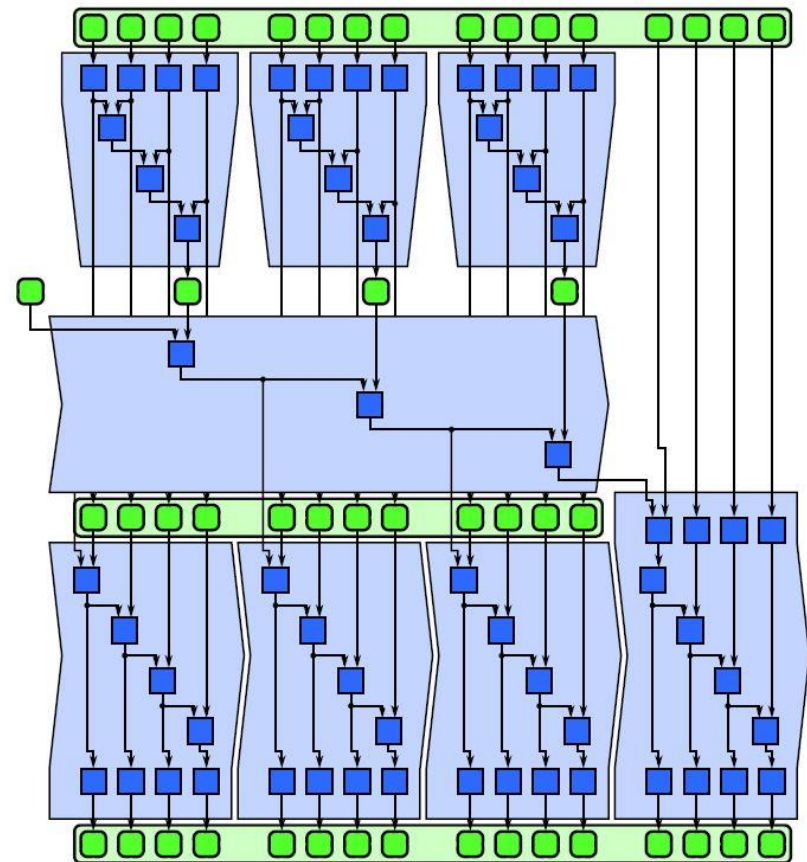



- ❑ Three phase scan with tiling





- Just like reduce, we can also fuse the **map** pattern with the **scan** pattern



- ❑ We can sort an array via a pair of a map and a reduce
- ❑ Map each element into a vector containing just that element
  - ▶  $<>$  is the merge operation:  
 $[1,3,5,7] <> [2,6,15] = [1,2,3,5,6,7,15]$
  - ▶  $[]$  is the empty list
- ❑ How fast is this?

Start with [14,3,4,8,7,52,1]

Map to [[14],[3],[4],[8],[7],[52],[1]]

Reduce:

$$[14] <> ([3] <> ([4] <> ([8] <> ([7] <> ([52] <> [1])))))$$
$$= [14] <> ([3] <> ([4] <> ([8] <> ([7] <> [1,52]))))$$
$$= [14] <> ([3] <> ([4] <> ([8] <> [1,7,52])))$$
$$= [14] <> ([3] <> ([4] <> [1,7,8,52]))$$
$$= [14] <> ([3] <> [1,4,7,8,52])$$
$$= [14] <> [1,3,4,7,8,52]$$
$$= [1,3,4,7,8,14,52]$$

- ❑ How long did that take?
- ❑ We did  $O(n)$  merges...but each one took  $O(n)$  time
- ❑  $O(n^2)$
- ❑ We wanted merge sort, but instead we got insertion sort!

Start with [14,3,4,8,7,52,1]

Map to [[14],[3],[4],[8],[7],[52],[1]]

Reduce:

$$((([14] <> [3]) <> ([4] <> [8]))) <> (([7] <> [52]) <> [1])$$
$$= ([3,14] <> [4,8]) <> ([7,52] <> [1])$$
$$= [3,4,8,14] <> [1,7,52]$$
$$= [1,3,4,7,8,14,52]$$

- ❑ Even if we only had a single processor this is better
  - ▶ We do  $O(\log n)$  merges
  - ▶ Each one is  $O(n)$
  - ▶ So  $O(n \cdot \log(n))$
- ❑ But opportunity for parallelism is not so great
  - ▶  $O(n)$  assuming sequential merge
  - ▶ Takeaway: the shape of reduction matters!