



058165 - PARALLEL COMPUTING

Fabrizio Ferrandi

a.a. 2022-2023

- ❑ “Structured Parallel Programming: Patterns for Efficient Computation,” Michael McCool, Arch Robinson, James Reinders, 1st edition, Morgan Kaufmann, ISBN: 978-0-12-415993-8, 2012

- ❑ Gather Pattern
 - ▶ Shifts, Zip, Unzip
- ❑ Scatter Pattern
 - ▶ Collision Rules: atomic, permutation, merge, priority
- ❑ Pack Pattern
 - ▶ Split, Unsplit, Bin
 - ▶ Fusing Map and Pack
 - ▶ Expand
- ❑ Partitioning Data
- ❑ AoS vs. SoA
- ❑ Example Implementation: AoS vs. SoA

- ❑ Performance is often more limited by data movement than by computation
 - ▶ Transferring data across memory layers is costly
 - locality is important to minimize data access times
 - data organization and layout can impact this
 - ▶ Transferring data across networks can take many cycles
 - attempting to minimize the # messages and overhead is important
 - ▶ Data movement also costs more in power
- ❑ For “data intensive” application, it is a good idea to design the data movement first
 - ▶ Design the computation around the data movements
 - ▶ Applications such as search and sorting are all about data movement and reorganization

- ❑ Remember we are looking to do things in parallel
- ❑ How to be faster than the sequential algorithm?
- ❑ Similar consistency issues arise as when dealing with computation parallelism
- ❑ Here we are concerned more with parallel data movement and management issues
- ❑ Might involve the creation of additional data structures (e.g., for holding intermediate data)

- ❑ Gather pattern creates a (source) collection of data by reading from another (input) data collection
 - ▶ Given a collection of (ordered) indices
 - ▶ Read data from the source collection at each index
 - ▶ Write data to the output collection in index order
- ❑ Transfers from source collection to output collection
 - ▶ Element type of output collection is the same as the source
 - ▶ Shape of the output collection is that of the index collection
 - same dimensionality
- ❑ Can be considered a combination of map and random serial read operations
 - ▶ Essentially does a number of random reads in parallel

```
1  template<typename Data, typename Idx>
2  void gather(
3      size_t n, // number of elements in data collection
4      size_t m, // number of elements in index collection
5      Data a[], // input data collection (n elements)
6      Data A[], // output data collection (m elements)
7      Idx idx[] // input index collection (m elements)
8  ) {
9      for (size_t i = 0; i < m; ++i) {
10         size_t j = idx[i]; // get ith index
11         assert(0 <= j && j < n); // check array bounds
12         A[i] = a[j]; // perform random read
13     }
14 }
```

Serial implementation of gather in pseudocode

```
1  template<typename Data, typename Idx>
2  void gather(
3      size_t n, // number of elements in data collection
4      size_t m, // number of elements in index collection
5      Data a[], // input data collection (n elements)
6      Data A[], // output data collection (m elements)
7      Idx idx[] // input index collection (m elements)
8  ) {
9      for (size_t i = 0; i < m; ++i) {
10         size_t j = idx[i]; // get ith index
11         assert(0 <= j && j < n); // check array bounds
12         A[i] = a[j]; // perform random read
13     }
14 }
```

Serial implementation of gather in pseudocode
Do you see opportunities for parallelism?

Gather: Serial Implementation

9

```
1  template<typename Data, typename Idx>
2  void gather(
3      size_t n, // number of elements in data collection
4      size_t m, // number of elements in index collection
5      Data a[], // input data collection (n elements)
6      Data A[], // output data collection (m elements)
7      Idx idx[] // input index collection (m elements)
8  ) {
9      for (size_t i = 0; i < m; ++i) {
10         size_t j = idx[i]; // get ith index
11         assert(0 <= j && j < n); // check array bounds
12         A[i] = a[j]; // perform random read
13     }
14 }
```


Parallelize over
for loop to
perform random
read

Serial implementation of gather in pseudocode
Are there any conflicts that arise?

Gather: Defined (parallel perspective)

10

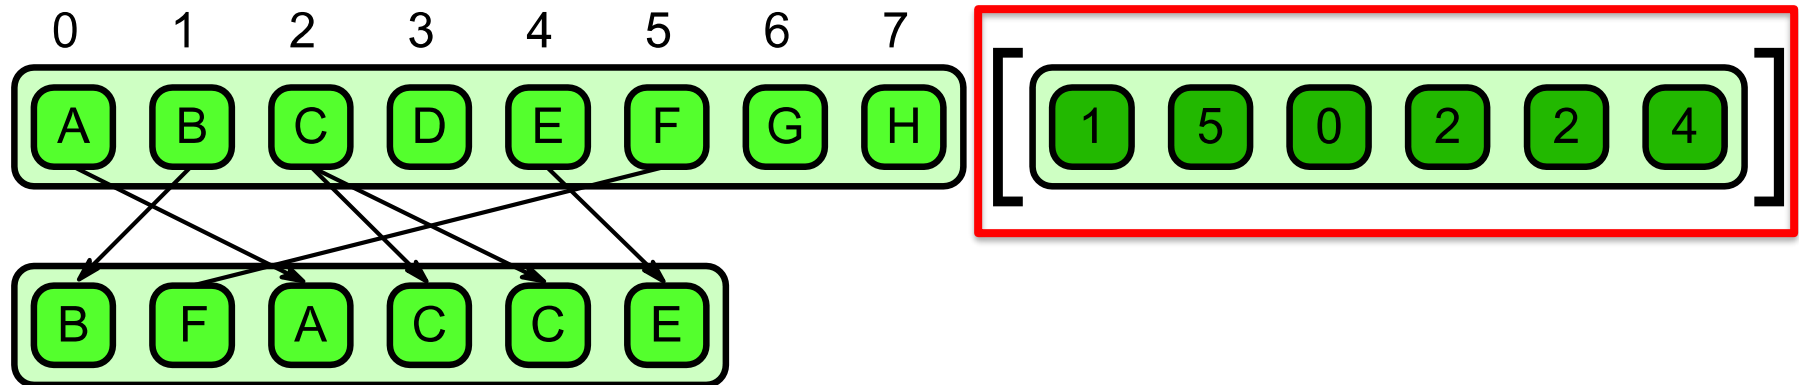
- Results from the combination of a map with a random read



- Simple pattern, but with many special cases that make the implementation more efficient

Gather: Defined

11




Given a **collection of read locations**

- address or array indices

Gather: Defined

12




Given a collection of read locations

□ address or array indices

and a **source array**

Gather: Defined

13




Given a collection of read locations

□ address or array indices

and a source array

gather all the data from the source array at the given locations and places them into an **output collection**

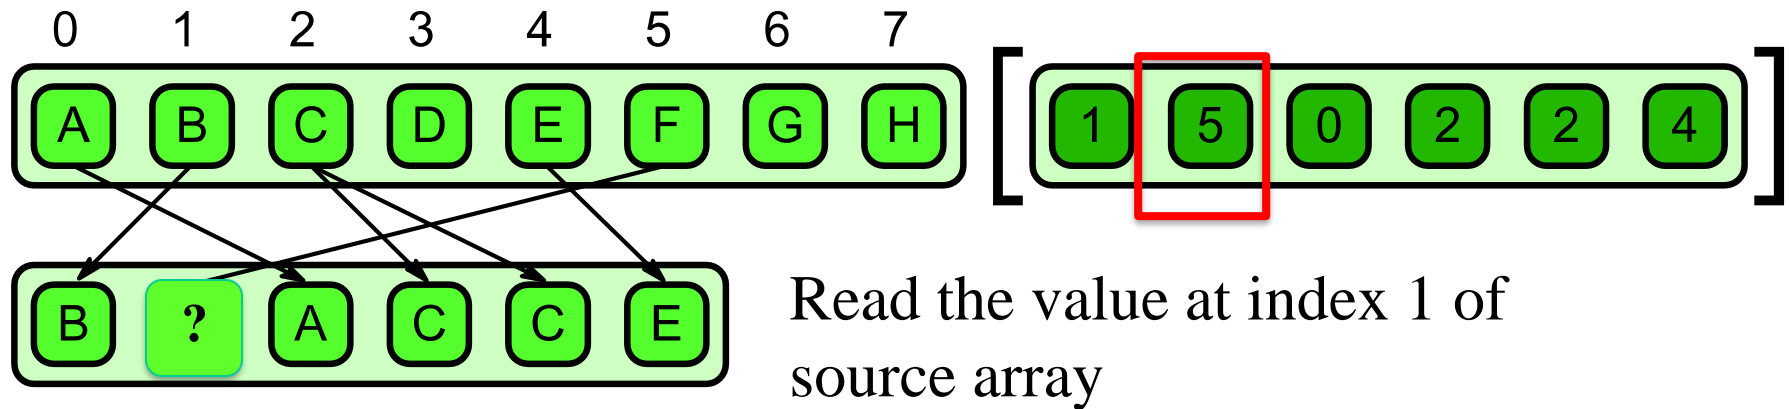


Given a collection of read locations

□ address or array indices

and a source array

gather all the data from the source array at the given locations and places them into an **output collection**




Given a collection of read locations

□ address or array indices

and a source array

gather all the data from the source array at the given locations and places them into an **output collection**



Given a collection of read locations

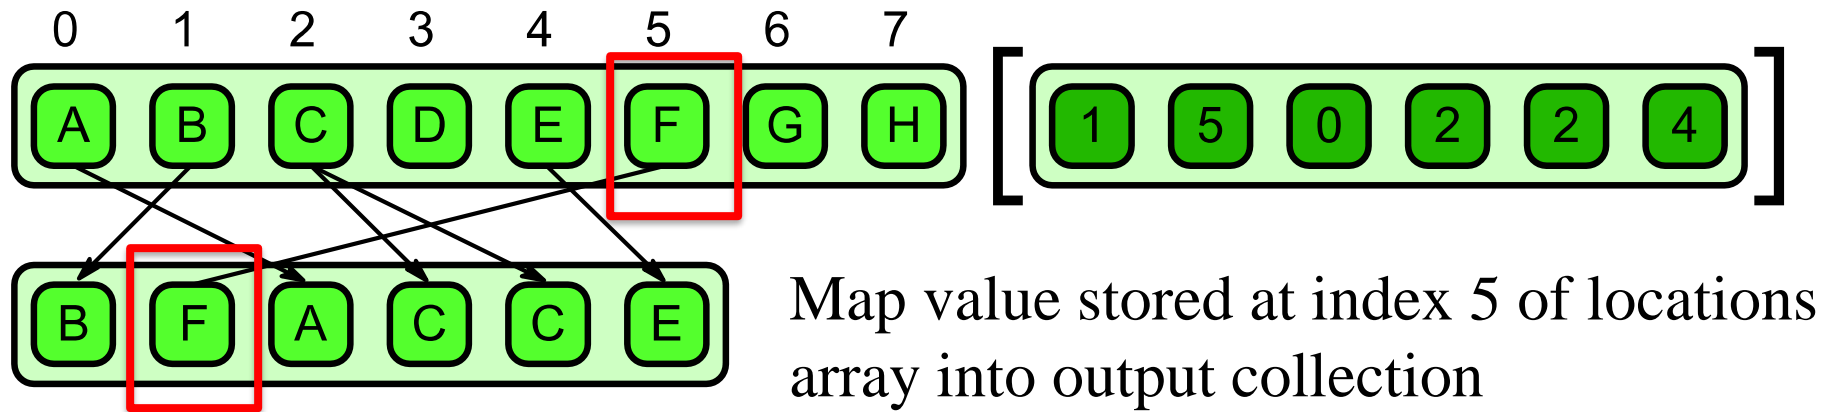
□ address or array indices

and a source array

gather all the data from the source array at the given locations and places them into an **output collection**

Gather: Defined

17



Given a collection of read locations

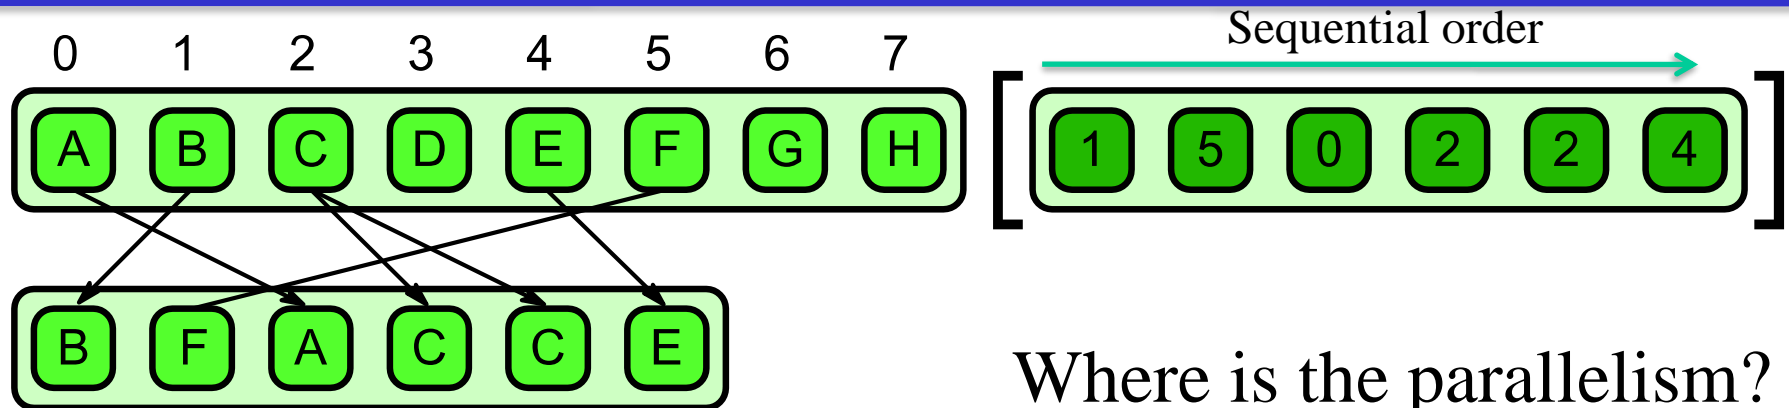
□ address or array indices

and a source array

gather all the data from the source array at the given locations and places them into an **output collection**

Gather: Defined

18



Where is the parallelism?

Given a collection of read locations

□ address or array indices

and a source array

gather all the data from the source array at the given locations and places them into an **output collection**

Quiz 1

19

Given the following locations and source array, use a gather to determine what values should go into the output collection:

0 1 2 3 4 5 6 7 8 9 10 11

3	7	0	1	4	0	0	4	5	3	1	0
---	---	---	---	---	---	---	---	---	---	---	---


1	9	6	9	3
---	---	---	---	---

?	?	?	?	?
---	---	---	---	---

Quiz 1


20

Given the following locations and source array, use a gather to determine what values should go into the output collection:



Gather: Array Size


21



- ❑ Output data collection has the same number of elements as the number of indices in the index collection
 - ▶ Same dimensionality

Gather: Array Size

22




- ❑ Output data collection has the same number of elements as the number of indices in the index collection
- ❑ Elements of the output collection are the same type as the input data collection

- ❑ Gather Pattern
 - ▶ Shifts, Zip, Unzip
- ❑ Scatter Pattern
 - ▶ Collision Rules: atomic, permutation, merge, priority
- ❑ Pack Pattern
 - ▶ Split, Unsplit, Bin
 - ▶ Fusing Map and Pack
 - ▶ Expand
- ❑ Partitioning Data
- ❑ AoS vs. SoA
- ❑ Example Implementation: AoS vs. SoA

Special Case of Gather: Shifts

24




- ❑ Moves data to the left or right in memory
- ❑ Data accesses are offset by fixed distances

- ❑ Regular data movement
- ❑ Variants from how boundary conditions handled
 - ▶ Requires “out of bounds” data at edge of the array
 - ▶ Options: default value, duplicate, rotate
- ❑ Shifts can be handled efficiently with vector instructions because of regularity
 - ▶ Shift multiple data elements at the same time
- ❑ Shifts can also take advantage of good data locality

- ❑ Gather Pattern
 - ▶ Shifts, Zip, Unzip
- ❑ Scatter Pattern
 - ▶ Collision Rules: atomic, permutation, merge, priority
- ❑ Pack Pattern
 - ▶ Split, Unsplit, Bin
 - ▶ Fusing Map and Pack
 - ▶ Expand
- ❑ Partitioning Data
- ❑ AoS vs. SoA
- ❑ Example Implementation: AoS vs. SoA


Special Case of Gather: Zip

27



Where is the parallelism?

❑ Function is to interleaves data (like a zipper)




- ❑ Given two separate arrays of real parts and imaginary parts
- ❑ Use zip to combine them into a sequence of real and imaginary pairs

- ❑ Can be generalized to more elements
- ❑ Can zip data of unlike types

- ❑ Gather Pattern
 - ▶ Shifts, Zip, **Unzip**
- ❑ Scatter Pattern
 - ▶ Collision Rules: atomic, permutation, merge, priority
- ❑ Pack Pattern
 - ▶ Split, Unsplit, Bin
 - ▶ Fusing Map and Pack
 - ▶ Expand
- ❑ Partitioning Data
- ❑ AoS vs. SoA
- ❑ Example Implementation: AoS vs. SoA

Special Case of Gather: Unzip

31




Where is the parallelism?

- ❑ Reverses a zip
- ❑ Extracts sub-arrays at certain offsets and strides from an input array

Unzip Example

32



- ❑ Given a sequence of complex numbers organized as pairs
- ❑ Use unzip to extract real and imaginary parts into separate arrays

- ❑ Gather Pattern
 - ▶ Shifts, Zip, Unzip
- ❑ Scatter Pattern
 - ▶ Collision Rules: atomic, permutation, merge, priority
- ❑ Pack Pattern
 - ▶ Split, Unsplit, Bin
 - ▶ Fusing Map and Pack
 - ▶ Expand
- ❑ Partitioning Data
- ❑ AoS vs. SoA
- ❑ Example Implementation: AoS vs. SoA

Gather

- ❑ Combination of map with random **reads**
- ❑ Read locations provided as input

Scatter

- ❑ Combination of map with random **writes**
- ❑ Write locations provided as input
- ❑ Race conditions ... Why?

Scatter: Serial Implementation

35

```
1  template<typename Data, typename Idx>
2  void scatter(
3      size_t n, // number of elements in output data collection
4      size_t m, // number of elements in input data and index collection
5      Data a[], // input data collection (m elements)
6      Data A[], // output data collection (n elements)
7      Idx idx[] // input index collection (m elements)
8  ) {
9      for (size_t i = 0; i < m; ++i) {
10         size_t j = idx[i]; // get ith index
11         assert(0 <= j && j < n); // check output array bounds
12         A[j] = a[i]; // perform random write
13     }
14 }
```

Serial implementation of scatter in pseudocode

```
1  template<typename Data, typename Idx>
2  void scatter(
3      size_t n, //number of elements in output data collection
4      size_t m, //number of elements in input data and index collection
5      Data a[], //input data collection (m elements)
6      Data A[], //output data collection (n elements)
7      Idx idx[] //input index collection (m elements)
8  ) {
9      for (size_t i = 0; i < m; ++i) {
10         size_t j = idx[i]; //get ith index
11         assert(0 <= j && j < n); //check output array bounds
12         A[j] = a[i]; //perform random write
13     }
14 }
```


Parallelize over
for loop to
perform random
write

Serial implementation of scatter in pseudocode

- ❑ Results from the combination of a map with a random write
- ❑ Writes to the same location are possible
- ❑ Parallel writes to the same location are **collisions**

Scatter: Defined


38



Given a collection of **input data**

Scatter: Defined

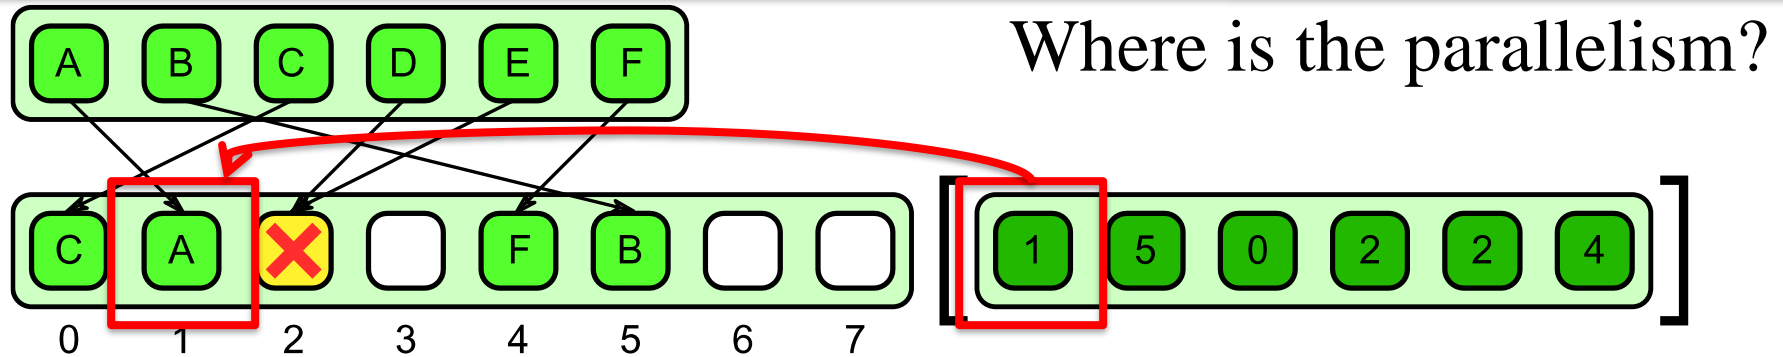
39



Given a collection of input data
and a collection of **write locations**

Scatter: Defined

40



Given a collection of input data
and a collection of write locations
scatter data to the **output collection**

Problems?

Does the output collection have to be larger in size?

Quiz 2

41

Given the following locations and source array, what values should go into the input collection:

0 1 2 3 4 5 6 7 8 9 10 11

3	7	0	1	4	0	0	4	5	3	1	0
---	---	---	---	---	---	---	---	---	---	---	---

2	4	1	5	5	0	4	2	1	2	1	4
---	---	---	---	---	---	---	---	---	---	---	---


?	?	?		?	?
---	---	---	--	---	---

Quiz 2

42


Given the following locations and source array, what values should go into the input collection:

0 1 2 3 4 5 6 7 8 9 10 11



Scatter: Race Conditions


43




Given a collection of input data
and a collection of write locations
scatter data to the output collection

Race Condition: Two (or more) values being written to the same location in output collection. Result is undefined unless enforce rules. **Need rules to resolve collisions!**

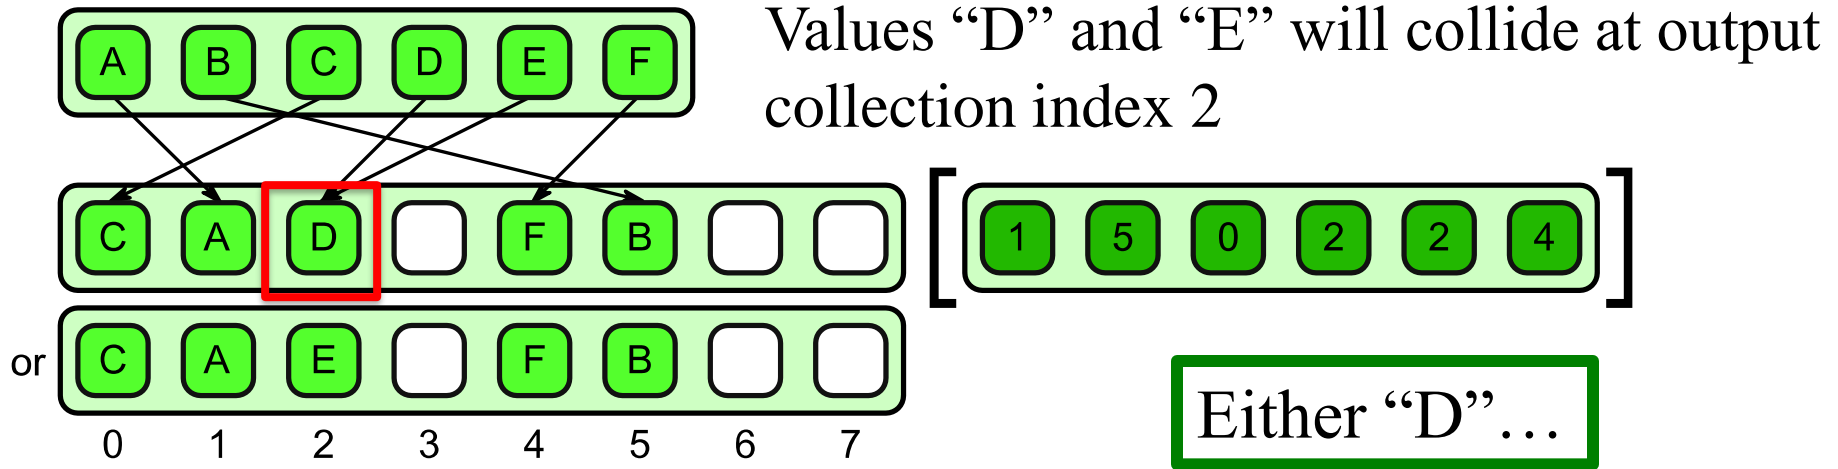
- ❑ Gather Pattern
 - ▶ Shifts, Zip, Unzip
- ❑ Scatter Pattern
 - ▶ Collision Rules: atomic, permutation, merge, priority
- ❑ Pack Pattern
 - ▶ Split, Unsplit, Bin
 - ▶ Fusing Map and Pack
 - ▶ Expand
- ❑ Partitioning Data
- ❑ AoS vs. SoA
- ❑ Example Implementation: AoS vs. SoA



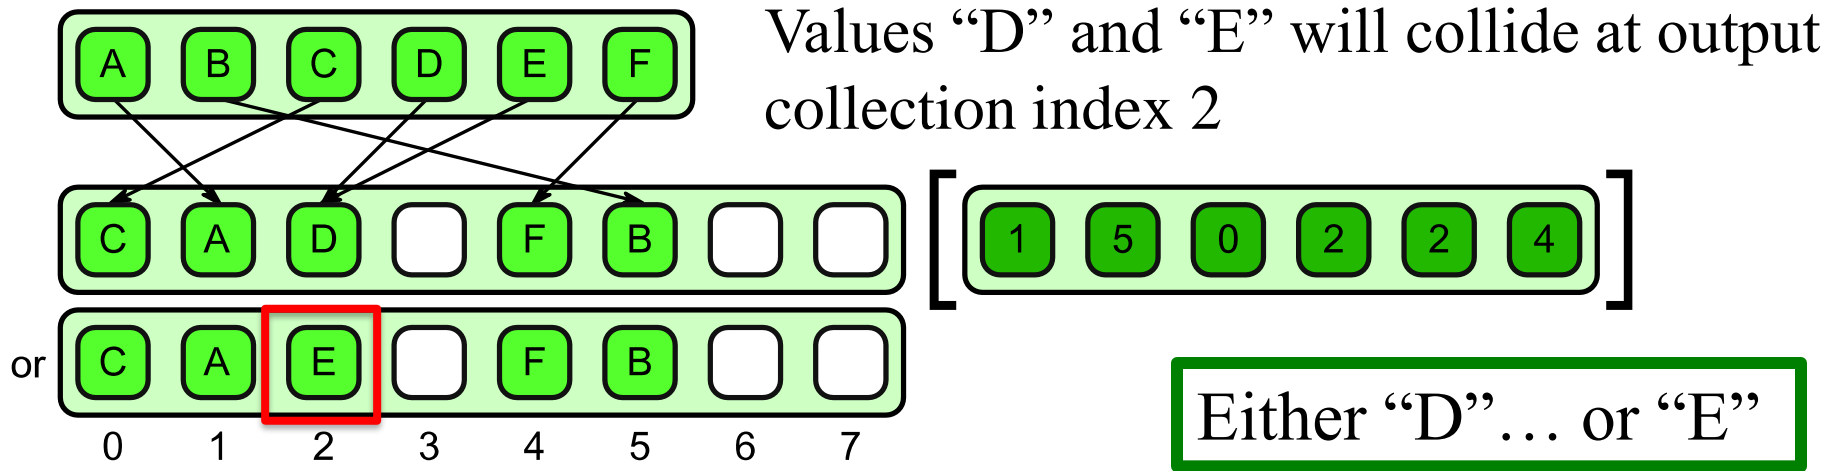
- ❑ Non-deterministic approach
- ❑ Upon collision, one and only one of the values written to a location will be written in its entirety



- ❑ Non-deterministic approach
- ❑ Upon collision, one and only one of the values written to a location will be written in its entirety




- ❑ Non-deterministic approach
- ❑ Upon collision, one and only one of the values written to a location will be written in its entirety
- ❑ No rule determines which of the input items will be retained



- ❑ Non-deterministic approach
- ❑ Upon collision, one and only one of the values written to a location will be written in its entirety
- ❑ No rule determines which of the input items will be retained

Collision Resolution: Permutation Scatter


49



- ❑ Pattern simply states that collisions are **illegal**
 - ▶ Output is a permutation of the input
- ❑ Check for collisions in advance
 - turn scatter into gather
- ❑ Examples
 - ▶ FFT scrambling, matrix/image transpose, unpacking

Collision Resolution: Merge Scatter


50




- ❑ Associative and commutative operators are provided to merge elements in case of a collision

Collision Resolution: Merge Scatter


51



- ❑ Associative and commutative operators are provided to merge elements in case of a collision




- ❑ Associative and commutative operators are provided to merge elements in case of a collision
- ❑ Use addition as the merge operator
- ❑ Both associative and commutative properties are required since scatters to a particular location could occur in any order



- ❑ Every element in the input array is assigned a priority based on its position
- ❑ Priority is used to decide which element is written in case of a collision
- ❑ Example
 - ▶ 3D graphics rendering

- ❑ Scatter is a more expensive than gather
 - ▶ Writing has cache line consequences
 - ▶ May cause additional reading due to cache conflicts
 - ▶ **False sharing** is a problem that arises
 - writes from different cores go to the same cache line
- ❑ Can avoid problems if addresses are known “in advance”
 - ▶ Allows optimizations to be applied
 - ▶ Convert addresses for a scatter into those for a gather
 - ▶ Useful if the same pattern of scatter address will be used repeatedly so the cost is amortized


- ❑ Gather Pattern
 - ▶ Shifts, Zip, Unzip
- ❑ Scatter Pattern
 - ▶ Collision Rules: atomic, permutation, merge, priority
- ❑ Pack Pattern
 - ▶ Split, Unsplit, Bin
 - ▶ Fusing Map and Pack
 - ▶ Expand
- ❑ Partitioning Data
- ❑ AoS vs. SoA
- ❑ Example Implementation: AoS vs. SoA



- ❑ Used to eliminate unused elements from a collection
- ❑ Retained elements are moved so they are contiguous in memory
- ❑ Goal is to improve the performance ... How?

Pack Algorithm


57



1. Convert input array of Booleans into integer 0's and 1's

Pack Algorithm


58



1. Convert input array of Booleans into integer 0's and 1's
2. Exclusive scan of this array with the addition operation

Pack Algorithm


59



1. Convert input array of Booleans into integer 0's and 1's
2. Exclusive scan of this array with the addition operation
3. Write values to output array based on offsets

Unpack: Defined

60




- ❑ Inverse of pack operation
- ❑ Given the same data on which elements were kept and which were discarded, spread elements back in their original locations

- ❑ Gather Pattern
 - ▶ Shifts, Zip, Unzip
- ❑ Scatter Pattern
 - ▶ Collision Rules: atomic, permutation, merge, priority
- ❑ Pack Pattern
 - ▶ Split, Unsplit, Bin
 - ▶ Fusing Map and Pack
 - ▶ Expand
- ❑ Partitioning Data
- ❑ AoS vs. SoA
- ❑ Example Implementation: AoS vs. SoA

Generalization of Pack: Split

62




- Generalization of pack pattern
- Elements are moved to upper or lower half of output collection based on some state
- Does not lose information like pack

Upper half of output collection: values equal to 0

Generalization of Pack: Split

63




**Lower half of output
collection: values equal to 1**

- Generalization of pack pattern
- Elements are moved to upper or lower half of output collection based on some state
- Does not lose information like pack

Generalization of Pack: Unsplit


64



- ❑ Inverse of split
- ❑ Creates **output collection** based on original input collection

Generalization of Pack: Bin

65




- Generalized split to support more categories (>2)
- Examples
 - Radix sort
 - Pattern classification

4 different categories = 4 bins

- ❑ Gather Pattern
 - ▶ Shifts, Zip, Unzip
- ❑ Scatter Pattern
 - ▶ Collision Rules: atomic, permutation, merge, priority
- ❑ Pack Pattern
 - ▶ Split, Unsplit, Bin
 - ▶ **Fusing Map and Pack**
 - ▶ Expand
- ❑ Partitioning Data
- ❑ AoS vs. SoA
- ❑ Example Implementation: AoS vs. SoA

Fusion of Map and Pack

67



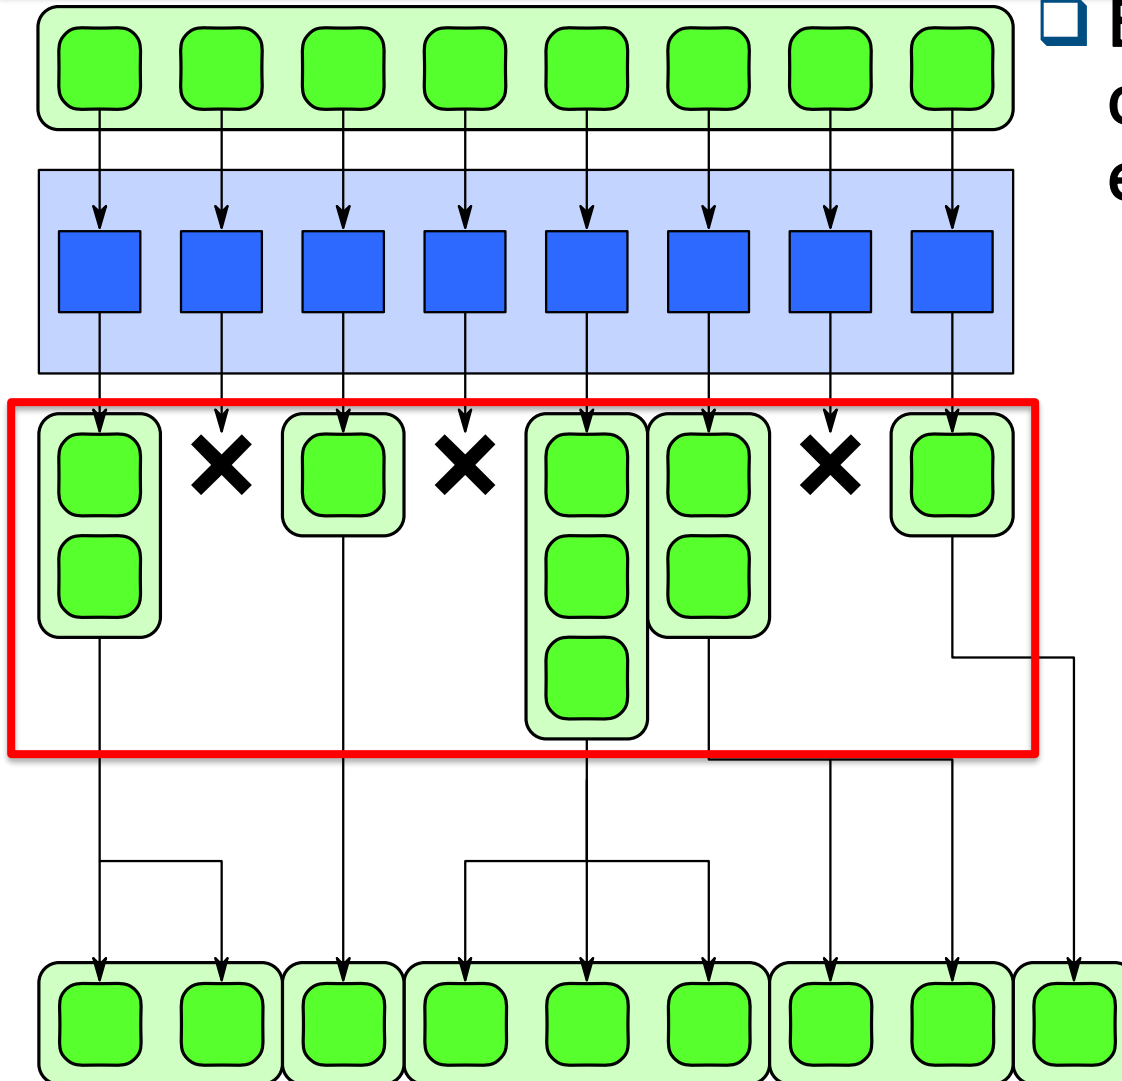
- ❑ Advantageous if most of the elements of a map are discarded
- ❑ **Map** checks pairs for collision
- ❑ **Pack** stores only actual collisions
- ❑ Output BW \sim results reported, not number of pairs tested
- ❑ Each element can output 0 or 1 element

- ❑ Gather Pattern
 - ▶ Shifts, Zip, Unzip
- ❑ Scatter Pattern
 - ▶ Collision Rules: atomic, permutation, merge, priority
- ❑ Pack Pattern
 - ▶ Split, Unsplit, Bin
 - ▶ Fusing Map and Pack
 - ▶ **Expand**
- ❑ Partitioning Data
- ❑ AoS vs. SoA
- ❑ Example Implementation: AoS vs. SoA

Generalization of Pack: Expand

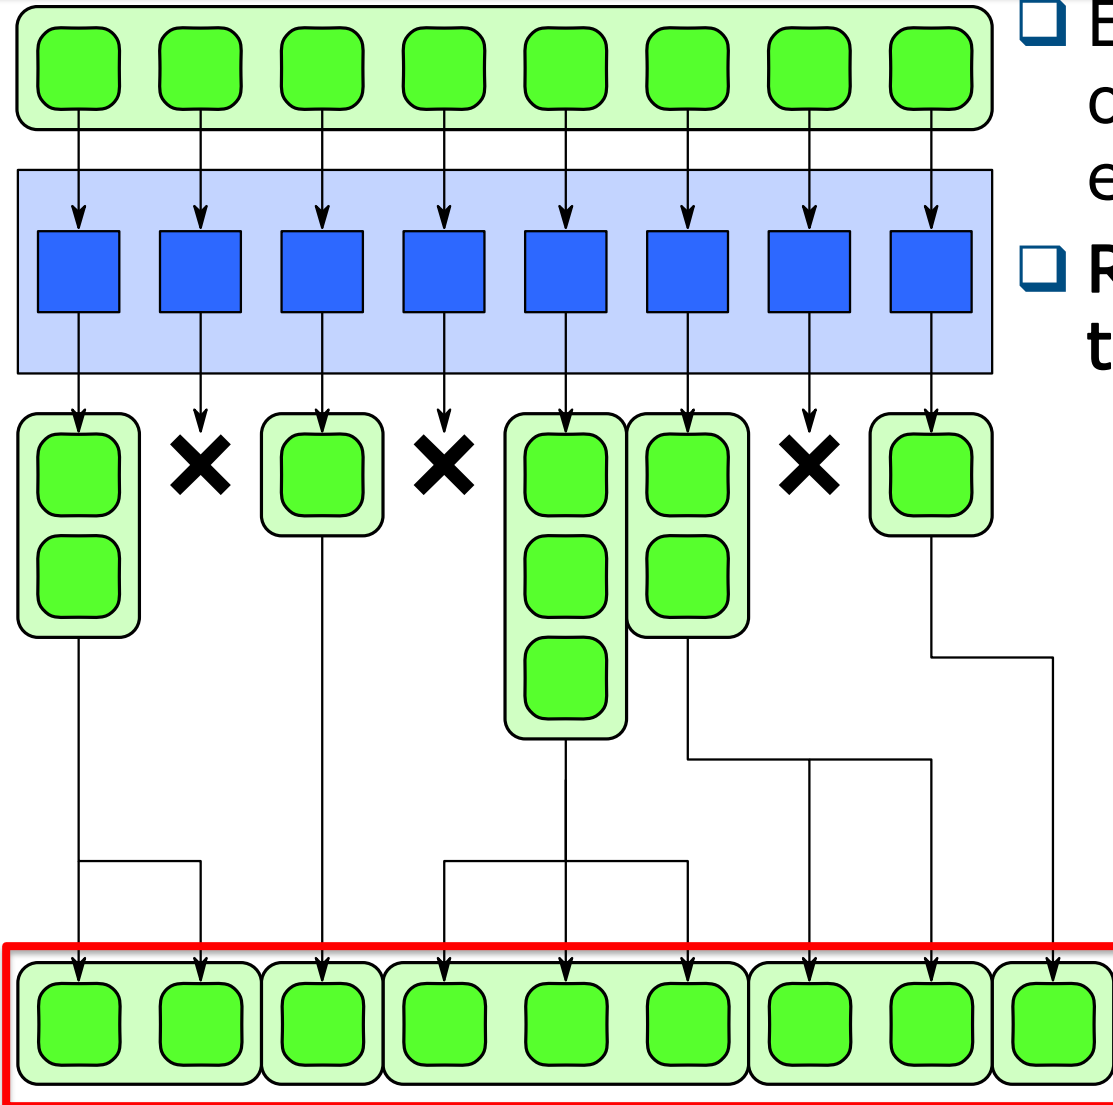
69

- Each element can output any number of elements



Generalization of Pack: Expand

70





- Each element can output any number of elements
- Results are fused together in order

- ❑ Gather Pattern
 - ▶ Shifts, Zip, Unzip
- ❑ Scatter Pattern
 - ▶ Collision Rules: atomic, permutation, merge, priority
- ❑ Pack Pattern
 - ▶ Split, Unsplit, Bin
 - ▶ Fusing Map and Pack
 - ▶ Expand
- ❑ Partitioning Data
- ❑ AoS vs. SoA
- ❑ Example Implementation: AoS vs. SoA



- ❑ Common strategy:
 1. Divide up the computational domain into sections
 2. Work on the sections individually
 3. Combine the results
- ❑ Methods
 - ▶ Divide-and-conquer
 - ▶ Fork-join (discussed in Chapter 8)
 - ▶ Geometric decomposition
 - ▶ Partitions
 - ▶ Segments

Partitioning

73





- Data is divided into
 - ▶ non-overlapping
 - ▶ equal-sized regions




- ❑ Data is divided into
 - ▶ non-overlapping
 - ▶ equal-sized regions

Avoid write conflicts and race conditions



- ❑ Data is divided into **non-uniform** non-overlapping regions
- ❑ Start of each segment can be marked using:
 - ▶ Array of integers
 - ▶ Array of Boolean flags


- ❑ Gather Pattern
 - ▶ Shifts, Zip, Unzip
- ❑ Scatter Pattern
 - ▶ Collision Rules: atomic, permutation, merge, priority
- ❑ Pack Pattern
 - ▶ Split, Unsplit, Bin
 - ▶ Fusing Map and Pack
 - ▶ Expand
- ❑ Partitioning Data
- ❑ AoS vs. SoA
- ❑ Example Implementation: AoS vs. SoA



- ☐ May lead to better cache utilization if data is accessed randomly

Structures of Arrays (SoA)

78



- Typically better for vectorization and avoidance of false sharing

Data Layout Options


Array of Structures (AoS), padding at end




Array of Structures (AoS), padding after each structure



Structure of Arrays (SoA), padding at end



Structure of Arrays (SoA), padding after each component



- ❑ Gather Pattern
 - ▶ Shifts, Zip, Unzip
- ❑ Scatter Pattern
 - ▶ Collision Rules: atomic, permutation, merge, priority
- ❑ Pack Pattern
 - ▶ Split, Unsplit, Bin
 - ▶ Fusing Map and Pack
 - ▶ Expand
- ❑ Partitioning Data
- ❑ AoS vs. SoA
- ❑ Example Implementation: AoS vs. SoA

AoS Code

```
struct node {  
    float x, y, z;  
};  
struct node NODES[1024];  
  
float dist[1024];  
for(i=0;i<1024;i+=16){  
    float x[16],y[16],z[16],d[16];  
    x[:] = NODES[i:16].x;  
    y[:] = NODES[i:16].y;  
    z[:] = NODES[i:16].z;  
    d[:] = sqrtf(x[:]*x[:] + y[:]*y[:] + z[:]*z[]);  
    dist[i:16] = d[:];  
}
```

SoA Code

```
struct node1 {  
    float x[1024], y[1024], z[1024];  
}  
struct node1 NODES1;  
  
float dist[1024];  
for(i=0;i<1024;i+=16){  
    float x[16],y[16],z[16],d[16];  
    x[:] = NODES1.x[i:16];  
    y[:] = NODES1.y[i:16];  
    z[:] = NODES1.z[i:16];  
    d[:] = sqrtf(x[:]*x[:] + y[:]*y[:] + z[:]*z[]);  
    dist[i:16] = d[:];  
}
```

```
struct node {  
    float x, y, z;  
};  
  
struct node NODES[1024];  
  
float dist[1024];  
for(i=0;i<1024;i+=16){  
    float x[16],y[16],z[16],d[16];  
    x[:] = NODES[i:16].x;  
    y[:] = NODES[i:16].y;  
    z[:] = NODES[i:16].z;  
    d[:] = sqrtf(x[:] * x[:] + y[:] * y[:] + z[:] * z[]);  
    dist[i:16] = d[:];  
}
```

- ❑ Most logical data organization layout
- ❑ Extremely difficult to access memory for reads (gathers) and writes (scatters)
- ❑ Prevents efficient vectorization

```
struct node1 {  
    float x[1024], y[1024], z[1024];  
}  
struct node1 NODES1;  
  
float dist[1024];  
for(i=0;i<1024;i+=16){  
    float x[16],y[16],z[16],d[16];  
    x[:] = NODES1.x[i:16];  
    y[:] = NODES1.y[i:16];  
    z[:] = NODES1.z[i:16];  
    d[:] = sqrtf(x[:] * x[:] + y[:] * y[:] + z[:] * z[]);  
    dist[i:16] = d[:];  
}
```

- ❑ Separate arrays for each structure-field keeps memory accesses contiguous when vectorization is performed over structure instances