



Parallel Patterns 4th part

058165 – Parallel Computing

Fabrizio Ferrandi

Politecnico di Milano

Dipartimento di Elettronica, Informazione e Bioingegneria

fabrizio.ferrandi@polimi.it

- ❑ “Structured Parallel Programming: Patterns for Efficient Computation,” Michael McCool, Arch Robinson, James Reinders, 1st edition, Morgan Kaufmann, ISBN: 978-0-12-415993-8, 2012

- ❑ What is the stencil pattern?
 - ▶ Update alternatives
 - ▶ 2D Jacobi iteration
- ❑ Implementing stencil with shift
- ❑ Stencil and cache optimizations
- ❑ Stencil and communication optimizations
- ❑ Recurrence

- ❑ A stencil pattern is a map where each output depends on a “neighborhood” of inputs
- ❑ These inputs are a set of fixed offsets relative to the output position
- ❑ A stencil output is a function of a “neighborhood” of elements in an input collection
 - ▶ Applies the stencil to select the inputs
- ❑ Data access patterns of stencils are regular
 - ▶ Stencil is the “shape” of “neighborhood”
 - ▶ Stencil remains the same

Serial Stencil Example (part 1)

5

```
1  template<
2      int NumOff,      // number of offsets
3      typename In,     // type of input locations
4      typename Out,    // type of output locations
5      typename F        // type of function / functor
6  >
7  void stencil(
8      int n,           // number of elements in data collection
9      const In a[],    // input data collection (n elements)
10     Out r[],         // output data collection (n elements)
11     In b,            // boundary value
12     F func,          // function / functor from neighborhood inputs to output
13     const int offsets[] // offsets (NumOffsets elements)
14 ) {
```

Serial Stencil Example (part 2)


6

```
15 // array to hold neighbors
16 In neighborhood[NumOfff];
17 // loop over all output locations
18 for (int i = 0; i < n; ++i) {
19     // loop over all offsets and gather neighborhood
20     for (int j = 0; j < NumOfff; ++j) {
21         // get index of jth input location
22         int k = i+offsets[j];
23         if (0 <= k && k < n) {
24             // read input location
25             neighborhood[j] = a[k];
26         } else {
27             // handle boundary case
28             neighborhood[j] = b;
29         }
30     }
31     // compute output value from input neighborhood
32     r[i] = func(neighborhood);
33 }
34 }
```

How would we
parallelize this?


What is the stencil pattern?

7




What is the stencil pattern?

8




What is the stencil pattern?

9




What is the stencil pattern?

10



What is the stencil pattern?


11



This stencil has
3
elements in the
neighborhood:
 $i-1, i, i+1$


What is the stencil pattern?

12



What is the stencil pattern?


13



And outputs to the i^{th} position of the output array


- ❑ Stencils can operate on one dimensional and multidimensional data
- ❑ Stencil neighborhoods can range from compact to sparse, square to cube, and anything else!
- ❑ It is the pattern of the stencil that determines how the stencil operates in an application

2-Dimensional Stencils




4-point stencil

Center cell (P)
is not used



5-point stencil

Center cell (P)
is used as well




9-point stencil


Center cell (C)
is used as well

Source: http://en.wikipedia.org/wiki/Stencil_code

3-Dimensional Stencils



6-point stencil
(7-point stencil)



24-point stencil
(25-point stencil)

Source: http://en.wikipedia.org/wiki/Stencil_code

□ Here is our array, A

A

0	0	0	0
0	9	7	0
0	6	4	0
0	0	0	0

- Here is our array A
- B is the output array
 - ▶ Initialize to all 0
- Apply a stencil operation to the inner square of the form:

$$B(i,j) = \text{avg} \left(\begin{array}{l} A(i,j), \\ A(i-1,j), A(i+1,j), \\ A(i,j-1), A(i,j+1) \end{array} \right)$$

) What is the stencil?

A


0	0	0	0
0	9	7	0
0	6	4	0
0	0	0	0

- 1) Average all blue squares

A

0	0	0	0
0	9	7	0
0	6	4	0
0	0	0	0

- 1) Average all blue squares
- 2) Store result in B



- 1) Average all blue squares
- 2) Store result in B
- 3) Repeat 1 and 2 for all green squares

A

0	0	0	0
0	9	7	0
0	6	4	0
0	0	0	0

Practice!

Stencil Pattern Practice

23

A

0	0	0	0
0	9	7	0
0	6	4	0
0	0	0	0

B

0	0	0	0
0	4.4	0	0
0	0	0	0
0	0	0	0

Stencil Pattern Practice

24

A

0	0	0	0
0	9	7	0
0	6	4	0
0	0	0	0

B

0	0	0	0
0	4.4	4.0	0
0	0	0	0
0	0	0	0

Stencil Pattern Practice

25

A

0	0	0	0
0	9	7	0
0	6	4	0
0	0	0	0

B

0	0	0	0
0	4.4	4.0	0
0	3.8	0	0
0	0	0	0

A

0	0	0	0
0	9	7	0
0	6	4	0
0	0	0	0

B

0	0	0	0
0	4.4	4.0	0
0	3.8	3.4	0
0	0	0	0

- ❑ Partitioning
- ❑ What is the stencil pattern?
 - ▶ Update alternatives
 - ▶ 2D Jacobi iteration
- ❑ Implementing stencil with shift
- ❑ Stencil and cache optimizations
- ❑ Stencil and communication optimizations
- ❑ Recurrence

Serial Stencil Example (part 1)

28

```
1  template<
2      int NumOff,      // number of offsets
3      typename In,     // type of input locations
4      typename Out,   // type of output locations
5      typename F       // type of function / functor
6  >
7  void stencil(
8      int n,           // number of elements in data collection
9      const In a[],    // input data collection (n elements)
10     Out r[],        // output data collection (n elements)
11     In b,             // boundary value
12     F func,           // function / functor from neighborhood inputs to output
13     const int offsets[] // offsets (NumOffsets elements)
14 ) {
```

Serial Stencil Example (part 2)

29


```
-15    // array to hold neighbors
16    In neighborhood[NumOff];
17    // loop over all output locations
18    for (int i = 0; i < n; ++i) {
19        // loop over all offsets and gather neighborhood
20        for (int j = 0; j < NumOff; ++j) {
21            // get index of jth input location
22            int k = i+offsets[j];
23            if (0 <= k && k < n) {
24                // read input location
25                neighborhood[j] = a[k];
26            } else {
27                // handle boundary case
28                neighborhood[j] = b;
29            }
30        }
31        // compute output value from input neighborhood
32        a[i] r[i] = func(neighborhood);
33    }
34 }
```

How would we
parallelize this?

Updates occur in place!!!

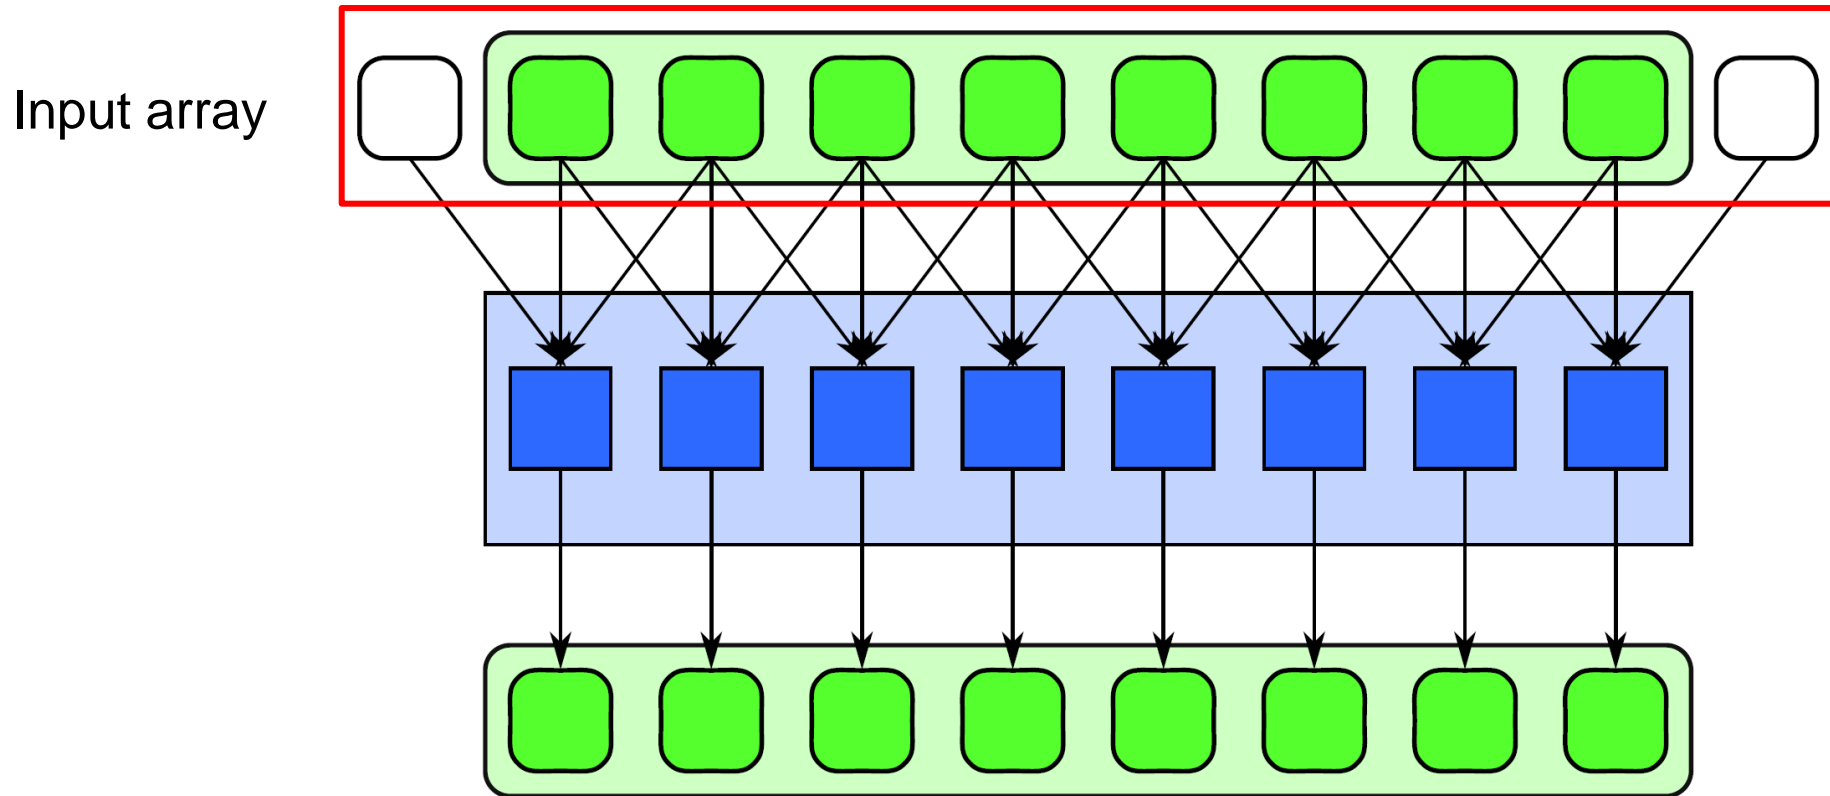
Stencil Pattern with In Place Update

30




Stencil Pattern with In Place Update

31



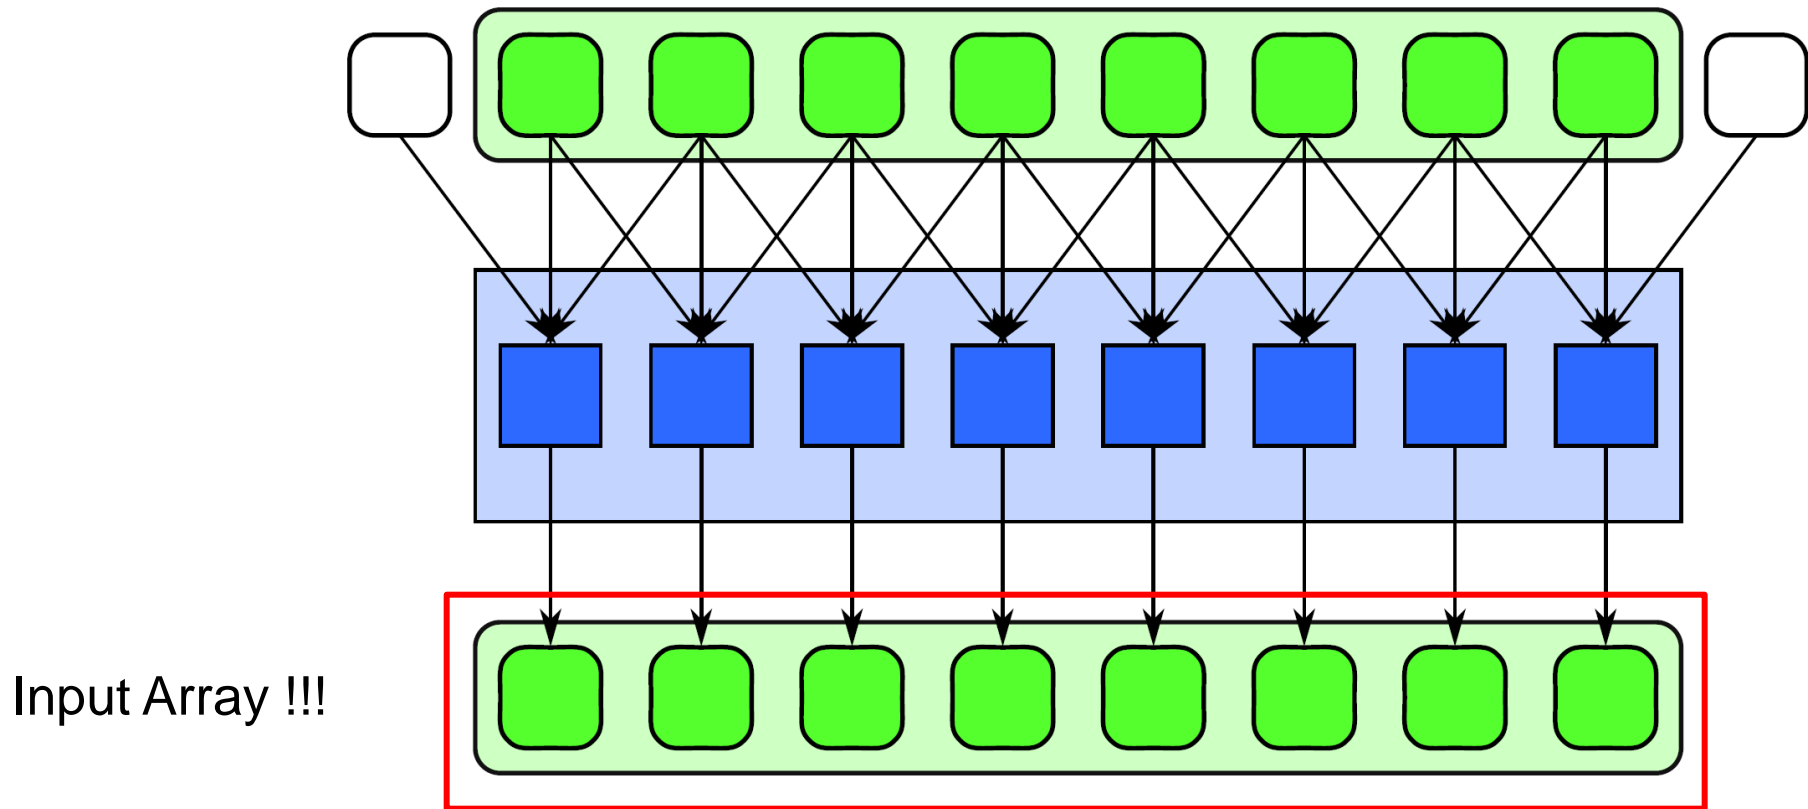
Stencil Pattern with In Place Update

32



Stencil Pattern with In Place Update

33



□ Here is our array, A

A

0	0	0	0
0	9	7	0
0	6	4	0
0	0	0	0

- ❑ Here is our array A
- ❑ Update A in place
- ❑ Apply a stencil operation to the inner square of the form:

$$A(i,j) = \text{avg} \left(\begin{array}{l} A(i,j), \\ A(i-1,j), A(i+1,j), \\ A(i,j-1), A(i,j+1) \end{array} \right)$$

What is the stencil?

A

0	0	0	0
0	9	7	0
0	6	4	0
0	0	0	0

- 1) Average all blue squares

0	0	0	0
0	9	7	0
0	6	4	0
0	0	0	0

- 1) Average all blue squares
- 2) Store result in red square

0	0	0	0
0	9	7	0
0	6	4	0
0	0	0	0

- 1) Average all blue squares
- 2) Store result in red square
- 3) Repeat 1 and 2 for all green squares

0	0	0	0
0	9	7	0
0	6	4	0
0	0	0	0

Practice!

Stencil Pattern Practice

40

A

0	0	0	0
0	9	7	0
0	6	4	0
0	0	0	0

A

0	0	0	0
0	4.4	7	0
0	6	4	0
0	0	0	0

What is the stencil pattern?

42

A

0	0	0	0
0	4.4	7	0
0	6	4	0
0	0	0	0

What is the stencil pattern?

43

A

0	0	0	0
0	4.4	3.08	0
0	6	4	0
0	0	0	0

What is the stencil pattern?

44

A

0	0	0	0
0	4.4	3.08	0
0	6	4	0
0	0	0	0

What is the stencil pattern?

45

A

0	0	0	0
0	4.4	3.08	0
0	2.88	4	0
0	0	0	0

What is the stencil pattern?

46

A

0	0	0	0
0	4.4	3.08	0
0	2.88	4	0
0	0	0	0

What is the stencil pattern?

47

A

0	0	0	0
0	4.4	3.08	0
0	2.88	1.992	0
0	0	0	0

Input

9	7
6	4

Separate
output
array



Output

4.4	4.0
3.8	3.4

Input

9	7
6	4

Updates
occur in
place



Output

4.4	3.08
2.88	1.992

Which is correct?

49

Input

9	7
6	4



Output

4.4	3.08
2.88	1.992

Is this output incorrect?

- ❑ What is the stencil pattern?
 - ▶ Update alternatives
 - ▶ **2D Jacobi iteration**
- ❑ Implementing stencil with shift
- ❑ Stencil and cache optimizations
- ❑ Stencil and communication optimizations
- ❑ Recurrence

- ❑ Iterative codes are ones that update their data in steps
 - ▶ At each step, a new value of an element is computed using a formula based on other elements
 - ▶ Once all elements are updated, the computation proceeds to the next step or completes
- ❑ Iterative codes are most commonly found in computer simulations of physical systems for scientific and engineering applications
 - ▶ Computational fluid dynamics
 - ▶ Electromagnetics modeling
- ❑ They are often applied to solve partial differential equations
 - ▶ Jacobi iteration
 - ▶ Gauss-Seidel iteration
 - ▶ Successive over relaxation

- ❑ Stencils essentially define which elements are used in the update formula
- ❑ Because the data is organized in a regular manner, stencils can be applied across the data uniformly

Simple 2D Example

53

❑ Consider the following code

```
for k=1, 1000
```

```
  for i=1, N-2
```


```
    for j = 1, N-2
```

```
      a[i][j] = 0.25 * (a[i][j] + a[i-1][j]  
                        + a[i+1][j]  
                        + a[i][j-1]  
                        + a[i][j+1])
```

```
    }
```

```
  }
```

```
}
```




5-point stencil

Do you see anything interesting?

How would you parallelize?

- ❑ Consider a 2D array of elements
- ❑ Initialize each array element to some value
- ❑ At each step, update each array element to the arithmetic mean of its N, S, E, W neighbors
- ❑ Iterate until array values converge
- ❑ Here we are using a 4-point stencil



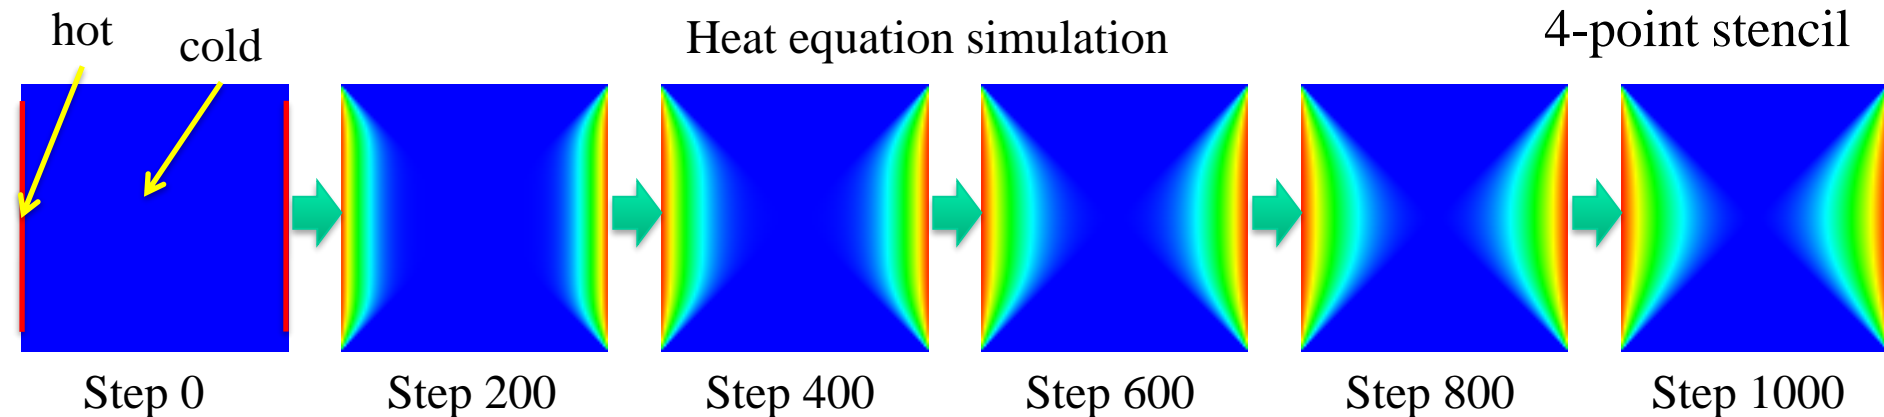

4-point stencil

- ❑ It is different from before because we want to update all array elements simultaneously ... How?

2-Dimension Jacobi Iteration

55

- ❑ Consider a 2D array of elements
- ❑ Initialize each array element to some value
- ❑ At each step, update each array element to the arithmetic mean of its N, S, E, W neighbors
- ❑ Iterate until array values converge



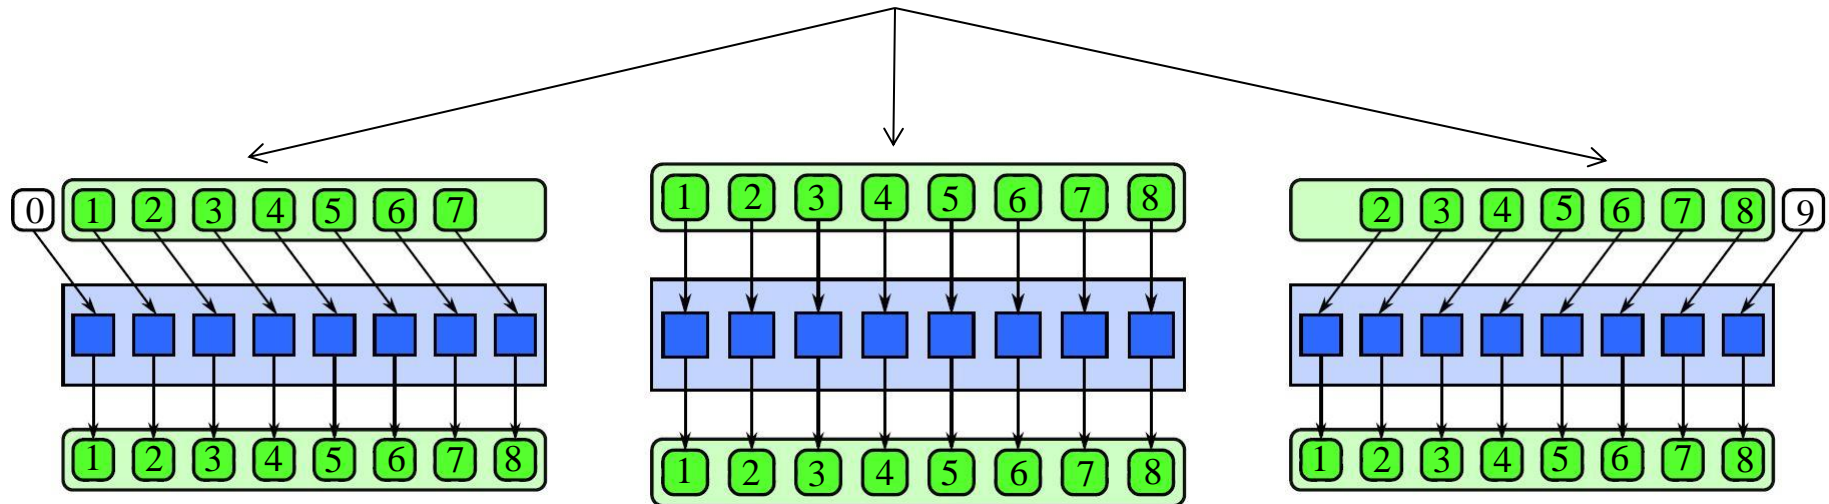
- ❑ What is the stencil pattern?
 - ▶ Update alternatives
 - ▶ 2D Jacobi iteration
- ❑ Implementing stencil with shift
- ❑ Stencil and cache optimizations
- ❑ Stencil and communication optimizations
- ❑ Recurrence

- ❑ One possible implementation of the stencil pattern includes shifting the input data
- ❑ For each offset in the stencil, we gather a new input vector by **shifting** the original input by the offset amount

Implementing Stencil with Shift

58

All input arrays are derived from the same original input array



- ❑ This implementation is only beneficial for one dimensional stencils or the memory-contiguous dimension of a multidimensional stencil
- ❑ Memory traffic to external memory is not reduced with shifts
- ❑ But, shifts allow vectorization of the data reads, which may reduce the total number of instructions

- ❑ Partitioning
- ❑ What is the stencil pattern?
 - ▶ Update alternatives
 - ▶ 2D Jacobi iteration
- ❑ Implementing stencil with shift
- ❑ Stencil and cache optimizations
- ❑ Stencil and communication optimizations
- ❑ Recurrence

- ❑ Assuming 2D array where rows are contiguous in memory...
 - ▶ Horizontally related data will tend to belong to the same cache line
 - ▶ Vertical offset accesses will most likely result in cache misses


- ❑ Assigning rows to cores:
 - ▶ Maximizes horizontal data locality
 - ▶ Assuming vertical offsets in stencil, this will create redundant reads of adjacent rows from each core
- ❑ Assigning columns to cores:
 - ▶ Redundantly read data from same cache line
 - ▶ Create false sharing as cores write to same cache line

- ❑ Assigning “strips” to each core can be a better solution
- ❑ **Strip-mining:** an optimization in a stencil computation that groups elements in a way that avoids redundant memory accesses and aligns memory accesses with cache lines

- ❑ A strip's size is a multiple of a cache line in width, and the height of the 2D array
- ❑ Strip widths are in increments of the cache line size so as to avoid false sharing and redundant reads
- ❑ Each strip is processed serially from top to bottom within each core

Stencil and Cache Optimizations

65



- ❑ What is the stencil pattern?
 - ▶ Update alternatives
 - ▶ 2D Jacobi iteration
- ❑ Implementing stencil with shift
- ❑ Stencil and cache optimizations
- ❑ Stencil and communication optimizations
- ❑ Recurrence

Stencil and Communication Optimizations

67

- ❑ When data is distributed, ghost cells must be **explicitly** communicated among nodes between loop iterations

- ❑ Darker cells are PE 0's ghost cells
- ❑ After first iteration of stencil computation
 - PE 0 must request PE 1 & PE 2's stencil results
 - PE 0 can perform another iteration of stencil

0	0	0	0
0	PE 0 1	PE 1 0	0
0	PE 2 1	PE 3 1	0
0	0	0	0

- ❑ Generally better to replicate ghost cells in each local memory and swap after each iteration than to share memory
 - ▶ Fine-grained sharing can lead to increased communication cost

- ❑ Halo: set of all ghost cells
- ❑ Halo must contain all neighbors needed for one iteration
- ❑ Larger halo (**deep halo**)
 - ▶ Trade off
 - less communications and more independence, but...
 - more redundant computation and more memory used
- ❑ **Latency Hiding**: Compute interior of stencil while waiting for ghost cell updates